


CSE113: Parallel Programming

Feb, 21, 2024

- **Topics:**

- Workstealing



Announcements

- HW 1 is completely graded
 - Let us know by the end of today if there are any issues!
- Starting on grading HW 2
 - Give us a week or so
- Grading midterm on Friday

Announcements

- HW 3 is out
 - Last day to turn it in is Tuesday Feb. 27
 - Plenty of time to get help
 - Office hours
 - Piazza
 - Etc.
- You should be able to do part 2 after today's lecture

Announcements

- Planned ~2 more lectures on concurrent data structures
 - Today:
 - Workstealing
 - Monday:
 - General Concurrent Sets

Announcements

- Green computing faculty talks!
 - Wednesdays at 11
 - Room E2 180
 - Good chance to see the academic job experience, feel free to join!

Previous quiz + Review


Previous quiz + Review

Input/output queues use atomic increments and decrements to protect against threads that are trying to concurrently enqueue and dequeue

☐ True

☐ False

Implementation



Thread 0:
enq(6);


Thread 1:
enq(7);

What happens if a thread wants
to add an element?

Think concurrently:

```
reserved_index = atomic_fetch_add(&end, 1);
```


Implementation



Thread 0:
enq(6);

Thread 1:
enq(7);

What happens if a thread wants to add an element?

Think concurrently:

```
reserved_index = atomic_fetch_add(&end, 1);
```

Implementation

*does it matter which order
threads add their data?*

reserved
T0



reserved
T1

end

Thread 0:
`enq(6);`

Thread 1:
`enq(7);`

What happens if a thread wants
to add an element?

Think concurrently:

```
reserved_index = atomic_fetch_add(&end, 1);
```

Implementation

*does it matter which order
threads add their data?*

reserved
T0



end

Thread 0:
enq(6);

Thread 1:
enq(7);

What happens if a thread wants
to add an element?

Think concurrently:

```
reserved_index = atomic_fetch_add(&end, 1);
```

Implementation

*does it matter which order
threads add their data? No!
Because there are no deqs!*

reserved
T0



end

Thread 0:
enq(6);

Thread 1:
enq(7);

What happens if a thread wants
to add an element?


Think concurrently:

```
reserved_index = atomic_fetch_add(&end, 1);
```

Previous quiz + Review

Write a few questions about the pros and cons of using a specialized concurrent queue (e.g. an IO queue) and a fully general concurrent queue.

```
class InputOutputQueue {  
    private:  
        atomic_int front;  
        atomic_int end;  
        int list[SIZE];  
  
    public:  
        InputOutputQueue() {  
            front = end = 0;  
        }  
  
        void enq(int x) {  
            int reserved_index = atomic_fetch_add(&end, 1);  
            list[reserved_index] = x;  
        }  
  
        void deq() {  
            int reserved_index = atomic_fetch_add(&front, 1);  
            return list[reserved_index];  
        }  
  
        int size() {  
            return end.load() - front.load();  
        }  
}
```



```
class ProdConsQueue {  
    private:  
        atomic_int head;  
        atomic_int tail;  
        int buffer[SIZE];  
  
    public:  
        void enq(int x) {  
            // wait for there to be room  
            // store value at head  
            // increment head  
        }  
        int deq() {  
            // wait while queue is empty  
            // get value at tail  
            // increment tail  
        }  
}
```

Previous quiz + Review

The performance of an application using a producer-consumer queue depends most on:

- ☐ If the queue is implemented using mutex or not
- ☐ The rate at which the consumer enqueues elements
- ☐ The rate at which the producer enqueues elements

Previous quiz + Review

A circular buffer is:

- ☐ A useful data representation for fixed-length queues
- ☐ Part of the C++ standard library
- ☐ A special type of memory that is organized in circular patterns

Producer Consumer Queues

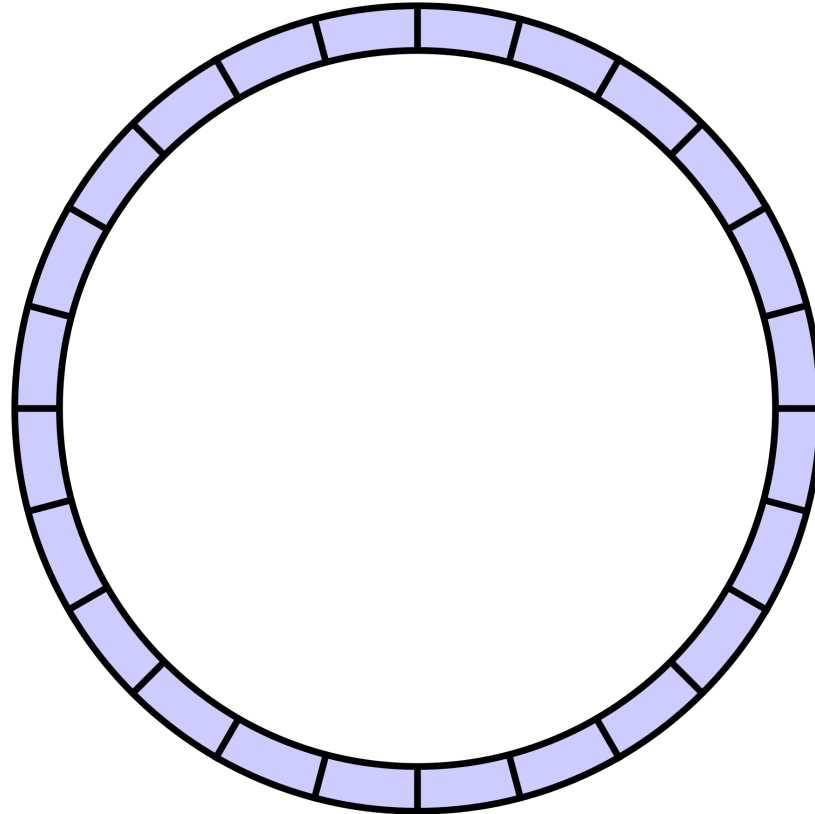
- Start with a fixed size array



We will use what is called a *circular buffer method*

Producer Consumer Queues

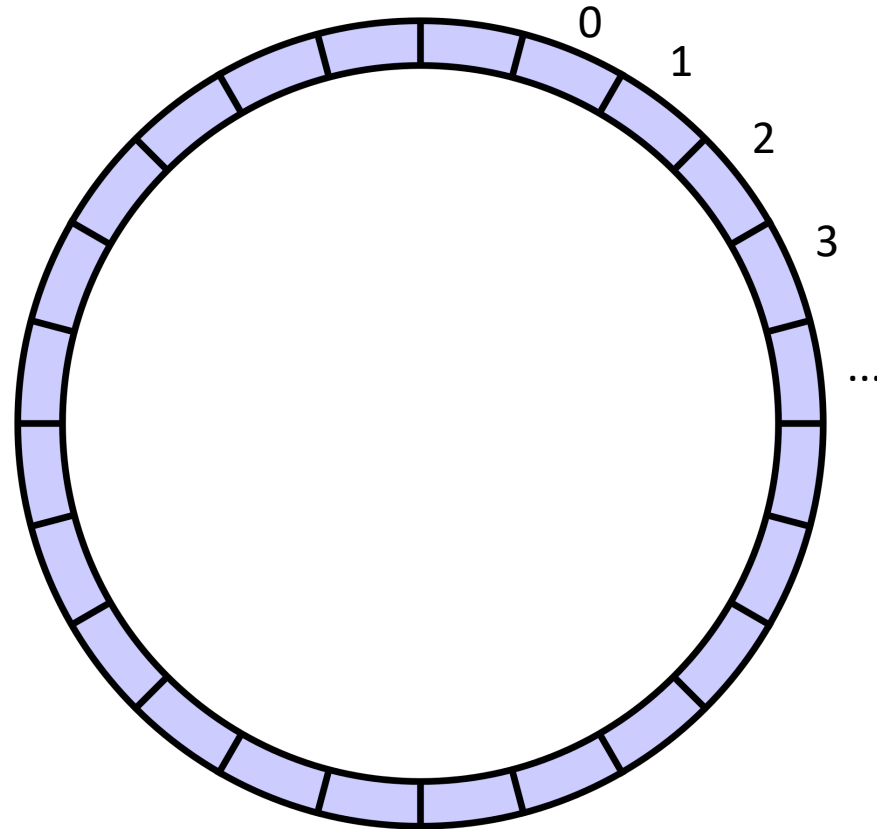
- Start with a fixed size array



conceptually it is a circle

Producer Consumer Queues


- Start with a fixed size array



conceptually it is a circle

Producer Consumer Queues

- Start with a fixed size array



indexes will
circulate in
order and
wrap around

conceptually it is a circle

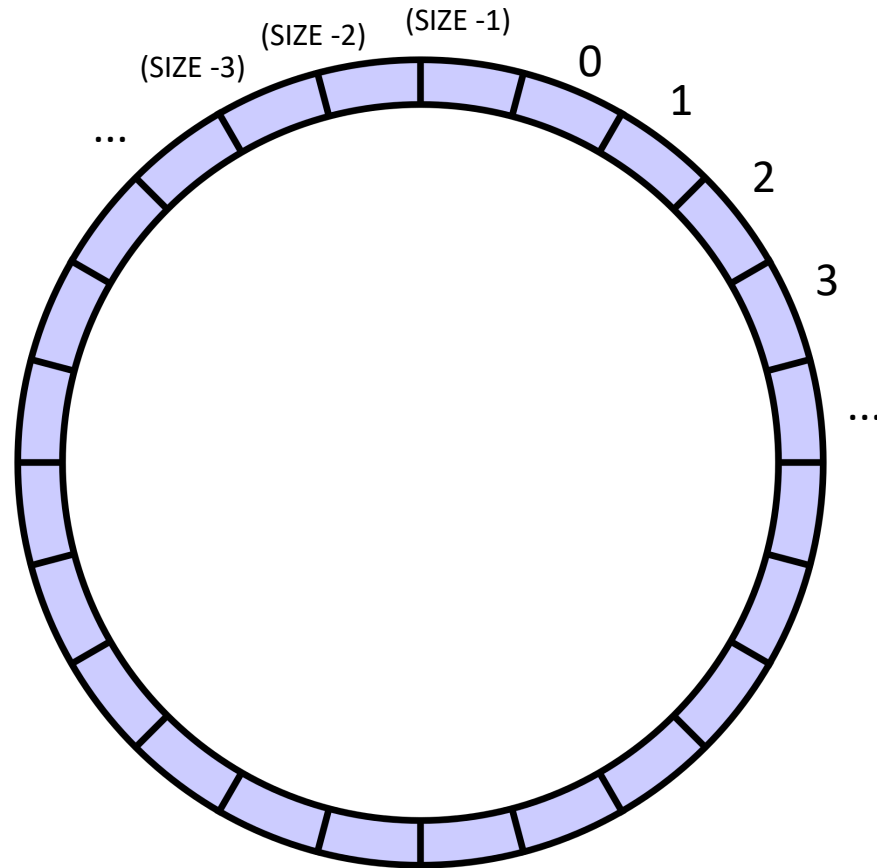
Producer Consumer Queues

- Start with a fixed size array

we will assume modular arithmetic:

if $x = (\text{SIZE} - 1)$ then
 $x + 1 == 0$;

conceptually it is a circle



indexes will
circulate in
order and
wrap around

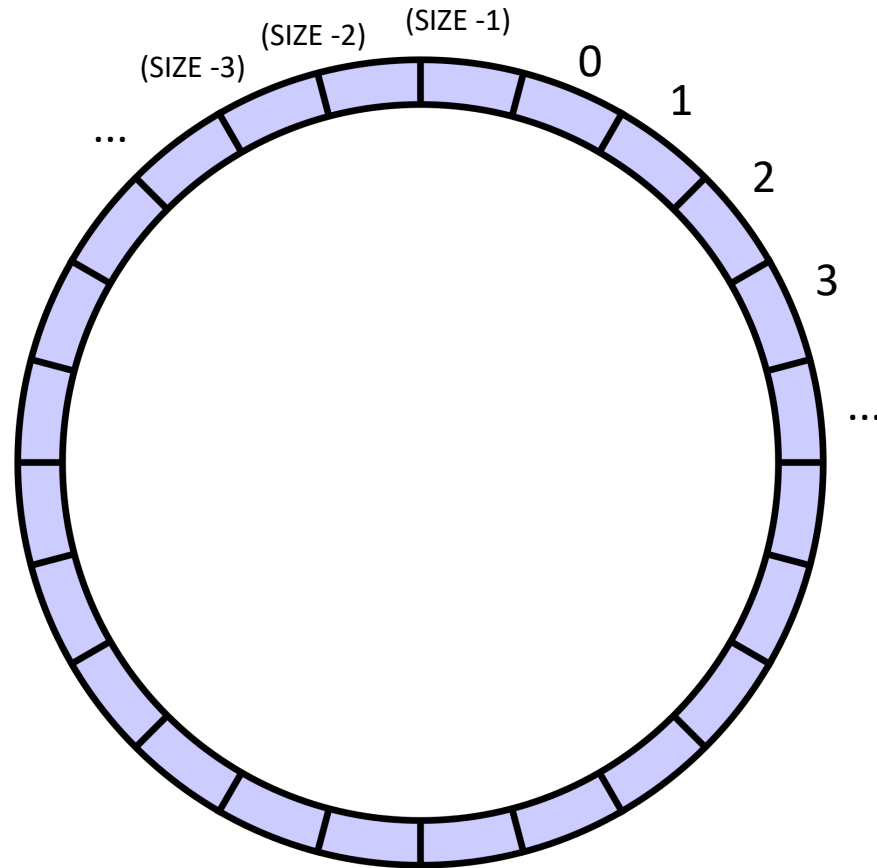
Producer Consumer Queues

- Start with a fixed size array

Two variables to keep track of
where to deq and enq:

head and tail

conceptually it is a circle



indexes will
circulate in
order and
wrap around

Producer Consumer Queues


- Start with a fixed size array

Two variables to keep track of where to deq and enq:

head and tail:

enq to the head, deq from the tail

conceptually it is a circle



indexes will circulate in order and wrap around

Producer Consumer Queues

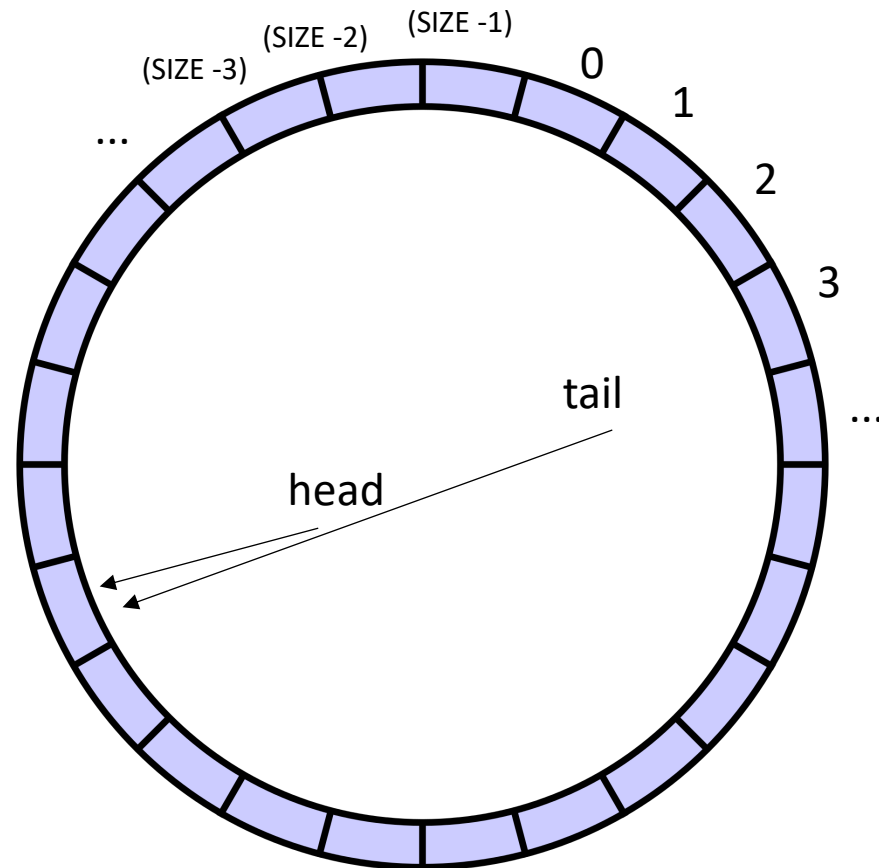
- Start with a fixed size array

Two variables to keep track of
where to deq and enq:

head and tail

Empty queue is when
 $\text{head} == \text{tail}$

conceptually it is a circle



indexes will
circulate in
order and
wrap around

Producer Consumer Queues

- Start with a fixed size array


Two variables to keep track of where to deq and enq:

head and tail

Empty queue is when
 $\text{head} == \text{tail}$

Full queue is when
 $\text{head} == \text{tail}?$

conceptually it is a circle



indexes will
circulate in
order and
wrap around

Producer Consumer Queues

- Start with a fixed size array

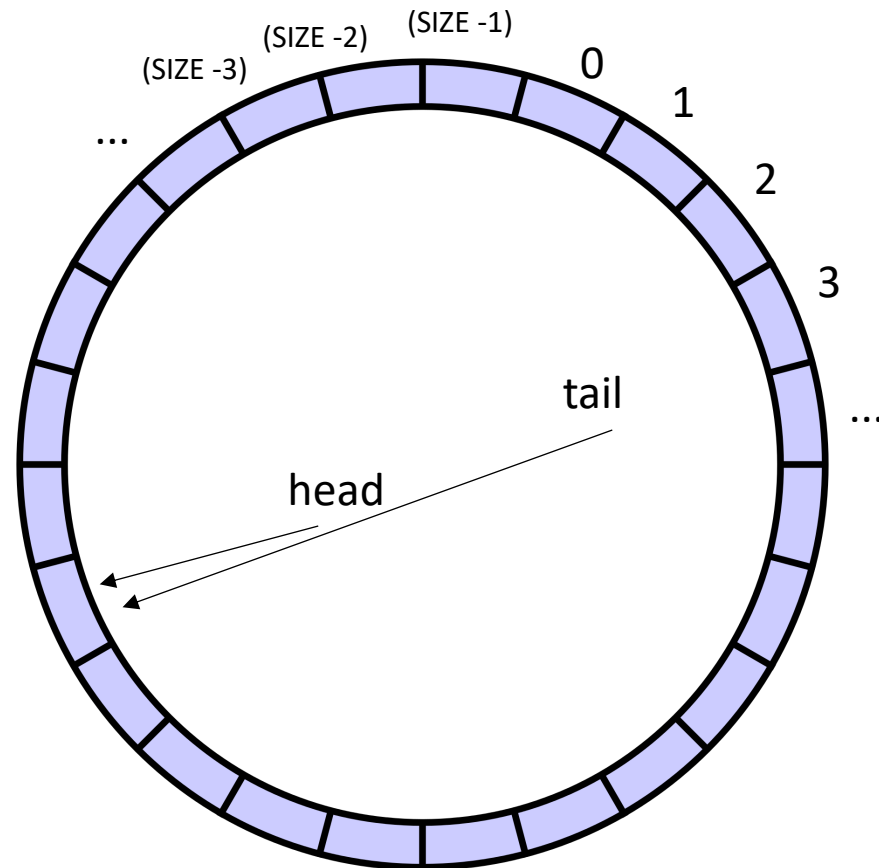
Two variables to keep track of where to deq and enq:

head and tail

Empty queue is when
 $\text{head} == \text{tail}$

Full queue is when
 $\text{head} == \text{tail}?$

conceptually it is a circle



indexes will
circulate in
order and
wrap around

but then
how to tell
full queue from
empty?

Producer Consumer Queues

- Start with a fixed size array


Two variables to keep track of where to deq and enq:

head and tail

Empty queue is when
 $\text{head} == \text{tail}$

Full queue is when
 $\text{head} + 1 == \text{tail}$


conceptually it is a circle



indexes will
circulate in
order and
wrap around

wasting one
location, but its okay...

review



```
class ProdConsQueue {  
    private:  
        atomic_int head;  
        atomic_int tail;  
        int buffer[SIZE];  
  
    public:  
        void enq(int x) {  
            // wait for there to be room  
            // store value at head  
            // increment head  
        }  
        int deq() {  
            // wait while queue is empty  
            // get value at tail  
            // increment tail  
        }  
}
```

On to new stuff!

- Work stealing

adds two arrays

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

adds elements with neighbors

```
for (int i = 0; i < SIZE; i++) {  
    a[i] += a[i+1]  
}
```


are they the same if you traverse them backwards?

adds two arrays

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

```
for (int i = SIZE-1; i >= 0; i--) {  
    a[i] = b[i] + c[i];  
}
```

adds elements with neighbors

```
for (int i = 0; i < SIZE; i++) {  
    a[i] += a[i+1]  
}
```

```
for (int i = SIZE-1; i >= 0; i--) {  
    a[i] += a[i+1]  
}
```

are they the same if you traverse them backwards?

adds two arrays

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

```
for (int i = SIZE-1; i >= 0; i--) {  
    a[i] = b[i] + c[i];  
}
```

adds elements with neighbors

```
for (int i = 0; i < SIZE; i++) {  
    a[i] += a[i+1]  
}
```

```
for (int i = SIZE-1; i >= 0; i--) {  
    a[i] += a[i+1]  
}
```

No!

adds two arrays

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

what about a random order?

```
for (pick i randomly) {  
    a[i] = b[i] + c[i];  
}
```

adds elements with neighbors

```
for (int i = 0; i < SIZE; i++) {  
    a[i] += a[i+1]  
}
```

```
for (pick i randomly) {  
    a[i] += a[i+1]  
}
```

adds two arrays

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

what about a random order?

```
for (pick i randomly) {  
    a[i] = b[i] + c[i];  
}
```

adds elements with neighbors

```
for (int i = 0; i < SIZE; i++) {  
    a[i] += a[i+1]  
}
```

```
for (pick i randomly) {  
    a[i] += a[i+1]  
}
```

No!

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

These are **DOALL** loops:

- Loop iterations are independent
- You can do them in ANY order and get the same results

```
for (int i = 0; i < SIZE; i++) {  
    a[i] = b[i] + c[i];  
}
```

These are **DOALL** loops:

- Loop iterations are independent
- You can do them in ANY order and get the same results
- Most importantly: you can do the iterations in parallel!
- Assign each thread a set of indices to compute

DOALL Loops

- Given a nest of For loops, can we make the outer-most loop parallel?
 - Safely
 - Efficiently

DOALL Loops

- We will consider a special type of for loop, common in scientific applications:
 - Operates on N dimensional arrays (only side-effects are array writes)
 - Array bases are disjoint and constant
 - Bounds, indexes are a function of loop variables, input variables and constants
 - Loops Increment by 1

```
for (int i = 0; i < dim1; i++) {  
    for (int j = 0; j < dim3; j++) {  
        for (int k = 0; k < dim2; k++) {  
            a[i][j] += b[i][k] * c[k][j];  
        }  
    }  
}
```

matrix multiplication
example

DOALL Loops

- We will consider a special type of for loop, common in scientific applications:
 - Operates on N dimensional arrays (only side-effects are array writes)
 - Array bases are disjoint and constant
 - Bounds, indexes are a function of loop variables, input variables and constants
 - Loops Increment by 1

DOALL Loops

- Given a nest of ***candidate*** For loops, determine if we can we make the outer-most loop parallel?
 - Safely
 - efficiently
- Criteria: every iteration of the outer-most loop must be *independent*
 - The loop can execute in any order, and produce the same result

Safety Criteria

- How do we check this?
 - If the property doesn't hold then there exists 2 iterations, such that if they are re-ordered, it causes different outcomes for the loop.
 - **Write-Write conflicts:** two distinct iterations write different values to the same location
 - **Read-Write conflicts:** two distinct iterations where one iteration reads from the location written to by another iteration.

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*
- the loop must produce the same result for any order of the iterations

```
for (i = 0; i < size; i++) {  
    a[index(i)] = loop(i);  
}
```

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*
- the loop must produce the same result for any order of the iterations

```
for (i = 0; i < size; i++) {  
    a[index(i)] = loop(i);  
}
```

index calculation based on the loop variable

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*
- the loop must produce the same result for any order of the iterations

```
for (i = 0; i < size; i++) {  
    a[index(i)] = loop(i);  
}
```

index calculation based on the loop variable
Computation to store in the memory location

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*
- the loop must produce the same result for any order of the iterations

```
for (i = 0; i < size; i++) {  
    a[index(i)] = loop(i);  
}
```

Write-write conflicts:

for two distinct iteration variables:

$i_x \neq i_y$

Check:

$\text{index}(i_x) \neq \text{index}(i_y)$

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*
- the loop must produce the same result for any order of the iterations

```
for (i = 0; i < size; i++) {  
    a[index(i)] = loop(i);  
}
```

Write-write conflicts:

for two distinct iteration variables:

$i_x \neq i_y$

Check:

$\text{index}(i_x) \neq \text{index}(i_y)$

Why?

Because if

$\text{index}(i_x) == \text{index}(i_y)$

then:

$a[\text{index}(i_x)]$ will equal
either $\text{loop}(i_x)$ or $\text{loop}(i_y)$
depending on the order

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*

```
for (i = 0; i < size; i++) {  
    a[write_index(i)] = a[read_index(i)] + loop(i);  
}
```

Read-write conflicts:

for two distinct iteration variables:

$i_x \neq i_y$

Check:

$\text{write_index}(i_x) \neq \text{read_index}(i_y)$

Safety Criteria

- Criteria: every iteration of the outer-most loop must be *independent*

```
for (i = 0; i < size; i++) {  
    a[write_index(i)] = a[read_index(i)] + loop(i);  
}
```

Read-write conflicts:

for two distinct iteration variables:

$i_x \neq i_y$

Check:

$\text{write_index}(i_x) \neq \text{read_index}(i_y)$

Why?

if i_x iteration happens first, then
iteration i_y reads an updated value.

if i_y happens first, then it reads the
original value

Examples:

```
for (i = 0; i < 128; i++) {  
    a[i] = a[i]*2;  
}
```

Examples:

```
for (i = 0; i < 128; i++) {  
    a[i] = a[i]*2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i] = a[0]*2;  
}
```

Examples:

```
for (i = 0; i < 128; i++) {  
    a[i] = a[i]*2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i] = a[0]*2;  
}
```

```
for (i = 1; i < 128; i++) {  
    a[i] = a[0]*2;  
}
```

Examples:

```
for (i = 0; i < 128; i++) {  
    a[i] = a[i] * 2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i] = a[0] * 2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i%64] = a[i] * 2;  
}
```

```
for (i = 1; i < 128; i++) {  
    a[i] = a[0] * 2;  
}
```

Examples:

```
for (i = 0; i < 128; i++) {  
    a[i] = a[i] * 2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i] = a[0] * 2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i%64] = a[i] * 2;  
}
```

```
for (i = 1; i < 128; i++) {  
    a[i] = a[0] * 2;  
}
```

```
for (i = 0; i < 128; i++) {  
    a[i%64] = a[i+64] * 2;  
}
```

Parallel Schedules

- Consider the following program:

There are 3 arrays: a , b , c .

We want to compute

```
for (int i = 0; i < SIZE; i++) {  
    c[i] = a[i] + b[i];  
}
```

Is this a DOALL loop?

Parallel Schedules

- Consider the following program:

There are 3 arrays: a , b , c .

We want to compute


```
for (int i = 0; i < SIZE; i++) {  
    c[i] = a[i] + b[i];  
}
```

Is this a DOALL loop?

How should we parallelize it?

Parallel Schedules

array a




+ + + + + + + +

array b



= = = = = = = =

array c

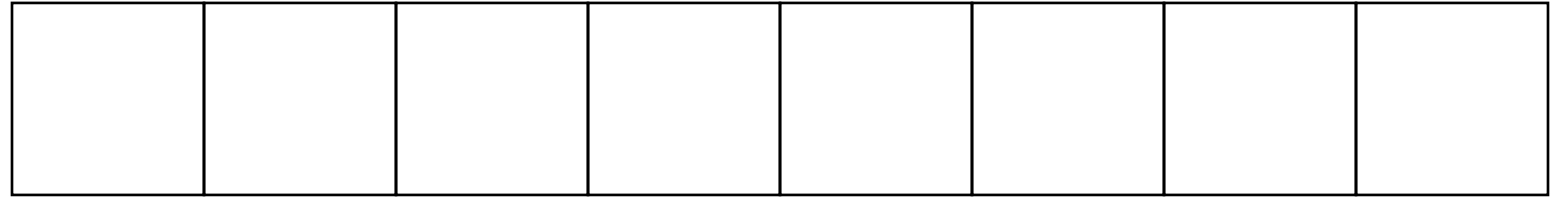


Parallel Schedules

Computation
can easily be
divided into
threads

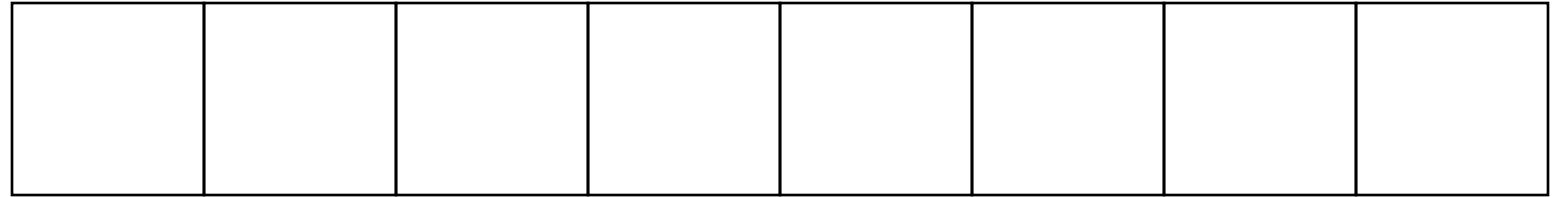
Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

array a



+ + + + + + + +

array b



= = = = = = = =

array c




Parallel Schedules

Computation
can easily be
divided into
threads


Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

array a



+ + + + + + + +

array b



= = = = = = = =

array c




Parallel Schedules

Computation
can easily be
divided into
threads


Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

array a



+ + + + + + + +

array b



= = = = = = = =

array c



Parallel Schedules

array a




+ + + + + + + +

array b



= = = = = = = =

array c



Computation
can easily be
divided into
threads

Thread 0 - Blue
Thread 1 - Yellow
Thread 2 - Green
Thread 3 - Orange

Parallel Schedules

- Which one is more efficient?

Parallel Schedules

- Which one is more efficient?
- These are called Parallel Schedules for DOALL Loops
- We will discuss several of them.

Schedule

- DOALL Loops
- **Parallel Schedules:**
 - **Static**
 - Global Worklists
 - Local Worklists

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
        // Each iteration takes roughly  
        // equal time  
    }  
    ...  
}
```

0	1	2	3	4	5	6	7		SIZE -1
---	---	---	---	---	---	---	---	--	---------

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
        // Each iteration takes roughly  
        // equal time  
    }  
    ...  
}
```

say $SIZE / NUM_THREADS = 4$


0	1	2	3	4	5	6	7		SIZE - 1
---	---	---	---	---	---	---	---	--	----------

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
        // Each iteration takes roughly  
        // equal time  
    }  
    ...  
}
```

say $SIZE / NUM_THREADS = 4$



Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
        // Each iteration takes roughly  
        // equal time  
    }  
    ...  
}
```

make a new function with the for loop inside. Pass all needed variables as arguments. Take an extra argument for a thread id

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
    // Each iteration takes roughly  
    // equal time  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads)  
{  
      
    for (int x = 0; x < SIZE; x++) {  
        // work based on x  
    }  
}
```

make a new function with the for loop inside. Pass all needed variables as arguments. Take an extra argument for a thread id

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
    // Each iteration takes roughly  
    // equal time  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    for (int x = 0; x < SIZE; x++) {  
        // work based on x  
    }  
}
```

determine chunk size in new function

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int x = 0; x < SIZE; x++) {  
    // Each iteration takes roughly  
    // equal time  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    int start = chunk_size * tid;  
    int end = start + chunk_size;  
    for (int x = start; x < end; x++) {  
        // work based on x  
    }  
}
```

Set new loop bounds

Static schedule

- Works well when loop iterations take similar amounts of time

```
void foo() {  
    ...  
    for (int t = 0; t < NUM_THREADS; t++) {  
        spawn(parallel_loop(..., t, NUM_THREADS))  
    }  
    join();  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    int start = chunk_size * tid;  
    int end = start + chunk_size;  
    for (int x = start; x < end; x++) {  
        // work based on x  
    }  
}
```

You will need to adapt the thread spawn, join
to C++

Spawn threads

Static schedule



- Example, 2 threads/cores, array of size 8

0	1	2	3	4	5	6	7
---	---	---	---	---	---	---	---

chunk_size = ?

0: start = ? 1: start = ?


0: end = ? 1: end = ?



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule

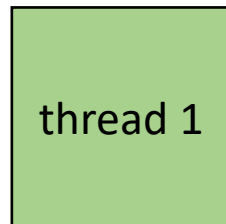

- Example, 2 threads/cores, array of size 8



chunk_size = 4

0: start = 0 1: start = 4

0: end = 4 1: end = 8



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule



- Example, 2 threads/cores, array of size 9

0	1	2	3	4	5	6	7	8
---	---	---	---	---	---	---	---	---

chunk_size = ?

0: start = ? 1: start = ?


0: end = ? 1: end = ?



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule



- Example, 2 threads/cores, array of size 9



chunk_size = 4

0: start = 0 1: start = 4


0: end = 4 1: end = 8



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule



- Example, 2 threads/cores, array of size 9



chunk_size = 4

0: start = 0 1: start = 4

0: end = 4 1: end = 8




```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    if (tid == num_threads - 1) {
        end = SIZE;
    }
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule

last thread gets more work



- Example, 2 threads/cores, array of size 9



chunk_size = 4

0: start = 0 1: start = 4

0: end = 4 1: end = ?




```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    if (tid == num_threads - 1) {
        end = SIZE;
    }
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule

- Example, 2 threads/cores, array of size 9

last thread gets more work

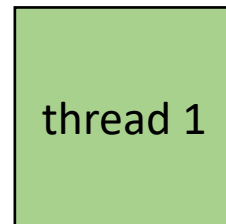
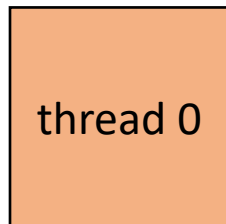
What is the worst case?



chunk_size = 4

0: start = 0 1: start = 4

0: end = 4 1: end = 9




```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size = SIZE / NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    if (tid == num_threads - 1) {
        end = SIZE;
    }
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```


End example

Static schedule


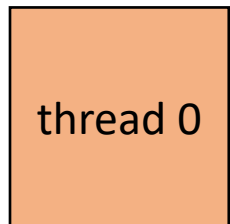
- Example, 2 threads/cores, array of size 9



chunk_size = 4

0: start = 0 1: start = 4

0: end = 4 1: end = 8

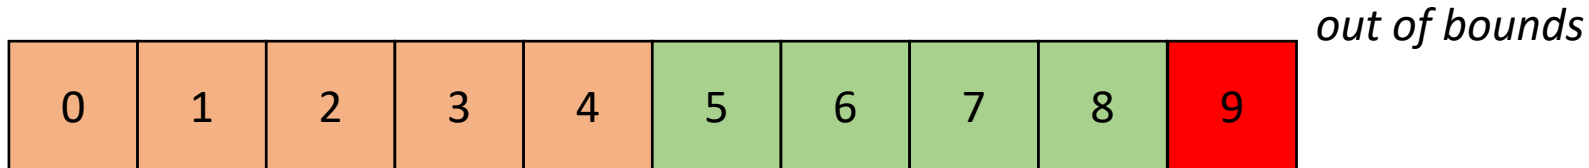


ceiling division, this will distribute uneven work in the last thread to all other threads

```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size =
        (SIZE+(NUM_THREADS-1))/NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule

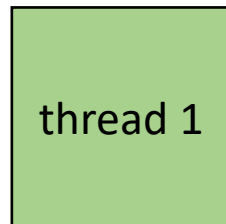

- Example, 2 threads/cores, array of size 9



chunk_size = 5

0: start = 0 1: start = 5


0: end = 5 1: end = 10



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size =
        (SIZE+(NUM_THREADS-1))/NUM_THREADS;
    int start = chunk_size * tid;
    int end = start + chunk_size;
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule


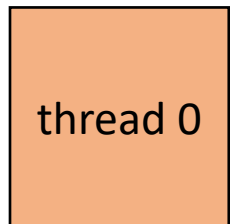
- Example, 2 threads/cores, array of size 9



chunk_size = 5

0: start = 0 1: start = 5


0: end = 5 1: end = 10



```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size =
        (SIZE + (NUM_THREADS - 1)) / NUM_THREADS;
    int start = chunk_size * tid;
    int end =
        min(start + chunk_size, SIZE)
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

Static schedule


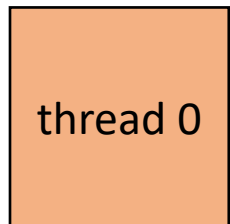
- Example, 2 threads/cores, array of size 9



chunk_size = 5

0: start = 0 1: start = 5

0: end = 5 1: end = 9



most threads do equal amounts of work, last thread may do less.

Which one is better/worse?

Max slowdown for last thread does all the extra work?

Max slowdown for ceiling?

```
void parallel_loop(..., int tid, int num_threads)
{
    int chunk_size =
        (SIZE + (NUM_THREADS - 1)) / NUM_THREADS;
    int start = chunk_size * tid;
    int end =
        min(start + chunk_size, SIZE)
    for (int x = start; x < end; x++) {
        // work based on x
    }
}
```

End example

Schedule

- DOALL Loops
- **Parallel Schedules:**
 - Static
 - **Global Worklists**
 - Local Worklists

Irregular parallelism in loops

- Tasks are not balanced
- Appears in lots of emerging workloads

Irregular parallelism in loops

- Tasks are not balanced
- Appears in lots of emerging workloads

social network analytics where threads are parallel across users



Irregular parallelism in loops

- Independent iterations have different amount of work to compute
- Threads with longer tasks take longer to compute.
- Threads with shorter tasks are under utilized.

```
for (x = 0; x < SIZE; x++) {  
    for (y = x; y < SIZE; y++) {  
        a[x,y] = b[x,y] + c[x,y];  
    }  
}
```

irregular (or unbalanced) parallelism:
each x iteration performs different
amount of work.

Irregular parallelism in loops

- Calculate imbalance cost if x is chunked:
 - Thread 1 takes iterations $0 - \text{SIZE}/2$
 - Thread 2 takes iterations $\text{SIZE}/2 - \text{SIZE}$

```
for (x = 0; x < SIZE; x++) {  
    for (y = x; y < SIZE; y++) {  
        a[x,y] = b[x,y] + c[x,y];  
    }  
}
```

Irregular parallelism in loops

- Calculate imbalance cost if x is chunked:
 - Thread 1 takes iterations 0 - SIZE/2
 - Thread 2 takes iterations SIZE/2 - SIZE

Calculate how much total work:

$$\text{total_work} = \sum_{n=0}^{\text{SIZE}} n$$

```
for (x = 0; x < SIZE; x++) {  
    for (y = x; y < SIZE; y++) {  
        a[x,y] = b[x,y] + c[x,y];  
    }  
}
```

Irregular parallelism in loops

- Calculate imbalance cost if x is chunked:
 - Thread 1 takes iterations 0 - SIZE/2
 - Thread 2 takes iterations SIZE/2 - SIZE

```
for (x = 0; x < SIZE; x++) {  
    for (y = x; y < SIZE; y++) {  
        a[x,y] = b[x,y] + c[x,y];  
    }  
}
```

Calculate how much total work:

$$\text{total_work} = \sum_{n=0}^{\text{SIZE}} n$$

Calculate work done by first thread:

$$\text{t1_work} = \sum_{n=0}^{\text{SIZE}/2} n$$

Irregular parallelism in loops

- Calculate imbalance cost if x is chunked:
 - Thread 1 takes iterations 0 - SIZE/2
 - Thread 2 takes iterations SIZE/2 - SIZE

```
for (x = 0; x < SIZE; x++) {  
    for (y = x; y < SIZE; y++) {  
        a[x,y] = b[x,y] + c[x,y];  
    }  
}
```

Calculate how much total work:

$$\text{total_work} = \sum_{n=0}^{\text{SIZE}} n$$

Calculate work done by first thread:

$$\text{t1_work} = \sum_{n=0}^{\text{SIZE}/2} n$$

Calculate work done by second thread:

$$\text{t2_work} = \text{total_work} - \text{t1_work}$$

Irregular parallelism in loops

Example: SIZE = 64

total_work = 2016

t2_work = 496

t1_work = 1520

t1 does ~3x more work than t2

Only provides ~1.3x speedup

Potential solution:

Have T1 do only ¼ of the iterations

Gives a better speedup of 1.77x

Not a feasible solution because often times load imbalance is not given by a static equation on loop bounds!

Calculate how much total work:

$$\text{total_work} = \sum_{n=0}^{\text{SIZE}} n$$

Calculate work done by first thread:

$$\text{t1_work} = \sum_{n=0}^{\text{SIZE}/2} n$$

Calculate work done by second thread:

$$\text{t2_work} = \text{total_work} - \text{t1_work}$$

Work stealing

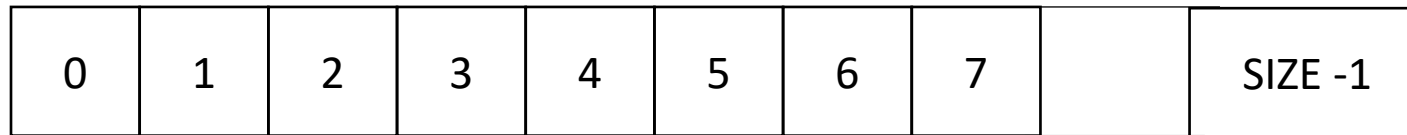
- Tasks are dynamically assigned to threads.

Work stealing - global implicit worklist



- Pros
 - Simple to implement
- Cons:
 - High contention on global counter
 - Potentially bad memory locality.

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




cannot color initially!





Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



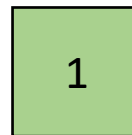

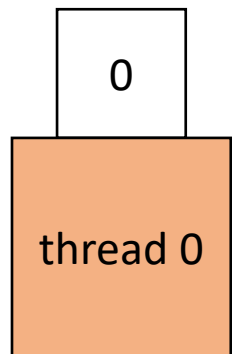

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



Work stealing - global implicit worklist

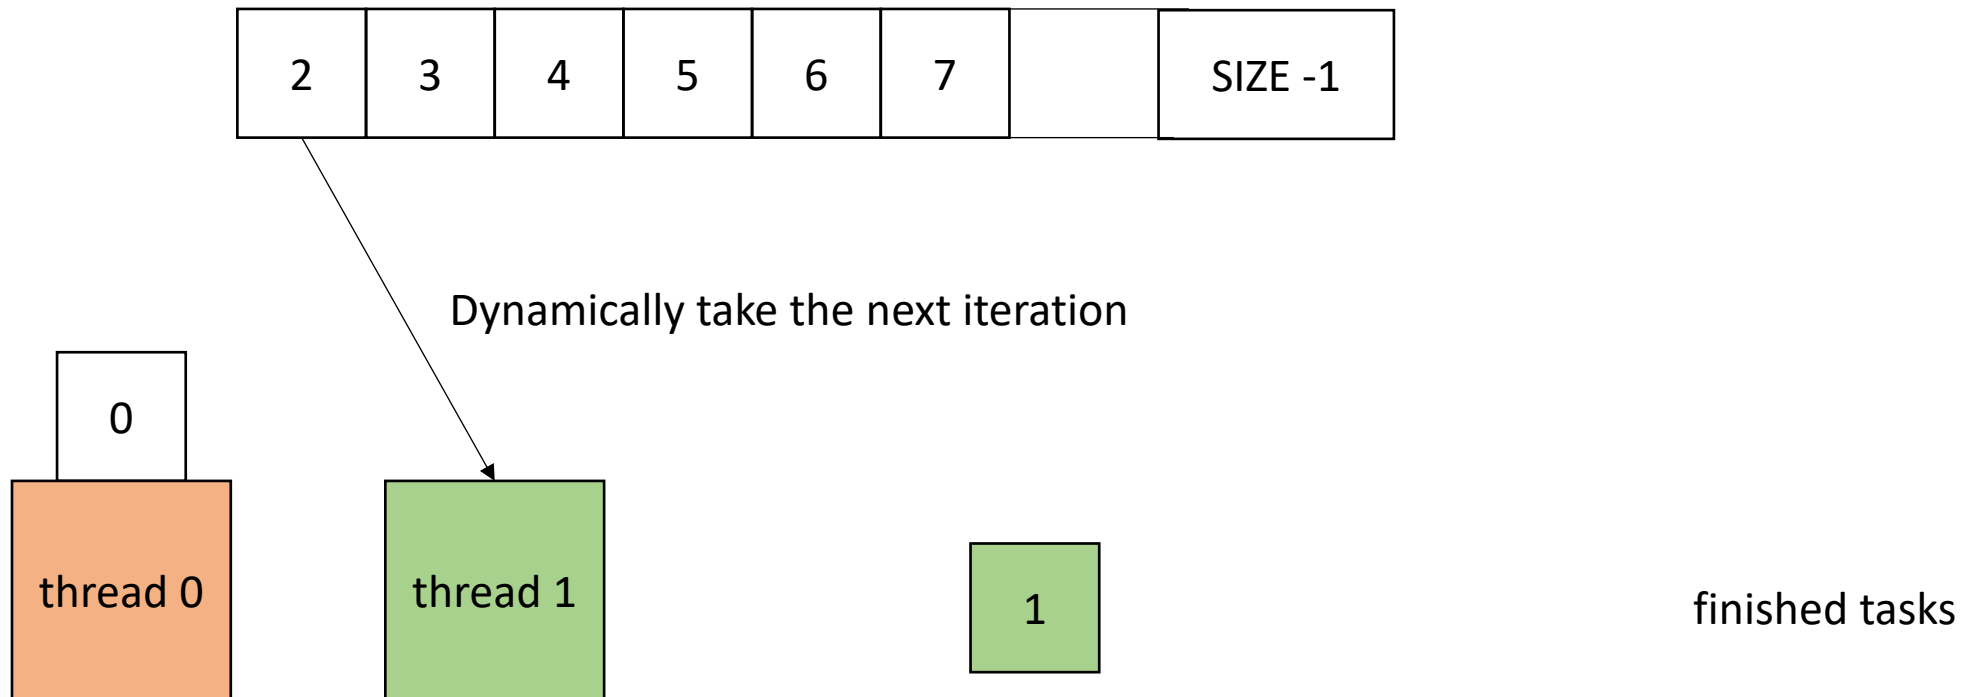
- Global worklist: threads take tasks (iterations) dynamically



finished tasks





Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



Work stealing - global implicit worklist

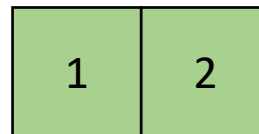

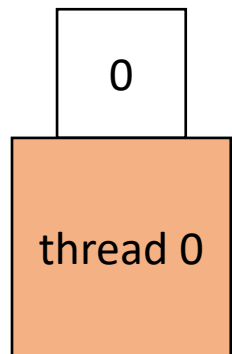

- Global worklist: threads take tasks (iterations) dynamically



finished tasks

Work stealing - global implicit worklist


- Global worklist: threads take tasks (iterations) dynamically



finished tasks





Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



Work stealing - global implicit worklist


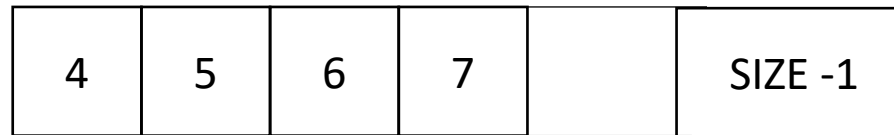
- Global worklist: threads take tasks (iterations) dynamically



finished tasks


Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically






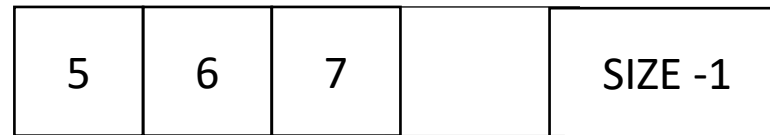
Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



finished tasks

End example

Work stealing - global implicit worklist

- How to implement

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
    ...  
}
```

Work stealing - global implicit worklist

- How to implement

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
void parallel_loop(...) {  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

Replicate code in a new function. Pass all needed variables as arguments.

Work stealing - global implicit worklist

- How to implement

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
atomic_int x(0);  
void parallel_loop(...) {  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

move loop variable to be a global atomic variable

Work stealing - global implicit worklist

- How to implement

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
atomic_int x(0);  
void parallel_loop(...) {  
    for (int local_x = ??  
        local_x < SIZE;  
        local_x = ??) {  
        // dynamic work based on x  
    }  
}
```

change loop bounds in new function to use a local variable using global variable.

Work stealing - global implicit worklist

- How to implement

These must be atomic updates!

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
atomic_int x(0);  
void parallel_loop(...) {  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
        // dynamic work based on x  
    }  
}
```

change loop bounds in new function to use a local variable using global variable.

Work stealing - global implicit worklist

- How to implement

```
void foo() {  
    ...  
    for (t = 0; x < THREADS; t++) {  
        spawn(parallel_loop);  
    }  
    join();  
    ...  
}
```


```
atomic_int x(0);  
void parallel_loop(...) {  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
        // dynamic work based on x  
    }  
}
```

Spawn threads in original function and join them afterwards

You will have to change to C++ syntax for the homework!

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically





x: 0

0 - local_x - UNDEF


1 - local_x - UNDEF

```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```



Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




x: 2
0 - local_x - 0
1 - local_x - 1


```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




x: 2
0 - local_x - 0
1 - local_x - 1




```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



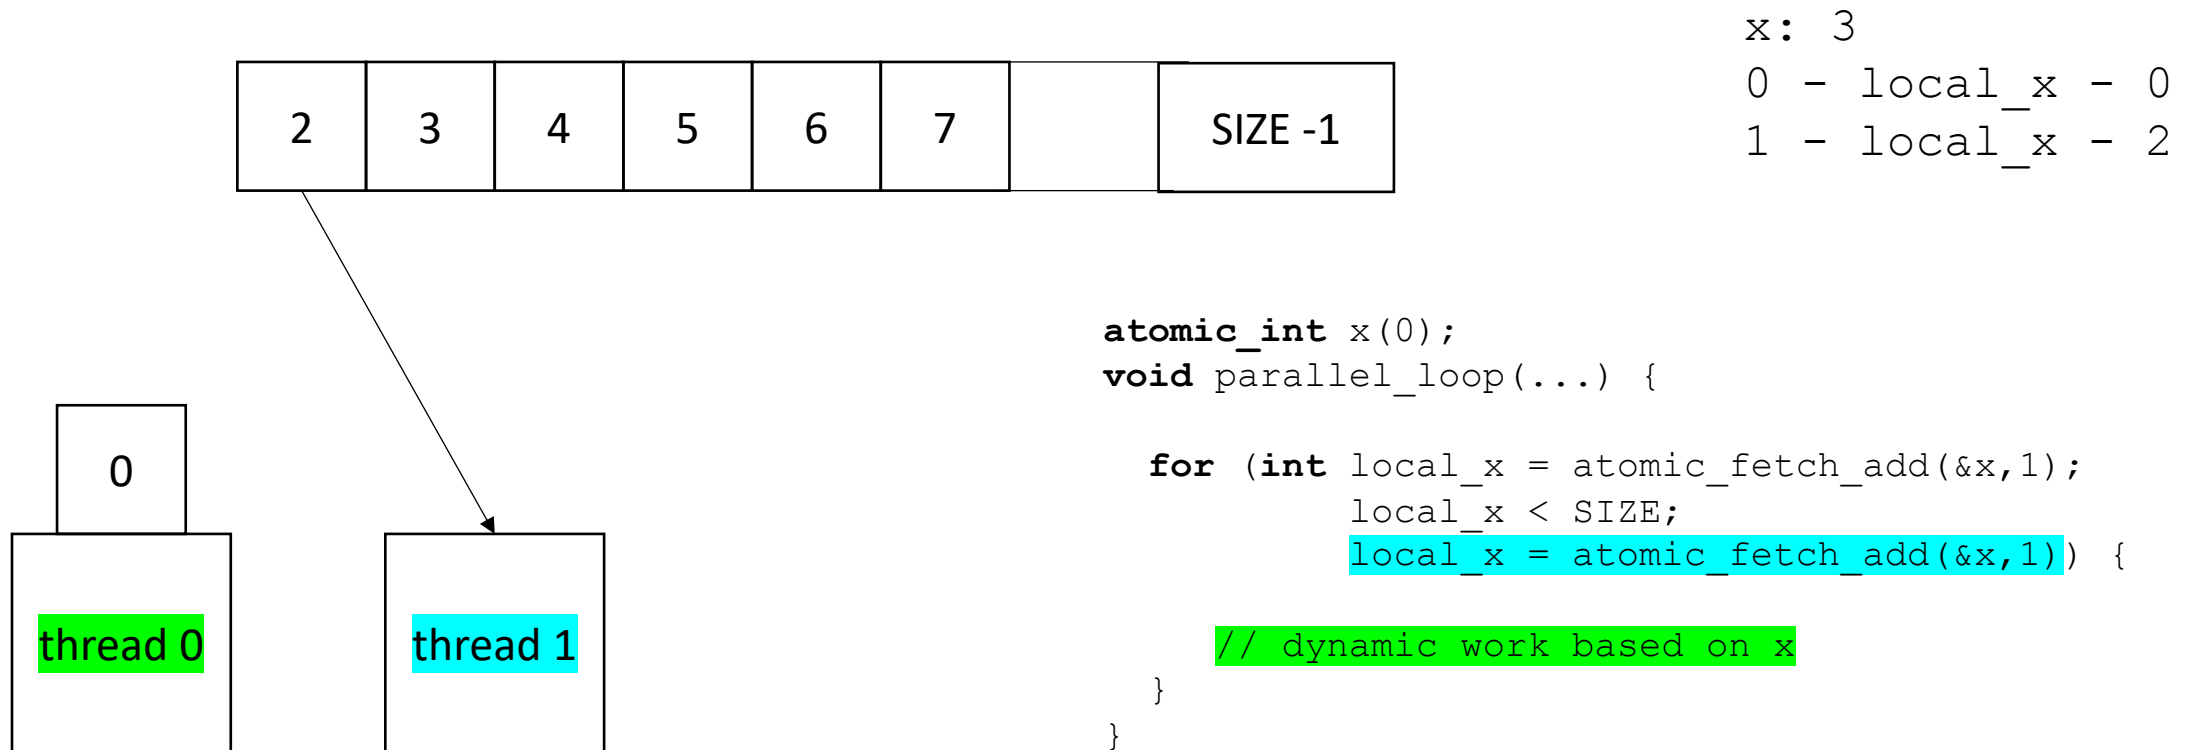
x: 2
0 - local_x - 0
1 - local_x - 1



```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```



Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




x: 3
0 - local_x - 0
1 - local_x - 2




```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



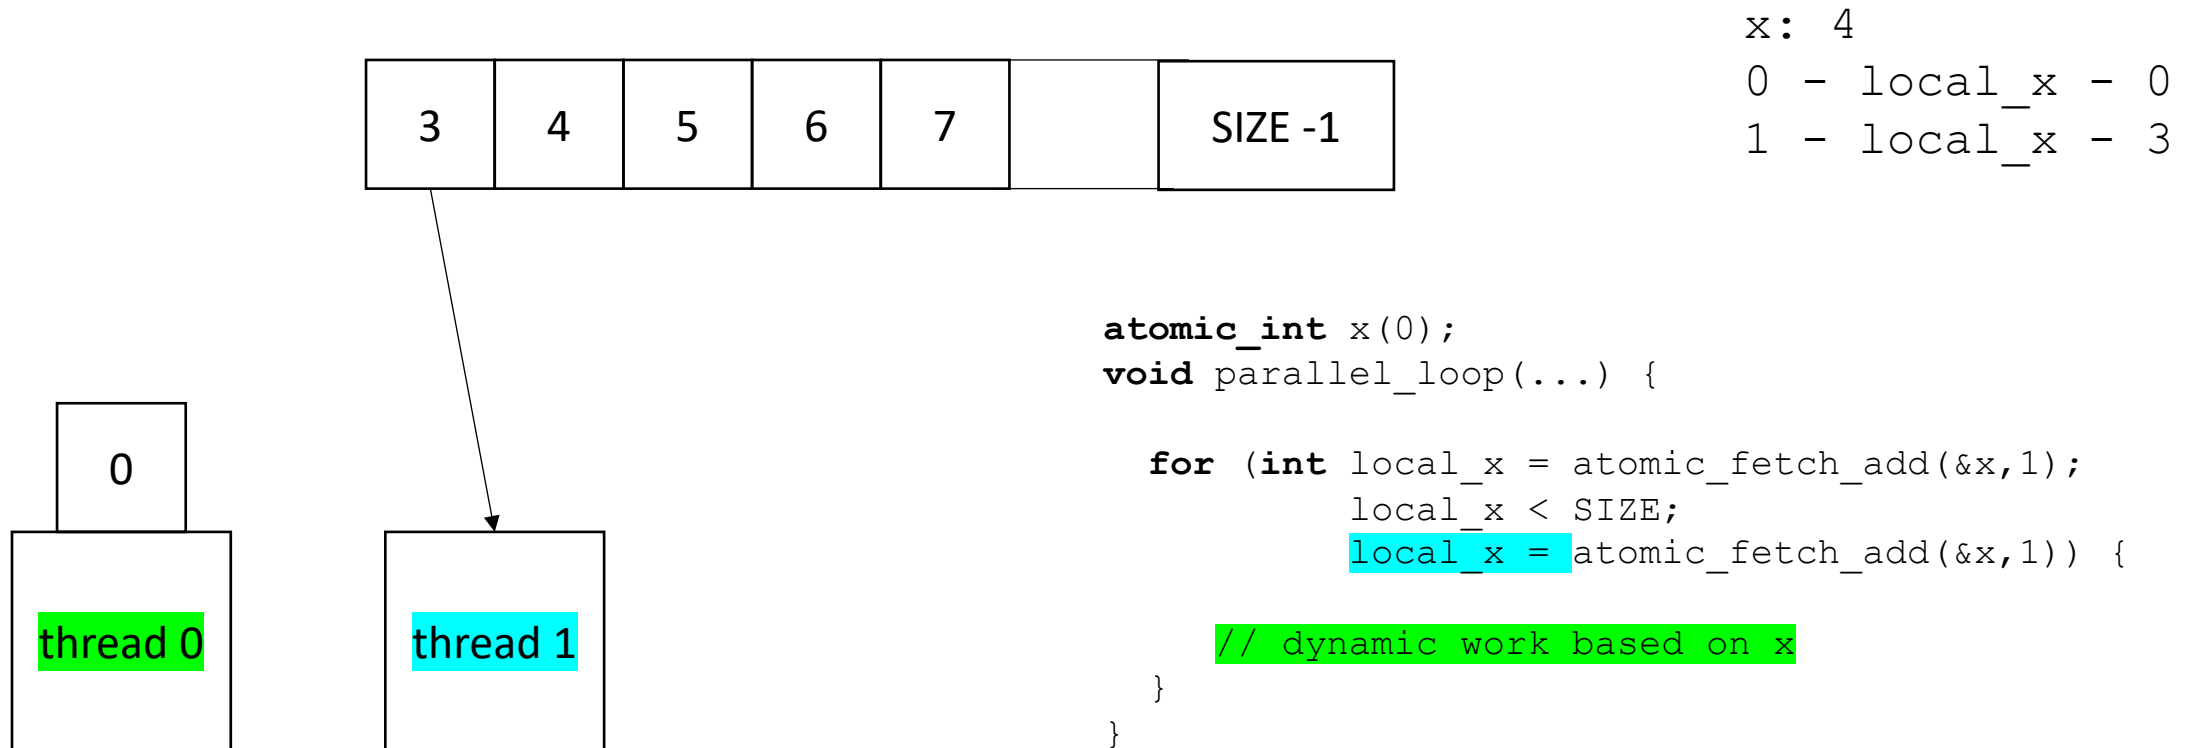
x: 3
0 - local_x - 0
1 - local_x - 2



```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```


Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




x: 4
0 - local_x - 0
1 - local_x - 3




```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```

Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically




x: 4
0 - local_x - 0
1 - local_x - 3



```
atomic_int x(0);  
void parallel_loop(...) {  
  
    for (int local_x = atomic_fetch_add(&x,1);  
         local_x < SIZE;  
         local_x = atomic_fetch_add(&x,1)) {  
  
        // dynamic work based on x  
    }  
}
```


Work stealing - global implicit worklist

- Global worklist: threads take tasks (iterations) dynamically



Work stealing - global implicit worklist


- Global worklist: threads take tasks (iterations) dynamically



x: 5

0 - local_x - 4

1 - local_x - 3



```
atomic_int x(0);
void parallel_loop(...) {
    for (int local_x = atomic_fetch_add(&x,1);
         local_x < SIZE;
         local_x = atomic_fetch_add(&x,1)) {
        // dynamic work based on x
    }
}
```


Schedule




- DOALL Loops
- **Parallel Schedules:**
 - Static
 - Global Worklists
 - **Local Worklists**

Work stealing - local worklists

- More difficult to implement
- low contention on local data-structures
- potentially better cache locality


Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread



Work stealing - local worklists

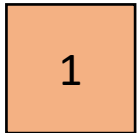
- local worklists: divide tasks into different worklists for each thread




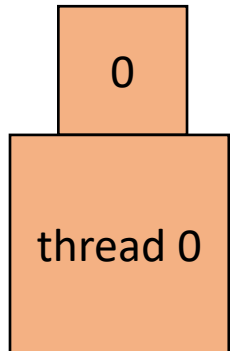
Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread

Queue 0



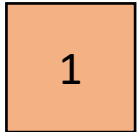
Queue 1






Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread

Queue 0




Queue 1



Work stealing - local worklists


- local worklists: divide tasks into different worklists for each thread






Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread

Queue 0






Queue 1




Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread

Queue 0




Queue 1



Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread




Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread

Queue 0

Queue 1



Work stealing - local worklists

- How to implement:

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
    ...  
}
```

Work stealing - local worklists

- How to implement:

```
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid) {  
  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

Make a new function, taking any variables used in loop body as args. Additionally take in a thread id

Work stealing - local worklists

- How to implement:

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid) {  
  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

Make a global array of concurrent queues

Work stealing - local worklists

- How to implement:

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

```
void parallel_loop(..., int tid) {  
  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

What type of queues?

Make a global array of concurrent queues

Work stealing - local worklists

- How to implement:

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    for (x = 0; x < SIZE; x++) {  
    // dynamic work based on x  
    }  
    ...  
}
```

Make a global array of concurrent queues


```
void parallel_loop(..., int tid) {  
  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

What type of queues?

We're going to use InputOutput Queues!


Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread



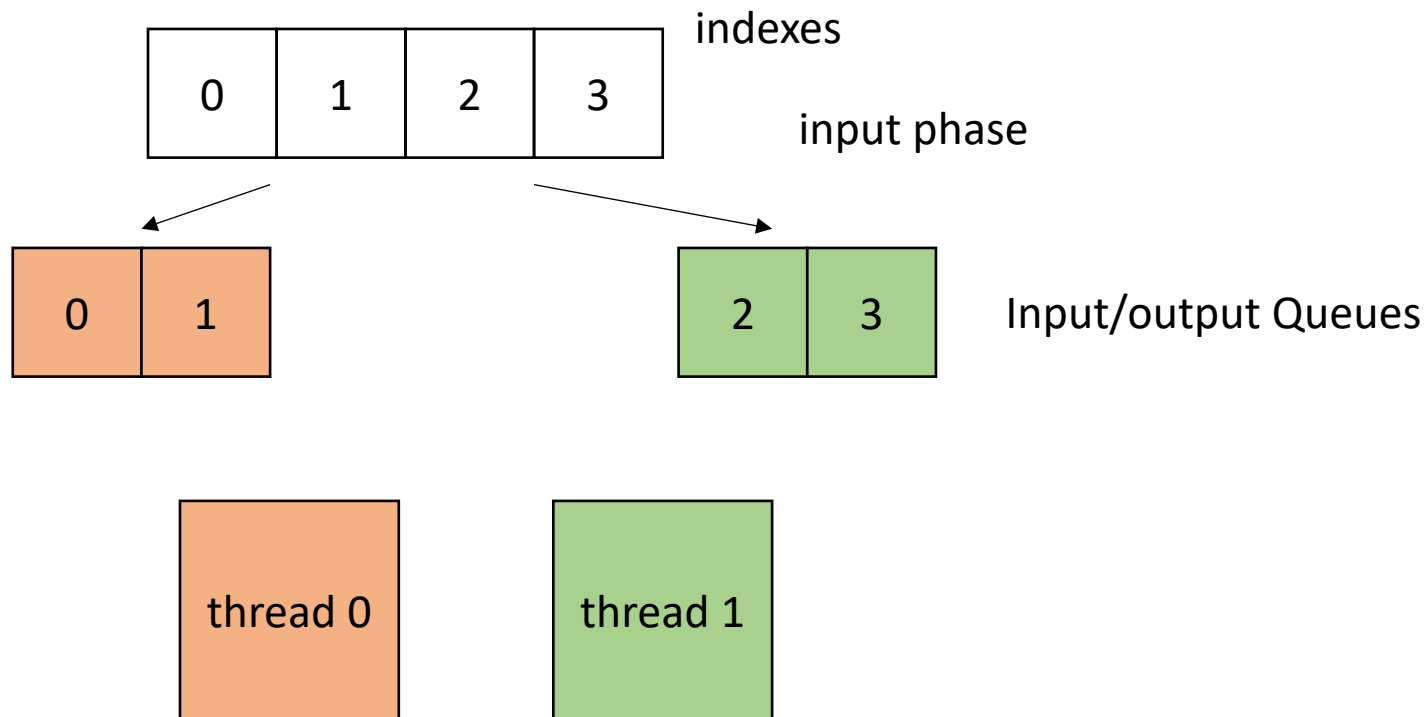
Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread



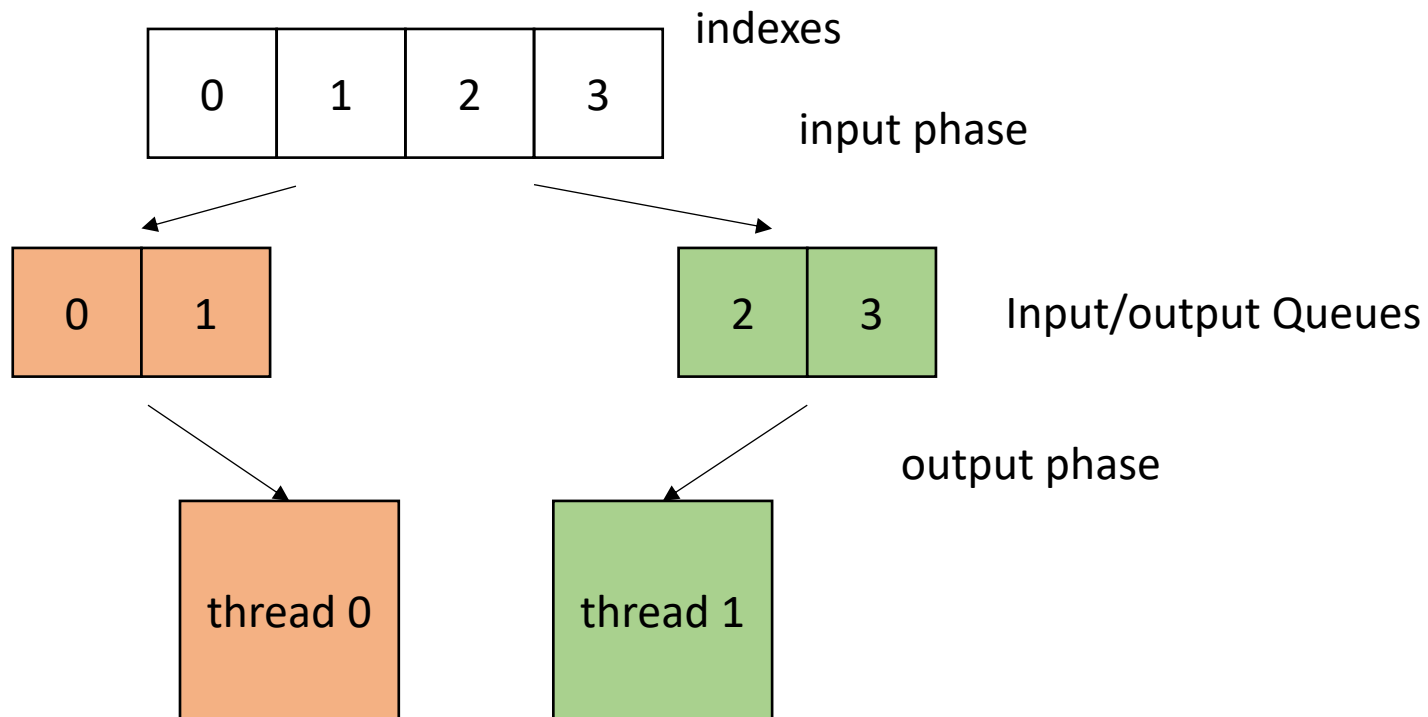
Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread



Work stealing - local worklists

- local worklists: divide tasks into different worklists for each thread



Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
  
    ...  
}
```

First we need to initialize the queues

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // Spawn threads to initialize  
    // join initializing threads  
    ...  
}
```

```
void parallel_enq(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    int start = chunk_size * tid;  
    int end = start + chunk_size;  
    for (int x = start; x < end; x++) {  
        cq[tid].enq(x);  
    }  
}
```

Just like the static schedule, except we are enqueueing

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // Spawn threads to initialize  
    // join initializing threads  
    ...  
}
```

Make sure to account for boundary conditions!

```
void parallel_enq(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    int start = chunk_size * tid;  
    int end = start + chunk_size;  
    for (int x = start; x < end; x++) {  
        cq[tid].enq(x);  
    }  
}
```

Just like the static schedule, except we are enqueueing

Work stealing - local worklists

- How to implement in a compiler:

```
NUM_THREADS = 2;  
SIZE = 4;  
CHUNK = 2;
```

x	0	1	2	3
tid	0	0	1	1

Make sure to account for boundary conditions!

```
void parallel_enq(..., int tid, int num_threads)  
{  
  
    int chunk_size = SIZE / NUM_THREADS;  
    int start = chunk_size * tid;  
    int end = start + chunk_size;  
    for (int x = start; x < end; x++) {  
        cq[tid].enq(x);  
    }  
}
```

Just like the static schedule, except we are enqueueing

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads) {  
    for (x = 0; x < SIZE; x++) {  
        // dynamic work based on x  
    }  
}
```

How do we modify the parallel loop?

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
}
```

loop until the queue is empty

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
}
```

loop until the queue is empty
Are we finished?

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
}
```

Track how many threads are finished

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
    while (finished_threads.load() != num_threads) {  
  
    }  
}
```

While there are threads that are still working

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    ...  
}
```

```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
    while (finished_threads.load() != num_threads) {  
        int target = // pick a thread to steal from  
        int task = cq[target].deq();  
    }  
}
```

pick a random target and steal a task

Work stealing - local worklists

```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    // join loop threads  
    ...  
}
```

```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
    while (finished_threads.load() != num_threads) {  
        int target = // pick a thread to steal from  
        int task = cq[target].deq();  
        if (task != -1) {  
            // perform task  
        }  
    }  
}
```

Work stealing - local worklists


```
concurrent_queues cq[NUM_THREADS];  
void foo() {  
    ...  
    // initialize queues  
    // join threads  
  
    // launch loop function  
    // join loop threads  
    ...  
}
```

join the threads




```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
    while (finished_threads.load() != num_threads) {  
        int target = // pick a thread to steal from  
        int task = cq[target].deq();  
        if (task != -1) {  
            // perform task  
        }  
    }  
}
```


Work stealing - local worklists

IOQueue 0




IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists





```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {



    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

IOQueue 0




IOQueue 1




```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads,1);  
    while (finished_threads.load() != num_threads) {  
        int target = // pick a thread to steal from  
        int task = cq[target].deq();  
        if (task != -1) {  
            // perform task  
        }  
    }  
}
```

Work stealing - local worklists

IOQueue 0



IOQueue 1





```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {


    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

IOQueue 0



IOQueue 1





```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {


    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

IOQueue 0



IOQueue 1




```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```



Work stealing - local worklists

finished_threads: 1

IOQueue 0



IOQueue 1




```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```



Work stealing - local worklists

finished_threads: 1

IOQueue 0



IOQueue 1




```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```




Work stealing - local worklists

finished_threads: 1

IOQueue 0



IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {


    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

finished_threads: 1

IOQueue 0

IOQueue 1




```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```


Work stealing - local worklists

finished_threads: 1

IOQueue 0



IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {


    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

finished_threads: 1

IOQueue 0

IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {


    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

finished_threads: 2

IOQueue 0

IOQueue 1



```
atomic_int finished_threads(0);  
void parallel_loop(..., int tid, int num_threads) {  
  
    int task = 0;  
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())  
    {  
        // dynamic work based on task  
    }  
    atomic_fetch_add(&finished_threads, 1);  
    while (finished_threads.load() != num_threads) {  
        int target = // pick a thread to steal from  
        int task = cq[target].deq();  
        if (task != -1) {  
            // perform task  
        }  
    }  
}
```


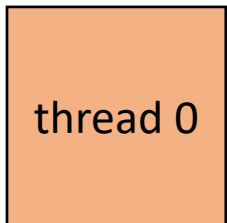
Work stealing - local worklists

finished_threads: 2

IOQueue 0

IOQueue 1

finished!



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {



    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

finished_threads: 2

IOQueue 0

IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {



    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```

Work stealing - local worklists

finished_threads: 2

IOQueue 0

IOQueue 1



```
atomic_int finished_threads(0);
void parallel_loop(..., int tid, int num_threads) {

    int task = 0;
    for (x = cq[tid].deq(); x != -1; x = cq[tid].deq())
    {
        // dynamic work based on task
    }
    atomic_fetch_add(&finished_threads,1);
    while (finished_threads.load() != num_threads) {
        int target = // pick a thread to steal from
        int task = cq[target].deq();
        if (task != -1) {
            // perform task
        }
    }
}
```


Next topic

- General concurrent set

Schedule

- **Concurrent set**
 - Coarse-grained lock
 - fine-grained lock
 - optimistic locking

Thanks to Roberto Palmieri (Lehigh University) and material from the text book for some of the slide content/ideas.

Set Interface

- Unordered collection of items
- No duplicates
- We will implement this as a sorted linked list


Set Interface

- Unordered collection of items
- No duplicates
- Methods
 - **add (x)** put **x** in set
 - **remove (x)** take **x** out of set
 - **contains (x)** tests if **x** in set

List Node

```
class Node {  
    public:  
        Value v;  
        int key;  
        Node *next;  
}
```

The List-Based Set



Sorted with Sentinel nodes
(min & max possible keys)

Sequential List Based Set

add(b)




remove(b)




Sequential List Based Set

add(b)



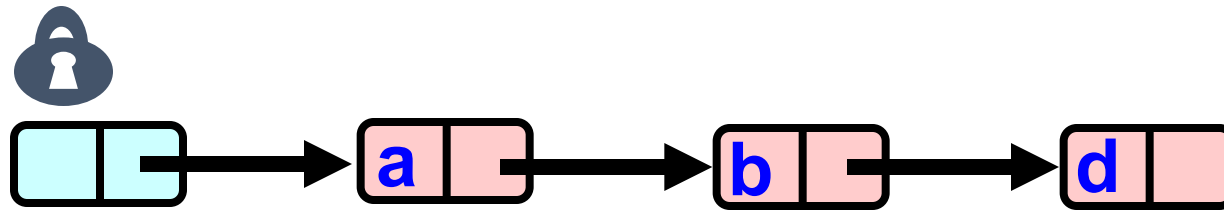
remove(b)



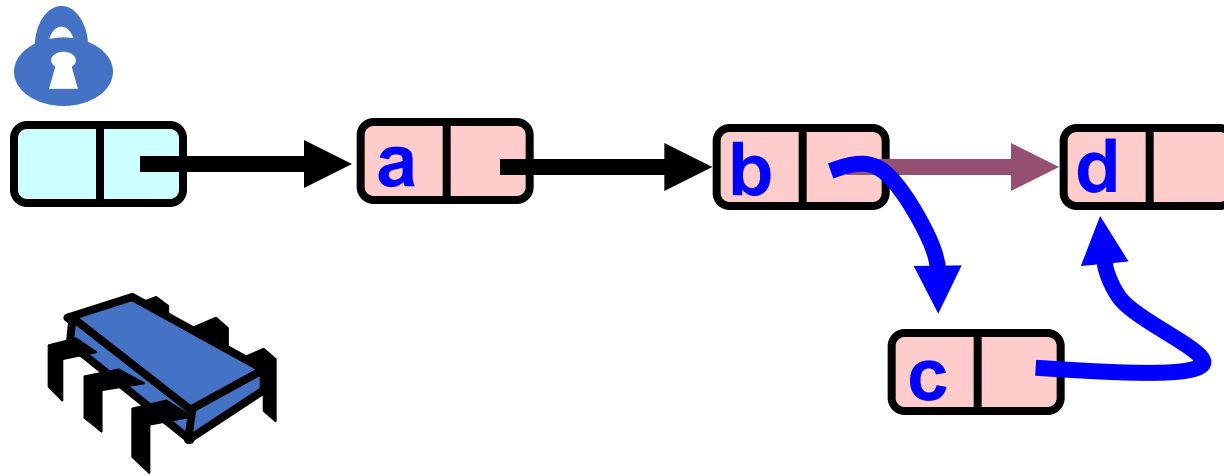
Schedule

- Concurrent set
 - **Coarse-grained lock**
 - fine-grained lock
 - optimistic locking

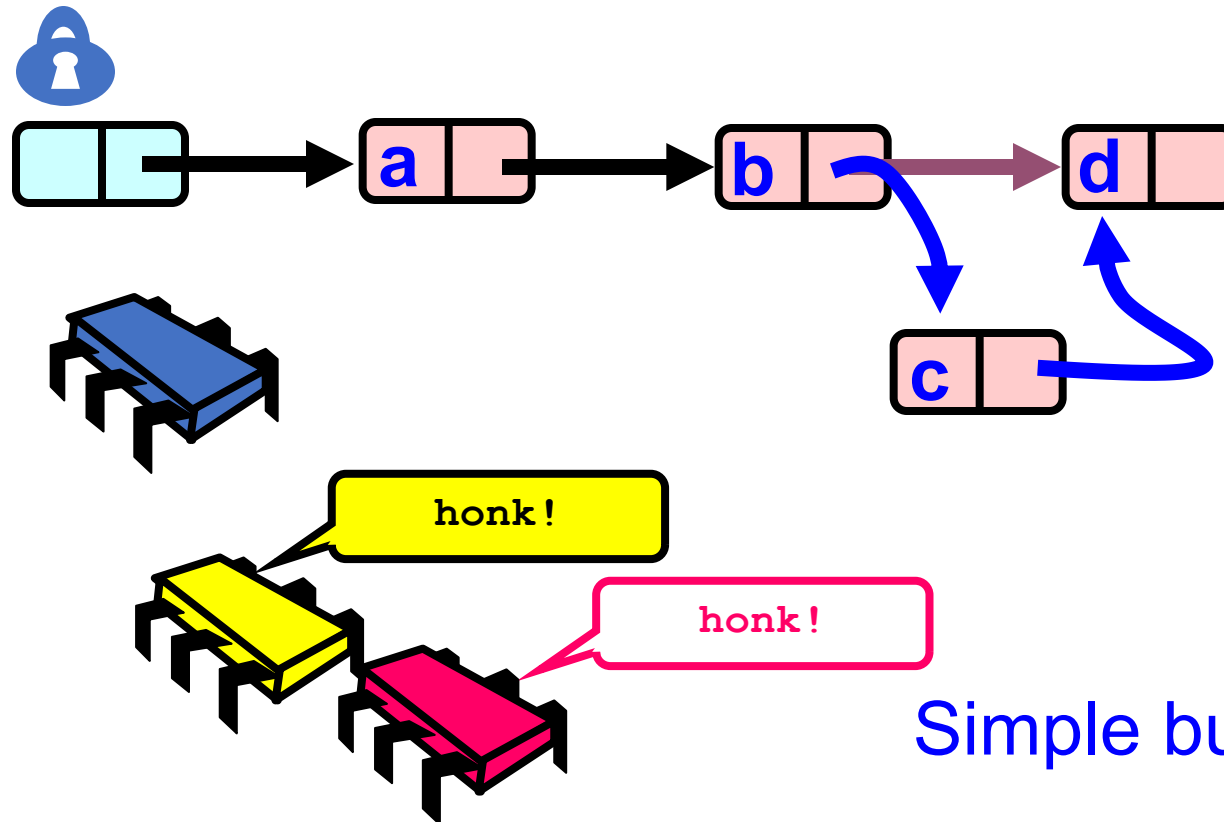
Coarse-Grained Locking



Coarse-Grained Locking



Coarse-Grained Locking



Simple but inefficient!



Schedule

- Concurrent set
 - Coarse-grained lock
 - **fine-grained lock**
 - optimistic locking


Fine-grained Locking

- Requires **careful** thought
- Split object into pieces
 - Each piece has own lock
 - Methods that work on disjoint pieces need not exclude each other


Hand-over-Hand locking




Hand-over-Hand locking




Hand-over-Hand locking




Hand-over-Hand locking




Hand-over-Hand locking




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes



Concurrent Removes




Uh, Oh





Uh, Oh

Bad news, c not removed



Problem



- To delete node c
 - Swing node b's next field to d
- Problem is,
 - ***Data conflict:***
 - Someone deleting b concurrently could direct a pointer to C




Insight

- If a node is locked
 - No one can delete node's *successor*
- If a thread locks
 - Node to be deleted
 - And its predecessor
 - Then it works


Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Hand-Over-Hand Again




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node




Removing a Node



Removing a Node



Removing a Node



Removing a Node



Adding Nodes




- To add node e
 - Must lock predecessor
 - Must lock successor
- Neither can be deleted

Drawbacks


- Better than coarse-grained lock
 - Threads can traverse in parallel
- Still not ideal
 - Long chain of acquire/release
 - Inefficient

```
void remove(Value v) {
    Node* pred = NULL, *curr = NULL;
    head.lock();
    pred = head;
    curr = pred.next();
    curr.lock();
    while (curr.value != v) {
        pred.unlock();
        pred = curr;
        curr = curr.next();
        curr.lock();
    }
    pred.next = curr.next;
    curr.unlock();
    pred.unlock();
}
```


```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```





```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```




```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```




```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```




```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```




```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```





```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```



```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```



```
void remove(Value v) {  
    Node* pred = NULL, *curr = NULL;  
    head.lock();  
    pred = head;  
    curr = pred.next();  
    curr.lock();  
    while (curr.value != v) {  
        pred.unlock();  
        pred = curr;  
        curr = curr.next();  
        curr.lock();  
    }  
    pred.next = curr.next;  
    curr.unlock();  
    pred.unlock();  
}
```



Schedule

- Concurrent set
 - Coarse-grained lock
 - fine-grained lock
 - **optimistic locking**

How can we improve

- Acquires and releases lock for every node traversed
 - If we have a long list to search, it can be bad!
 - reduces concurrency (traffic jams)

Optimistic Synchronization

Assume there will be no conflicts. Check before committing. If there was a conflict, try again.

Optimistic Synchronization

- Find nodes without locking


Optimistic Synchronization

- Find nodes without locking
- Lock nodes


Optimistic Synchronization

- Find nodes without locking
- Lock nodes
- Check that everything is OK


Optimistic: Traverse without Locking



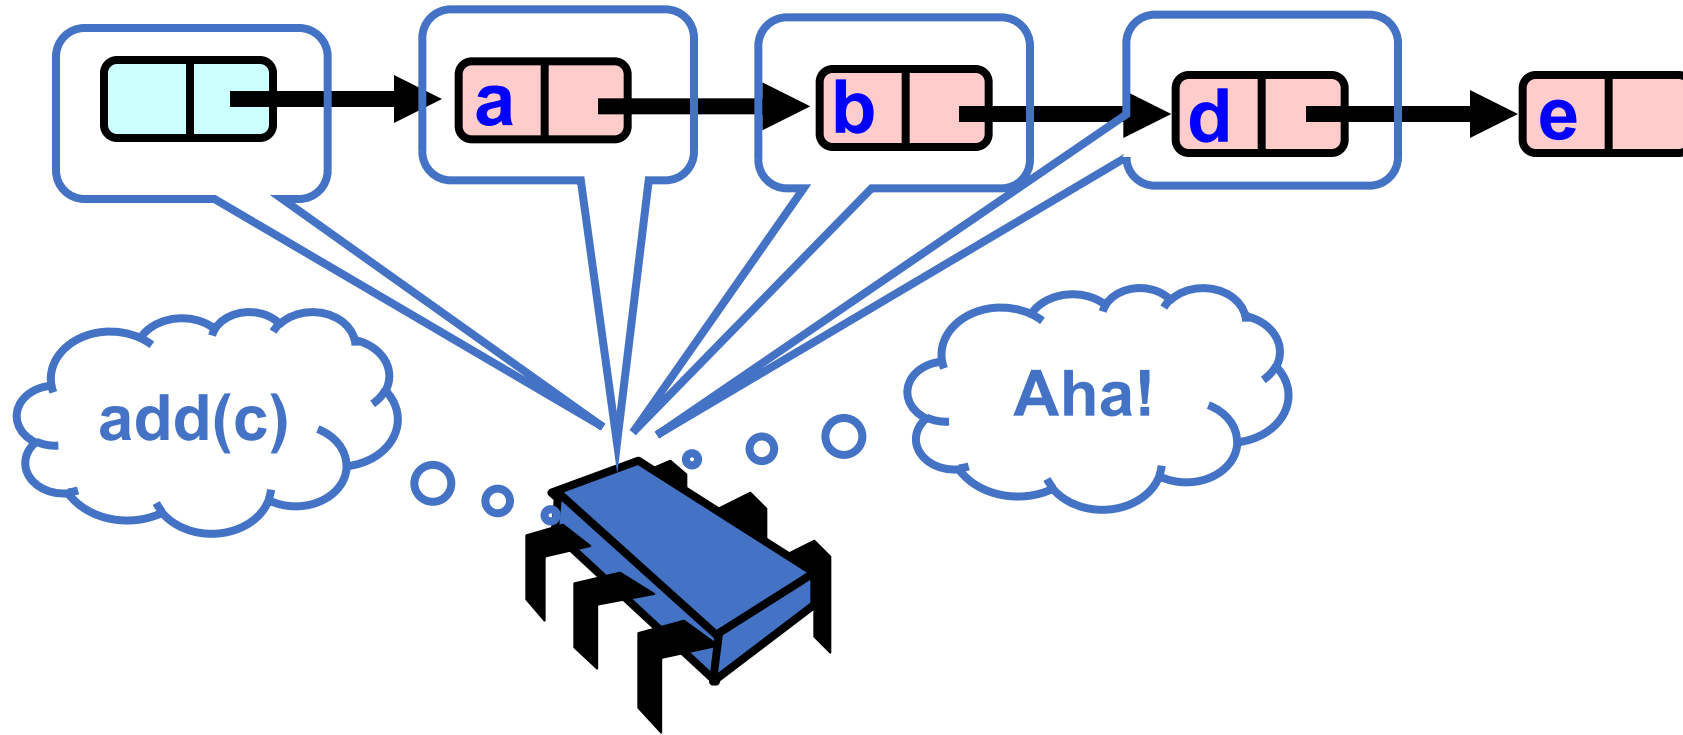
Optimistic: Lock and Load



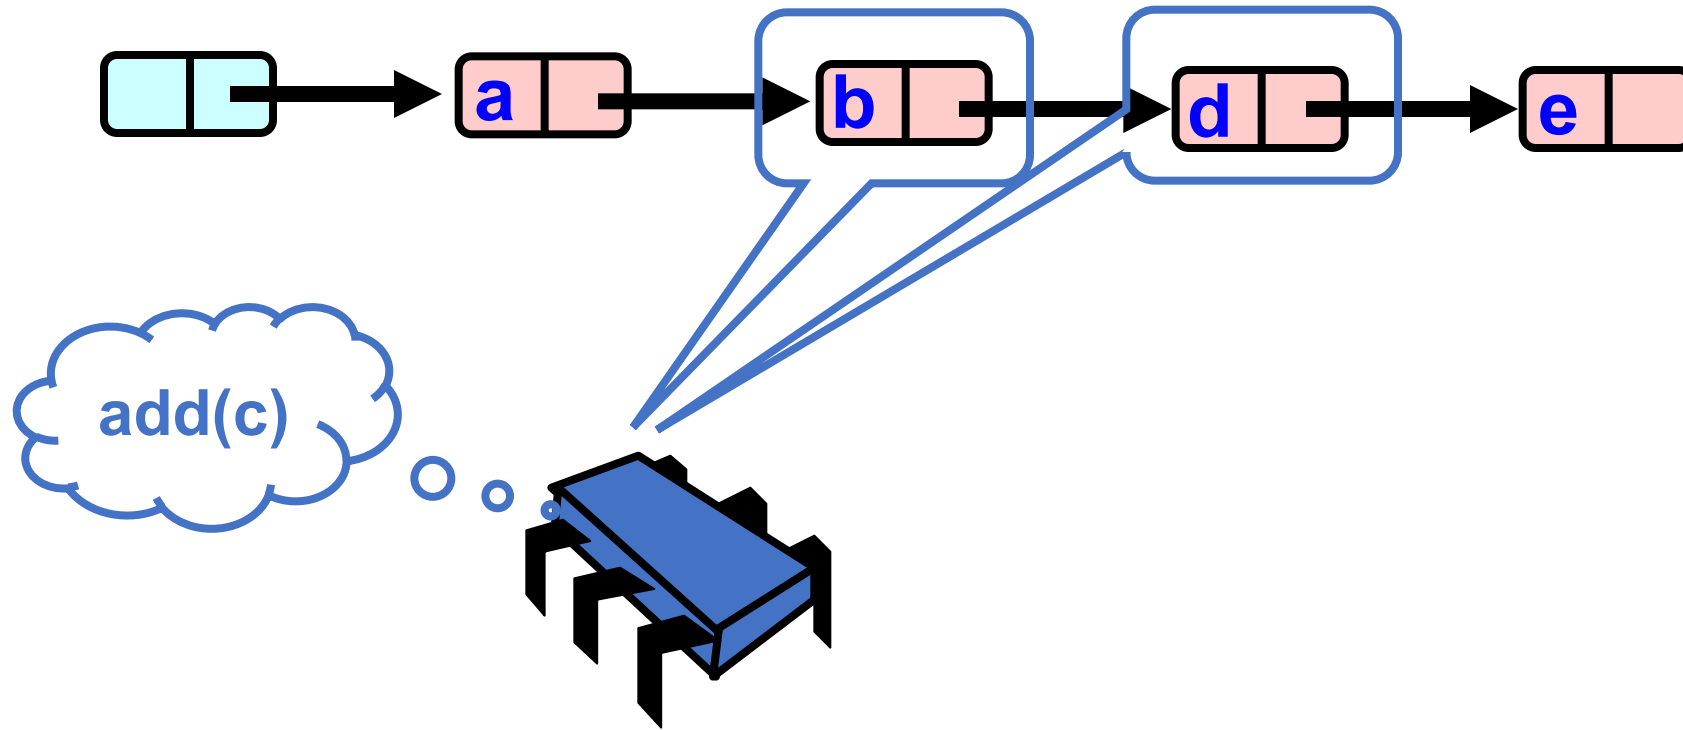
Optimistic: Lock and Load




What could go wrong?



What could go wrong?



What could go wrong?



Data conflict!

- Red thread has the lock on a node (so it can modify the node)
- Blue thread is traversing without locks
- What do we do?

Data conflict!

- Red thread has the lock on a node (so it can modify the node)
- Blue thread is traversing without locks
- What do we do? We decided that locking when traversing is too expensive.

Lock-free reasoning

- We can use atomic variables

Lock-free reasoning

- Default atomic accesses are documented to be sequentially consistent.

```
class Node {  
    public:  
        Value v;  
        int key;  
        Node *next;  
}
```


Lock-free reasoning

- Default atomic accesses are documented to be sequentially consistent.

```
class Node {  
    public:  
        Value v;  
        int key;  
        atomic<Node*> next;  
}
```

Create an atomic pointer type using C++ templates

Lock-free reasoning

- Default atomic accesses are documented to be sequentially consistent.

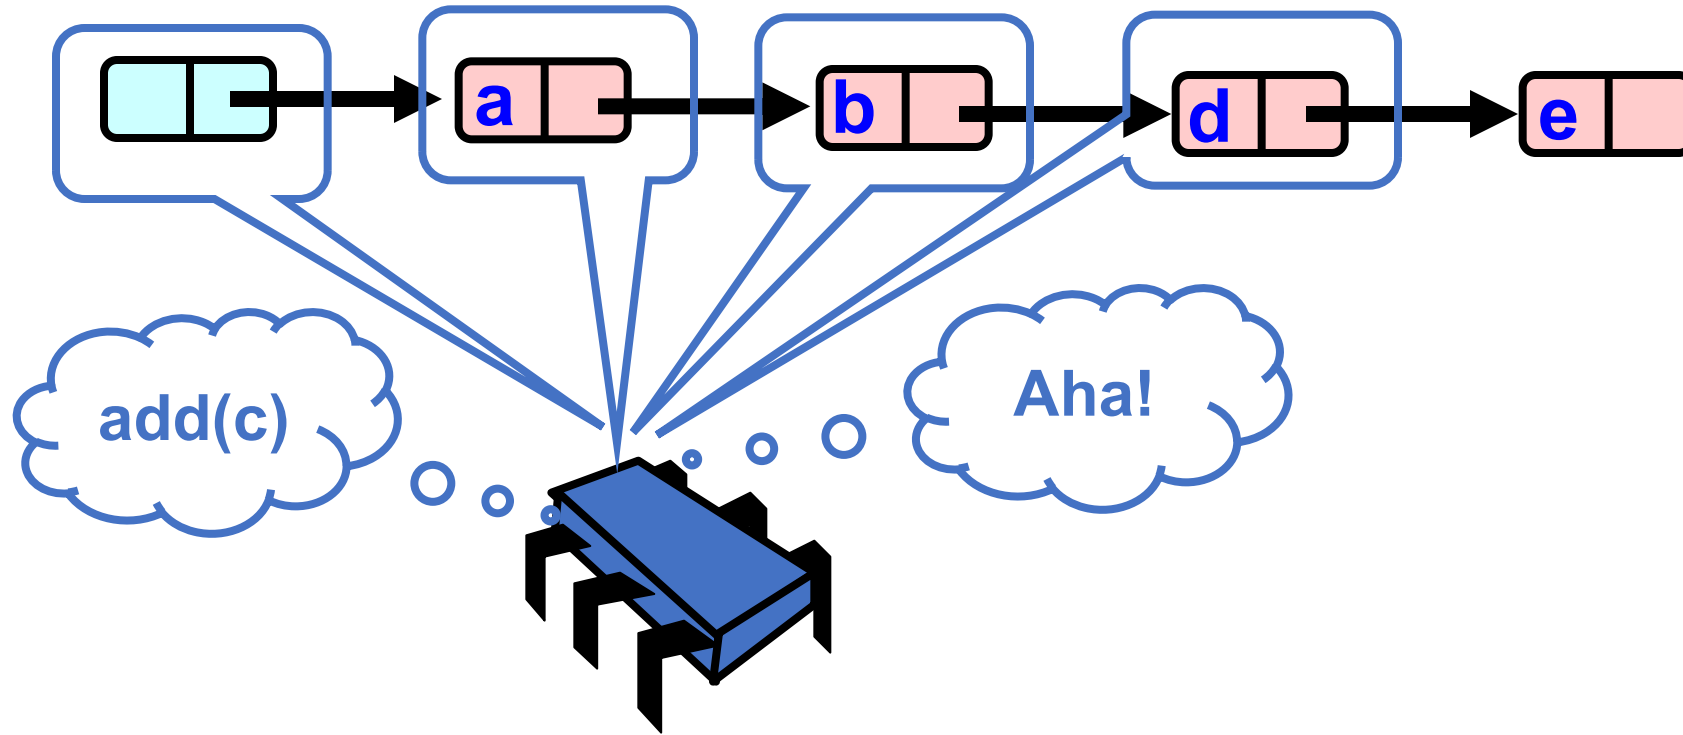
```
void traverse(node *n) {  
    while (n->next != NULL) {  
        n = n->next;  
    }  
}
```

Lock-free reasoning

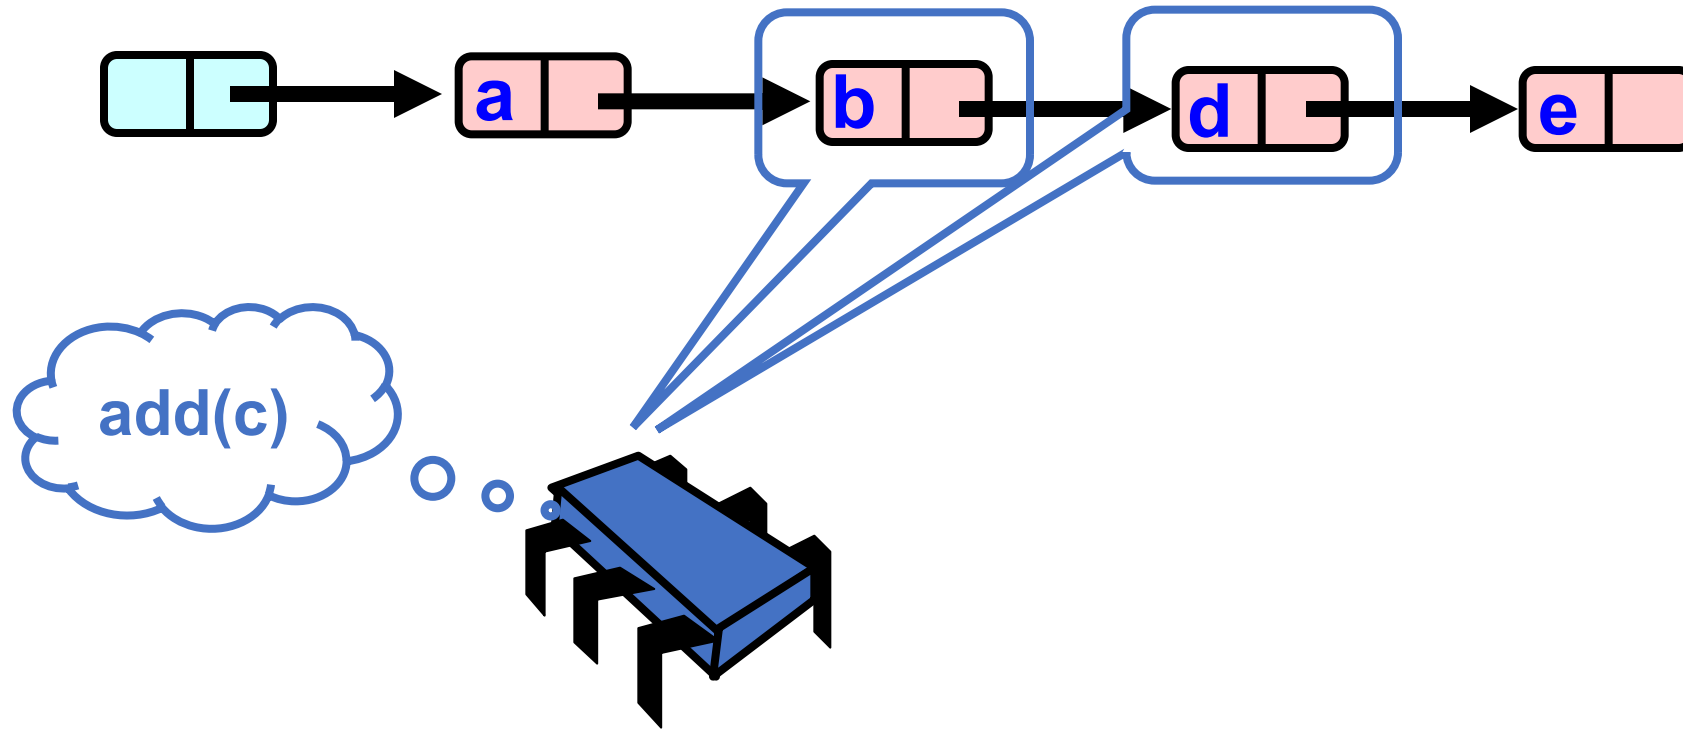
- Default atomic accesses are documented to be sequentially consistent.

```
void traverse(node *n) {  
    while (n->next.load() != NULL) {  
        n = n->next.load();  
    }  
}
```


What could go wrong?



What could go wrong?

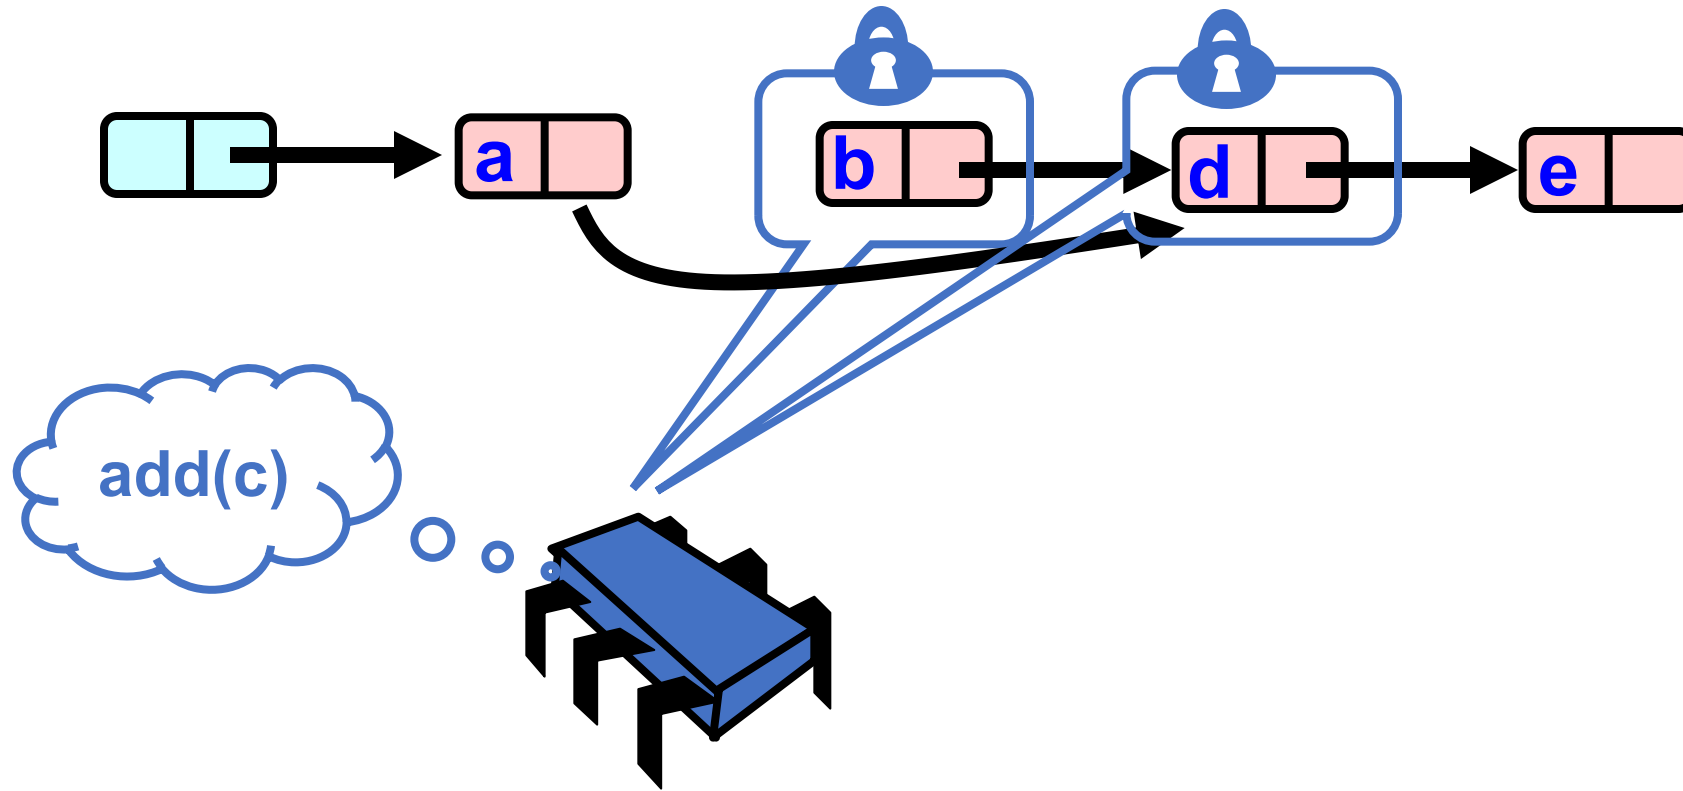


What could go wrong?

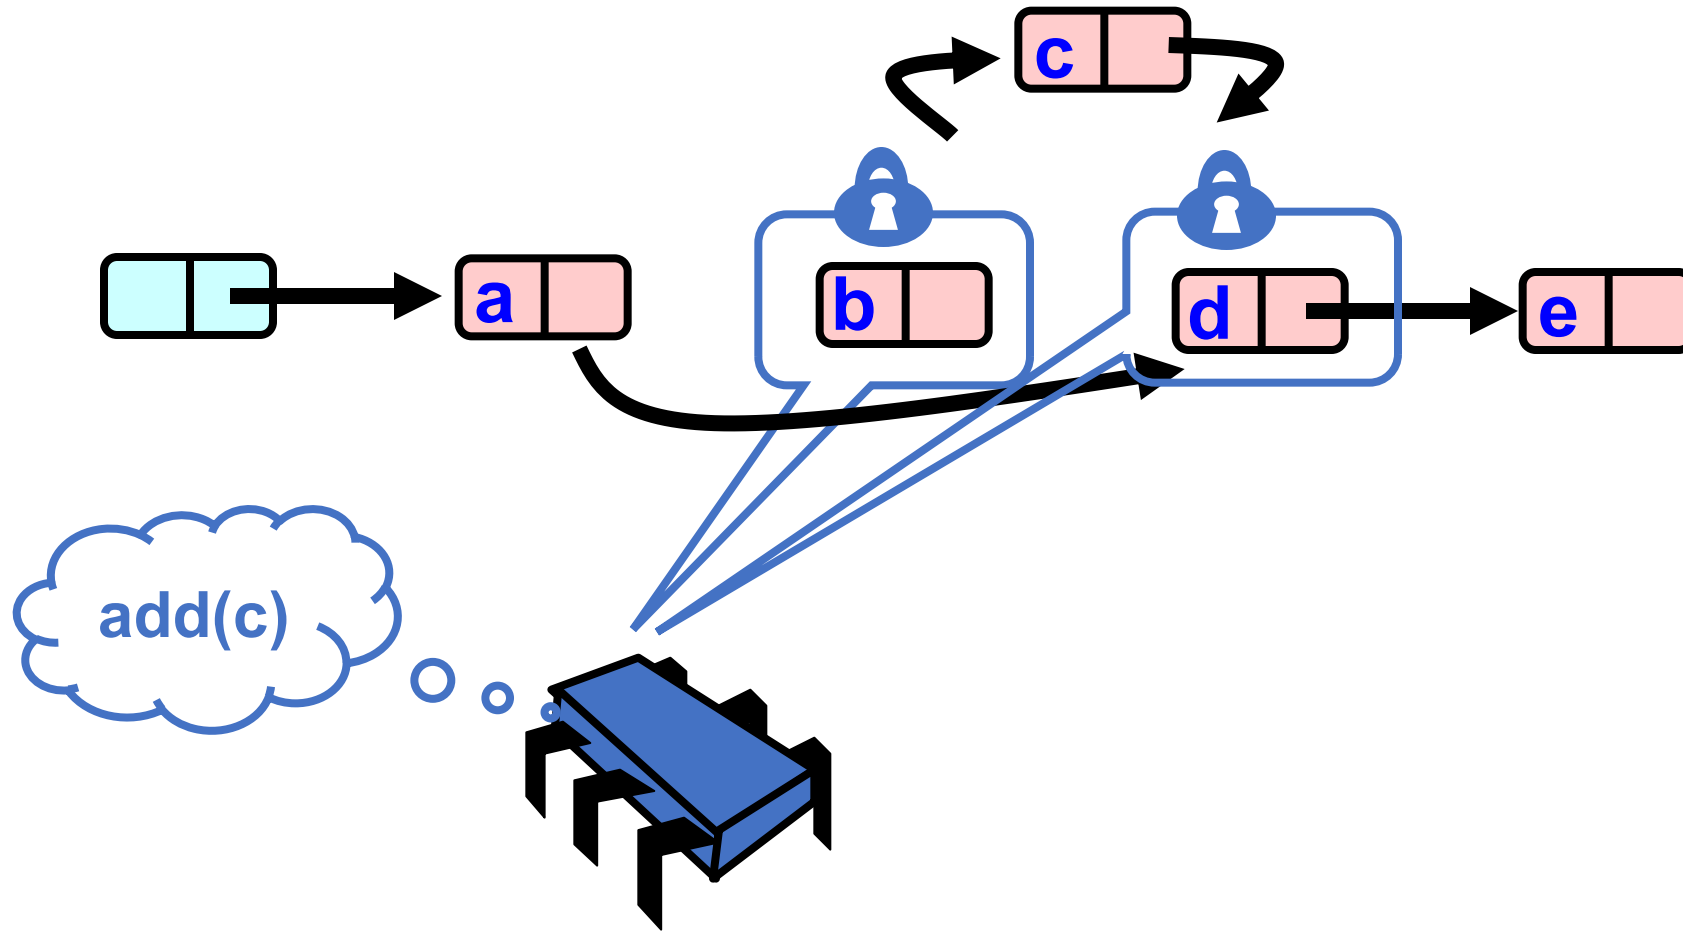


No more data conflict, but we do need to reason about interleavings and threads concurrent threads contending for values.

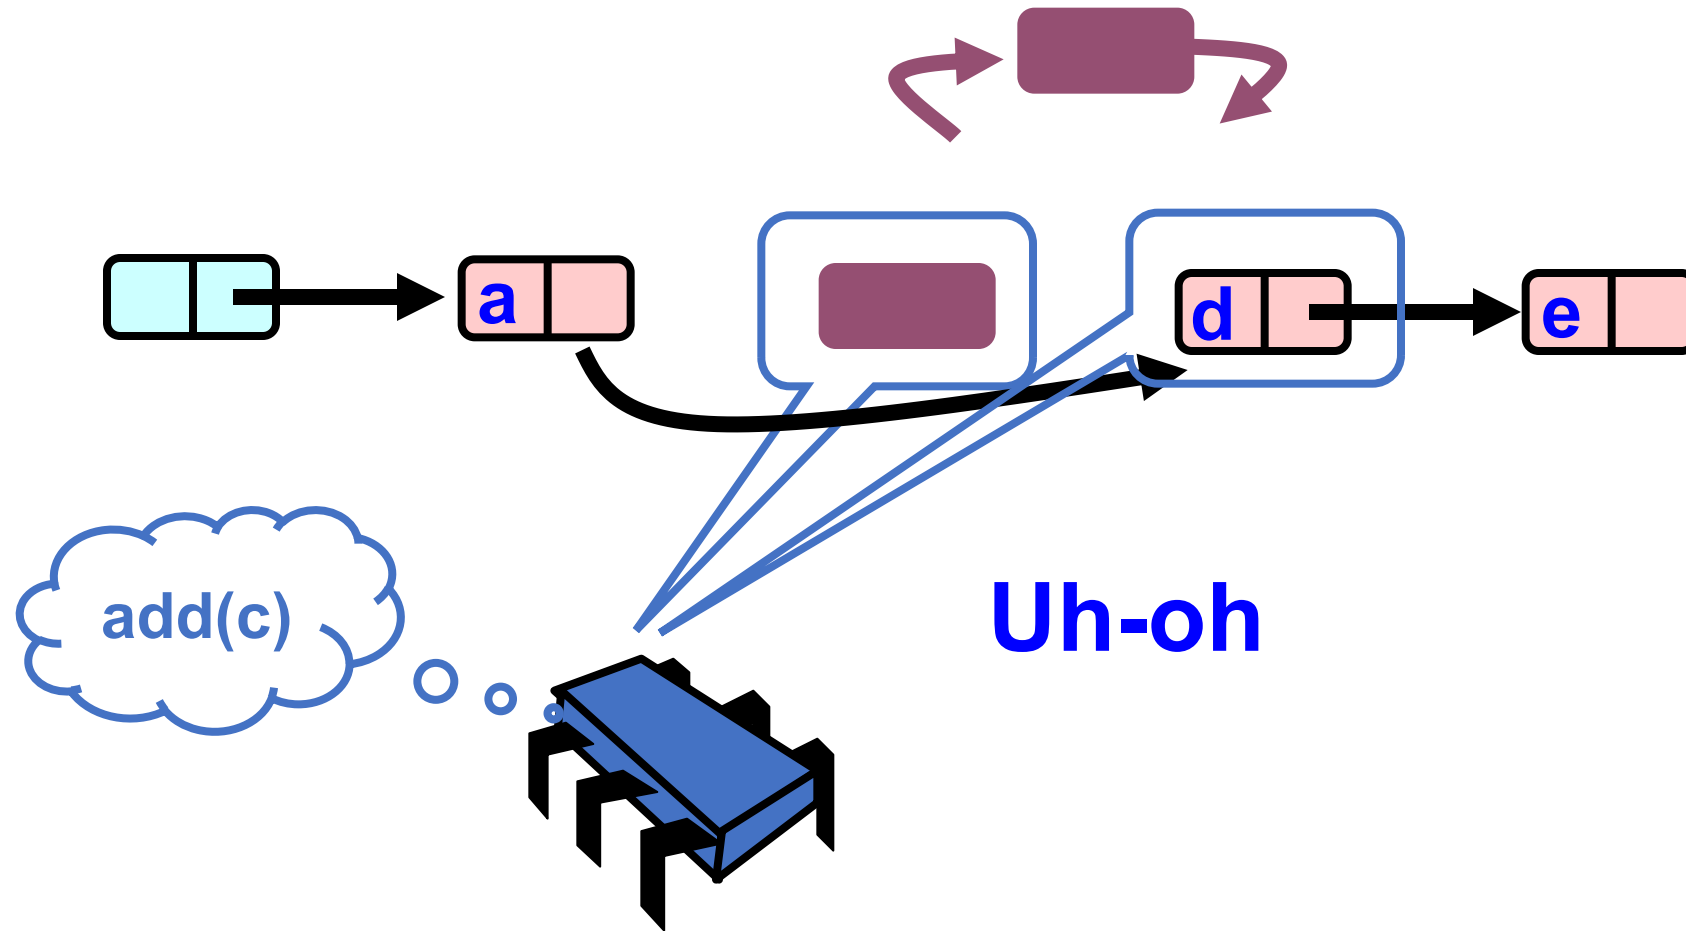
What could go wrong?



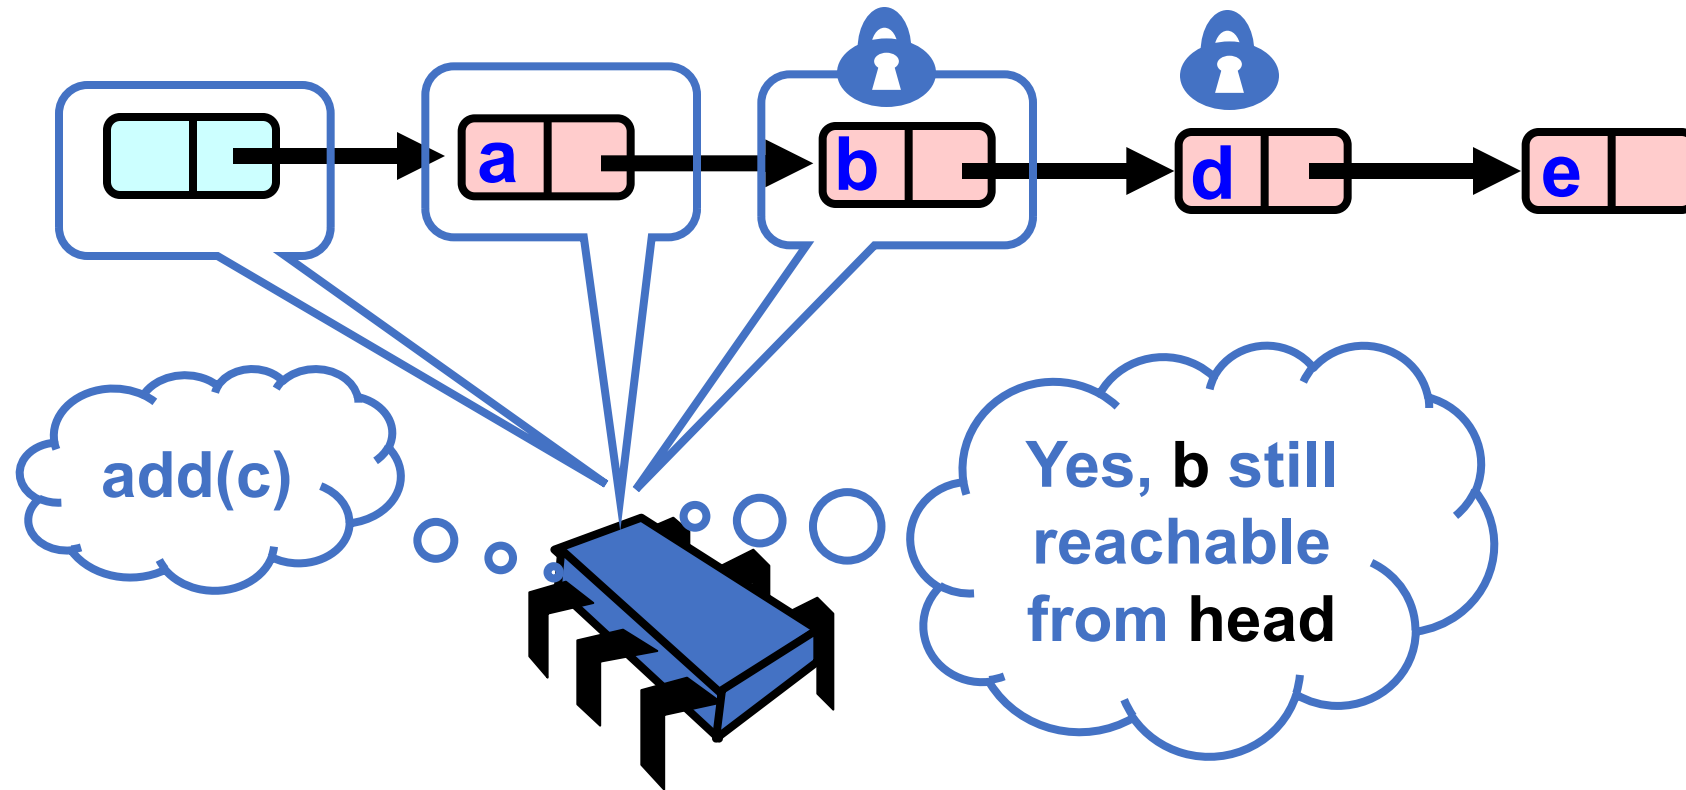
What could go wrong?



What could go wrong?



Validate – Part 1



What happens if failure?

- Ideas?

What happens if failure?


- Could try to recover? Back up a node?
 - Very tricky!
 - Just start over!

What happens if failure?


- Could try to recover? Back up a node?
 - Very tricky!
 - Just start over!
- Private method:
 - `try_remove`
 - remove loops on `try_remove` until it succeeds

What about deletion?


Can threads that remove a node delete it?




Can threads that remove a node delete it?




Can threads that remove a node delete it?



Can threads that remove a node delete it?



Our own garbage collector



Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:


Our own garbage collector



Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:





Our own garbage collector




Java's garbage collection will remove b

We are using a better™ language though...

maintain a list to delete:





Our own garbage collector



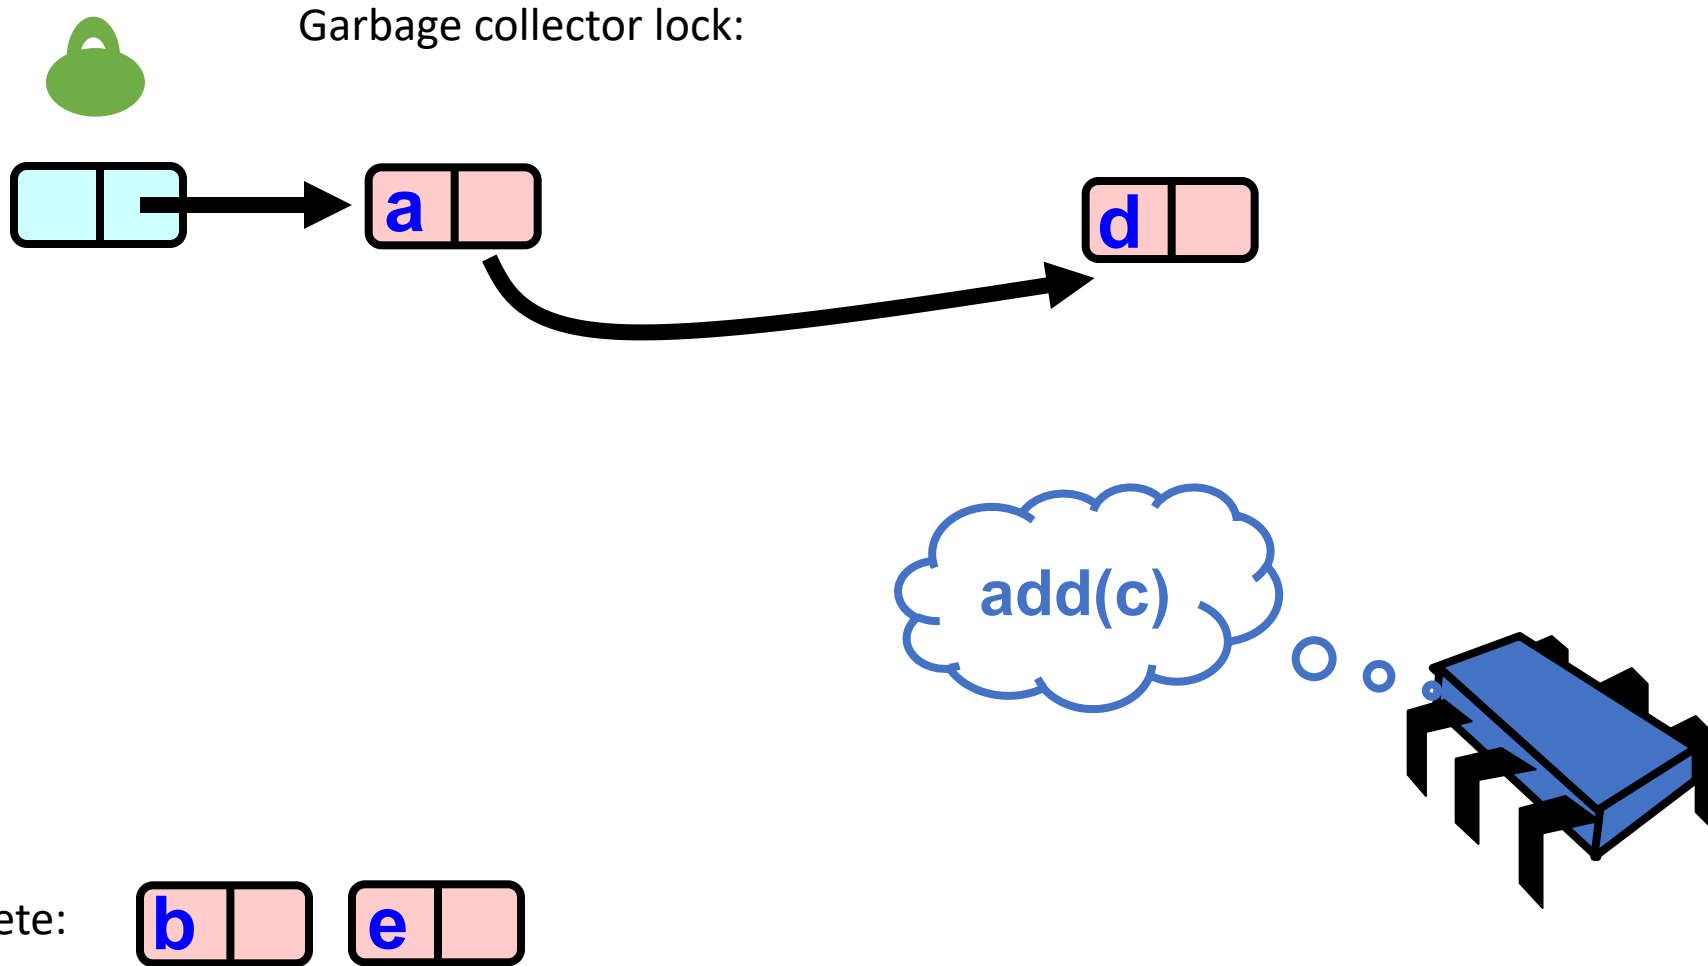
Java's garbage collection will remove b

We are using a better™ language though...


maintain a list to delete:




Our own garbage collector



Our own garbage collector



Similar to a reader/writer lock:
Allows an arbitrary number of threads that operate on the list
Only 1 garbage collector thread
Erases the list of nodes



maintain a list to delete:  

The text is followed by two pink boxes, each divided into two sections. The first box contains the letter 'b' in the left section, and the second box contains the letter 'e' in the left section. Both boxes have empty right sections.


Garbage collector lock

- Many strategies!
 - A big research area ~10 years ago
- **Strat 1:** Threads always try once to take the garbage collector lock:
 - if failed, no worries, the next operation will get a chance
 - if succeeded, then there was no contention
 - can starve garbage collection
- **Strat 2:** Wait until size grows to a threshold:
 - Wait on the lock (hope for a fair implementation!)
 - Can cause performance spikes

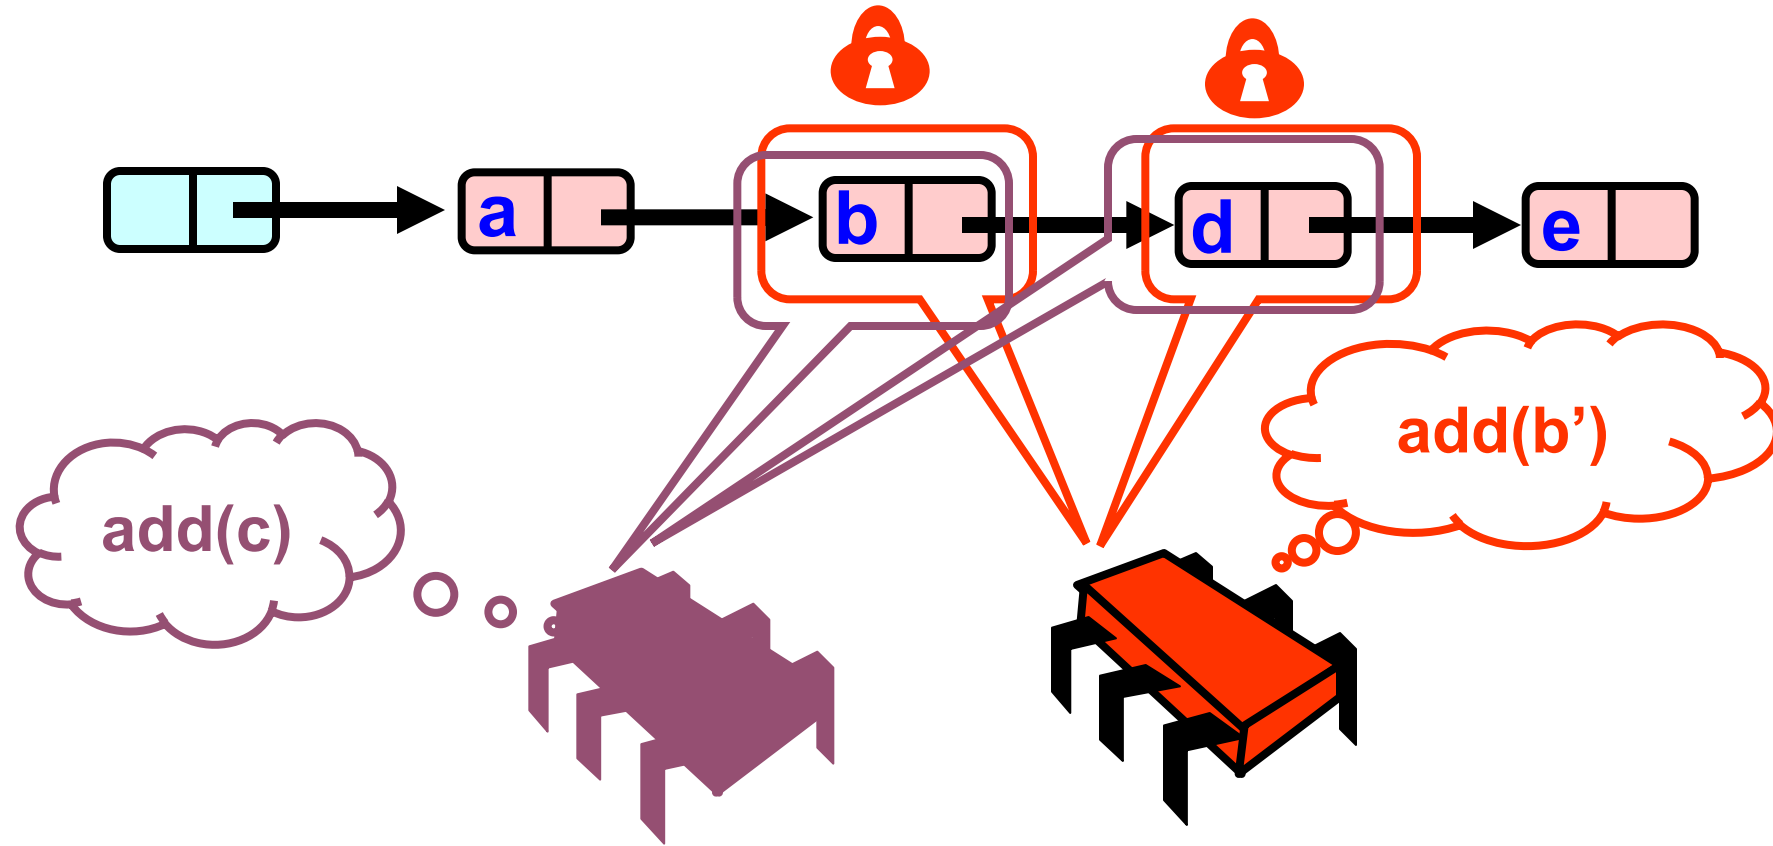
Back to the linked list

What if 2 threads try to add a node in the same position?


What Else Could Go Wrong?




What Else Could Go Wrong?



What Else Could Go Wrong?




What Else Could Go Wrong?

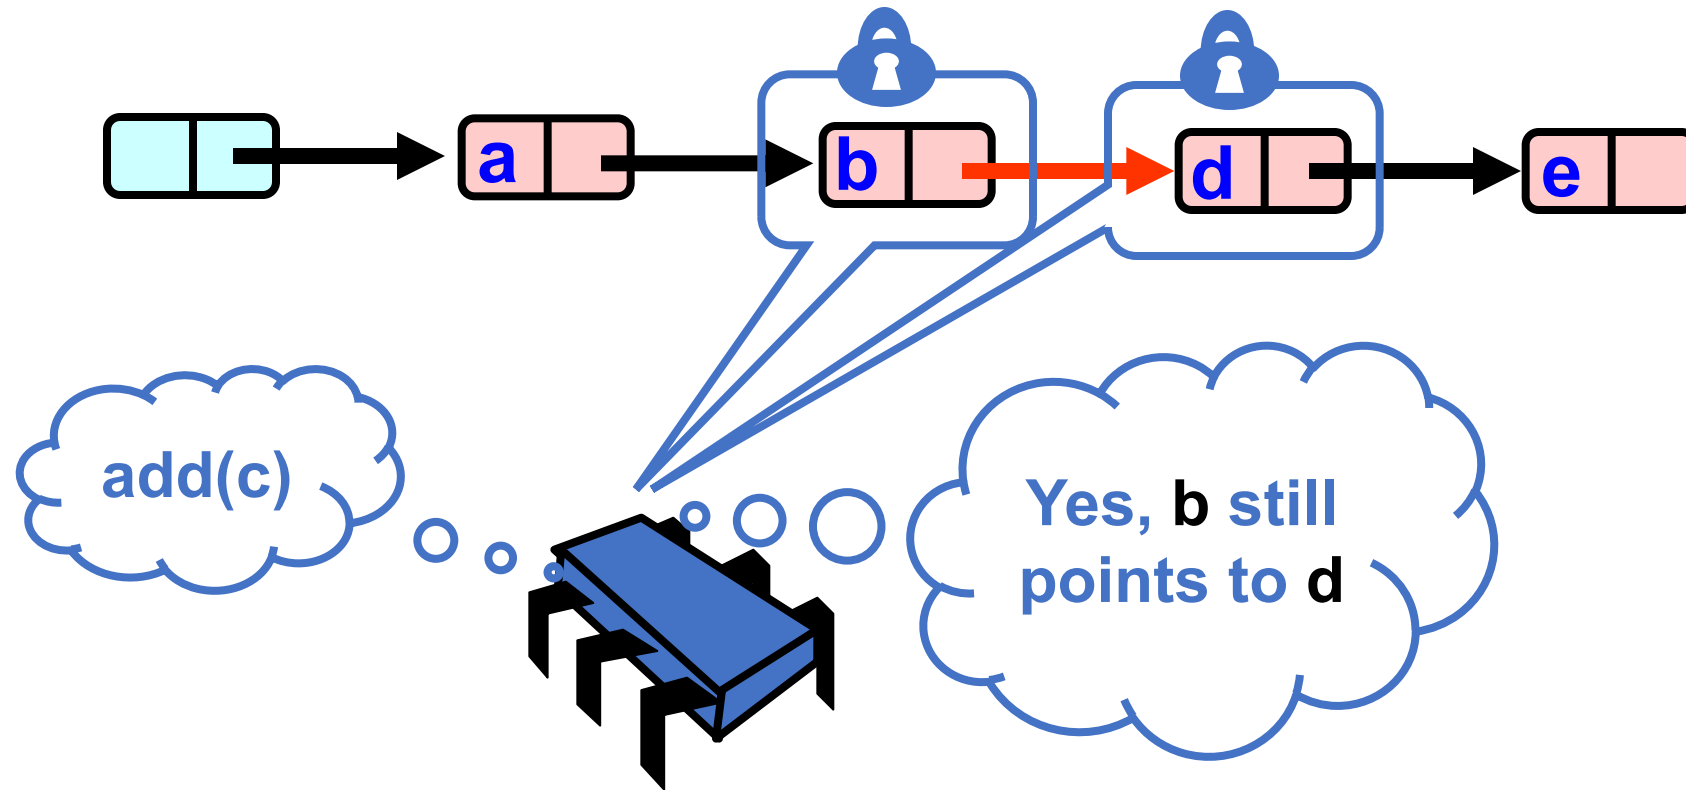


What Else Could Go Wrong?

Validation passes!



Validate Part 2 (while holding locks)



Summary

- We traverse without lock
 - Traversal may access nodes that are locked
 - Its okay because we have atomic pointers!
- We might traverse deleted nodes
 - Its okay because we validate after we obtain locks
 - Two validations:
 - our node is still reachable (it was not deleted)
 - Our insertion point is still valid (no thread has inserted in the meantime)
- We don't actually free node memory, but we put them in a list to be freed later

Summary

- We will make the set lock free next time!