

March 4, 2024

- Finish up GPUs
- Homework 4
- Start on memory models



# Announcements

## Grading

- Midterm is graded:
  - See me or a TA in office hours to discuss
  - Not providing tests back due to academic integrity violations we had last year
  - Gurpreet will be going over midterm in class on Wednesday (it will not be recorded or uploaded).
- Grades for HW 2 should be out by the end of the week

# Announcements

## **Wednesday:**

- I will be gone due to faculty recruiting
  - Jessica and Gurpreet will give the lecture
  - Half going over the midterm
  - Half going over barriers (needed for HW 5)

# Announcements

SETs are out, please do them! It helps us out a lot

Both bad things and good things!

# Announcements

## **Final**

- Allowed 3 pages of notes, front and back
- Similar to the midterm (50% longer)
- Time: Monday March 18: 7:30 – 10:30

# Announcements

## **HW 5**

- Released next week (March 11)
- You will have 9 days to do it. No extensions on this one.

# Previous quiz + Review

# Previous quiz + Review

What is a warp?

- 
- ☐ Group of 32 threads in a GPU
  - ☐ Group of 16 threads in a GPU
  - ☐ Group of 2 threads in a GPU
  - ☐ Group of some threads in a GPU
-





# Previous quiz + Review

Like the CPU cache, the Load/Store Unit reads in memory in chunks. Is this affirmation true or false?

---

☐ True

---

☐ False


4 cores are accessing memory. What can happen

**Read contiguous values**

Like the CPU cache, the Load/Store Unit reads in memory in chunks. 16 bytes

Can easily distribute the values to the threads

1 request to GPU memory



# Previous quiz + Review

What we could observe with the demonstrations made in class about memory accesses on GPUs?

# Previous quiz + Review

Why do we need to calculate  $\text{int } i = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$ ?

- ☐ To get the index in the matrix we want to compute in that thread
- ☐ To get the whole matrix we want to compute in that thread
- ☐ To get the index in the matrix we want to compute in that warp

# Chunked Pattern


the first element accessed  
by the 4 threads sharing a  
load store unit. What  
sort of access is this?

array a




Computation  
can easily be  
divided into  
threads

array b



Thread 0 - Blue  
Thread 1 - Yellow  
Thread 2 - Green  
Thread 3 - Orange

array c




How can we fix this

# Stride Pattern


array a



array b



array c



Computation  
can easily be  
divided into  
threads

Thread 0 - Blue  
Thread 1 - Yellow  
Thread 2 - Green  
Thread 3 - Orange

# Stride Pattern


What sort of pattern is this?

array a




Computation  
can easily be  
divided into  
threads

array b



Thread 0 - Blue  
Thread 1 - Yellow  
Thread 2 - Green  
Thread 3 - Orange

array c





# Homework notes

- Framerate will not be steady
  - This is the real world
- Cached pages
  - Incognito mode
  - TODO
- Serving webpage
- Small demo

```
openssl req -x509 -newkey rsa:4096 -keyout key.pem -out cert.pem -days 365 -nodes
```

# WebGPU

# WebGPU

- The language is wgsł
  - It is new, there are not many examples
  - Official specification is here: <https://www.w3.org/TR/WGSL/>

# WebGPU

- wgsi is NOT javascript
- Javascript is interpreted: not possible on GPUs
- wgsi is compiled
  - into Vulkan on Linux
  - into Metal on Apple
  - into HLSL on Windows
- No printing (so GPU code can be difficult to debug)

# WebGPU

- variables (optional types):

*var <name> = <value>;*

*var cluster\_dist = 3.0;*

*var <name> : <type> = <value>;*

*var cluster\_dist : f32 = 3.0;*

# WebGPU

- types:
  - i32
  - u32
  - f32
  - vec2<f32>
  - array<type>

```
struct Particle {  
    pos : vec2<f32>;  
};
```

```
struct Particles {  
    particles : array<Particle>;  
};
```

```
var index_pos : vec2<f32> = particlesA.particles[index].pos;
```

- structures

```
var index : u32 = GlobalInvocationID.x;
```

- Built-ins (global id) *you have one thread for each particle!*

# WebGPU

- Built in functions:
  - `arrayLength`
  - `sqrt`

# WebGPU

For loops:

```
for (var i : u32 = 0u; i < arrayLength(&particlesA.particles); i = i + 1u) {  
    ...  
}
```



# WebGPU

- Types can be frustrating
- But compiler errors will help you, and you can do casts.

Some final GPU thoughts

# primitive names

- Why threadIdx.x?
- Because you can pass in a 3D range (x,y,z). Historically threads mapped to pixels (or voxels)

# Consequences of Warps

- Branches can cause significant performance impacts, Demo

Locks on GPU?

# What is the right way to program GPUs?

- Still an open question!

# Moving on to memory consistency!

- One of my favorite topics!

# Moving on to memory consistency!

- One of my favorite topics!
- What do other people think?

*Look, memory ordering pretty much is  
the rocket science of CS,*

Linus Torvalds



# Memory Consistency

- We have been very strict about using atomic types in this class
  - and the methods (.load and .store)
  - why?
- Architectures do very strange things with memory loads and stores
- Compilers do too (but we won't talk too much about them today)
- C++ gives us sequential consistency if we use atomic types and operations
- What do we remember sequential consistency from?

# Sequential consistency for atomic memory

- Let's play our favorite game:

**Global variable:**

```
atomic_int x(0);  
atomic_int y(0);
```

**Thread 0:**

```
x.store(1);  
y.store(1);
```

*Is it possible for*

t0 == 0 and t1 == 1



**Thread 1:**

```
int t0 = y.load();  
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

```
x.store(1);
```

```
y.store(1);
```

*Is it possible for*

$t0 == 0$  and  $t1 == 1$



Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

```
int t0 = y.load();
```

```
int t1 = x.load();
```

Global variable:


```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

*Is it possible for*

$t0 == 0$  and  $t1 == 1$



Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

### Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

#### Thread 0:

```
x.store(1);  
y.store(1);
```

```
x.store(1);
```

```
y.store(1);
```

**How about:**

*Is it possible for*

$t0 == 1$  and  $t1 == 0$



#### Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

```
int t0 = y.load();
```

```
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
y.store(1);
```

```
x.store(1);
```

*no where for this one to go!*

**How about:**

*Is it possible for*

$t0 == 1$  and  $t1 == 0$

```
y.store(1);
```

```
int t0 = y.load();
```

```
int t1 = x.load();
```

Thread 1:

```
int t0 = y.load();  
int t1 = x.load();
```

**Global variable:**

```
atomic_int x(0);  
atomic_int y(0);
```

Another test

Can `t0 == t1 == 0`?

**Thread 0:**

```
x.store(1);  
int t0 = y.load();
```

**Thread 1:**

```
y.store(1);  
int t1 = x.load();
```





Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
x.store(1);  
int t0 = y.load();
```

```
x.store(1);
```

```
int t0 = y.load();
```

Thread 1:

```
y.store(1);  
int t1 = x.load();
```

```
y.store(1);
```

```
int t1 = x.load();
```

Global variable:

```
atomic_int x(0);  
atomic_int y(0);
```

Thread 0:

```
x.store(1);  
int t0 = y.load();
```

```
x.store(1);
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
y.store(1);  
int t1 = x.load();
```

```
int t0 = y.load();
```

```
y.store(1);
```

```
int t1 = x.load();
```

*no place for this one!*

# C++

- Plain atomic accesses are documented to be sequentially consistent (SC)
- Why wasn't SC very good for concurrent data structures?
  - Compossibility: two objects that are SC might not be SC when used together
  - Programs contain only 1 shared memory though; no reason to compose different main memories.

# What about ISAs?

- Remember, it is important for us to understand how our code executes on the architecture to write high performing programs
- Lets think about x86
  - Instructions:
    - `MOV %t0 [x]` - loads the value at x to register t0
    - `MOV [y] 1` - stores the value 1 to memory location y

**Global variable:**

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

**Thread 0:**

```
mov [x], 1  
mov %t0, [y]
```

**Thread 1:**

```
mov [y], 1  
mov %t1, [x]
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

```
mov [x], 1
```

```
mov %t0, [y]
```

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```

```
mov [y], 1
```

```
mov %t1, [x]
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```

mov %t1, [x]

mov [x], 1

mov %t0, [y]

mov [y], 1

no place for this event!

# What if we actually run this code?

- We'd like to be able to compile atomic instructions just to regular ISA loads and stores



# Schedule

- Memory consistency models:
  - **Total store order**
  - Relaxed memory consistency
  - Examples

Thread 0:

mov [x], 1

mov %t0, [y]

Core 0

Thread 1:

mov [y], 1

mov %t1, [x]

Core 1



Thread 0:

mov %t0, [y]

Core 0

mov [x], 1

execute first instruction  
what happens to the stores?

Thread 1:

mov %t1, [x]

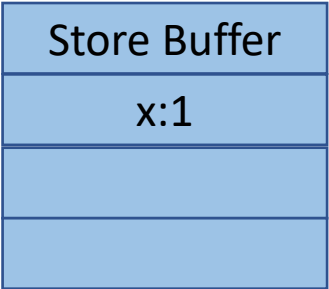

Core 1

mov [y], 1



Thread 0:

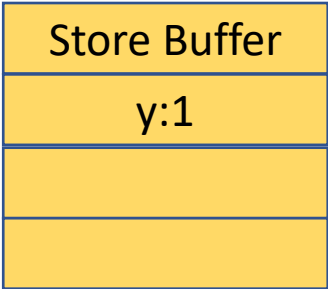

```
mov %t0, [y]
```



X86 cores contain a store buffer; holds stores before going to main memory

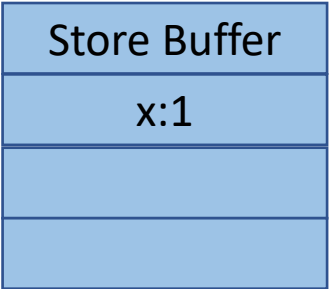
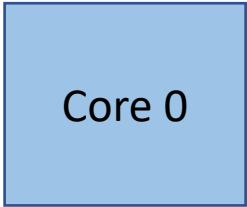
Thread 1:

```
mov %t1, [x]
```



Thread 0:

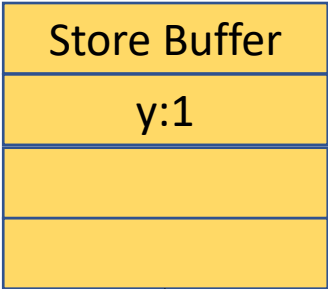
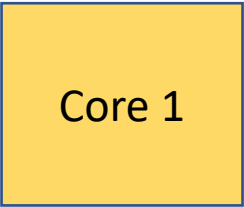
```
mov %t0, [y]
```



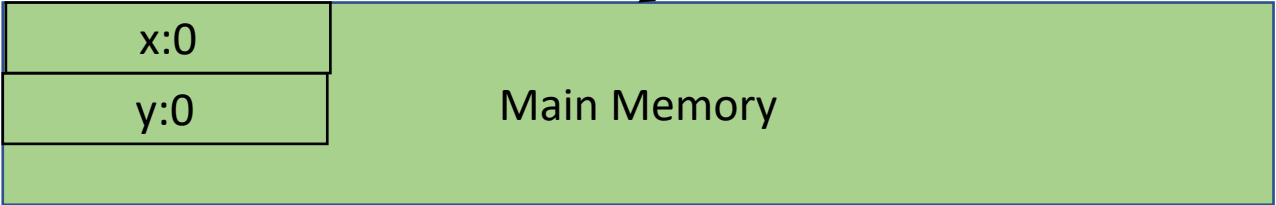
X86 cores contain a store buffer; holds stores before going to main memory

Thread 1:

```
mov %t1, [x]
```

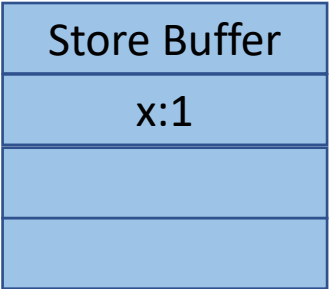
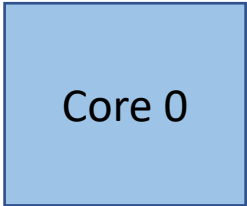


eventually they flush to main memory



Thread 0:

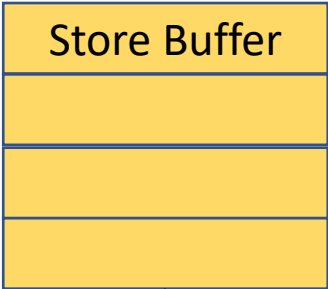
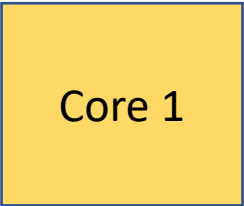
```
mov %t0, [y]
```




X86 cores contain a store buffer; holds stores before going to main memory

Thread 1:

```
mov %t1, [x]
```



eventually they flush to main memory



Thread 0:

mov [x], 1

mov %t0, [y]

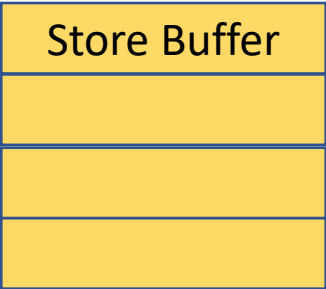
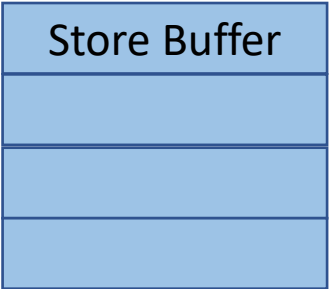
rewind

Thread 1:

mov [y], 1

mov %t1, [x]

Core 0



Core 1





Thread 0:

mov %t0, [y]

Thread 1:

mov %t1, [x]

execute first instruction





Thread 0:

mov %t0, [y]

Core 0

Store Buffer
x:1

values get stored in SB

Thread 1:

mov %t1, [x]

Store Buffer
y:1



Core 1

x:0	Main Memory
y:0	

Thread 0:

Thread 1:

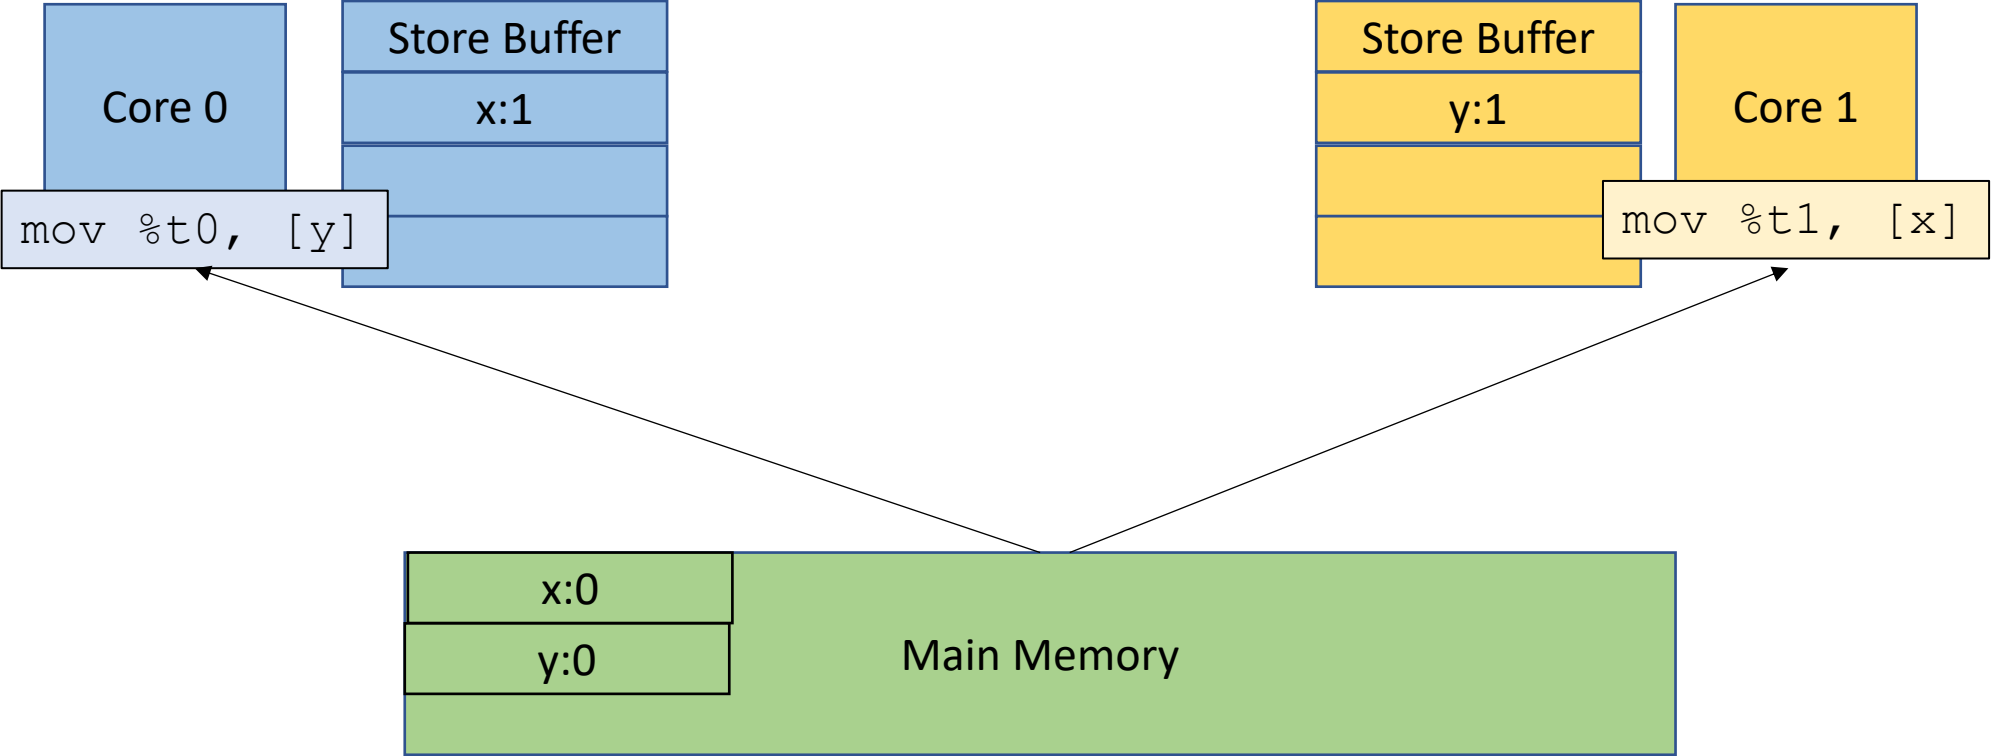
Execute next instruction



Thread 0:

Thread 1:

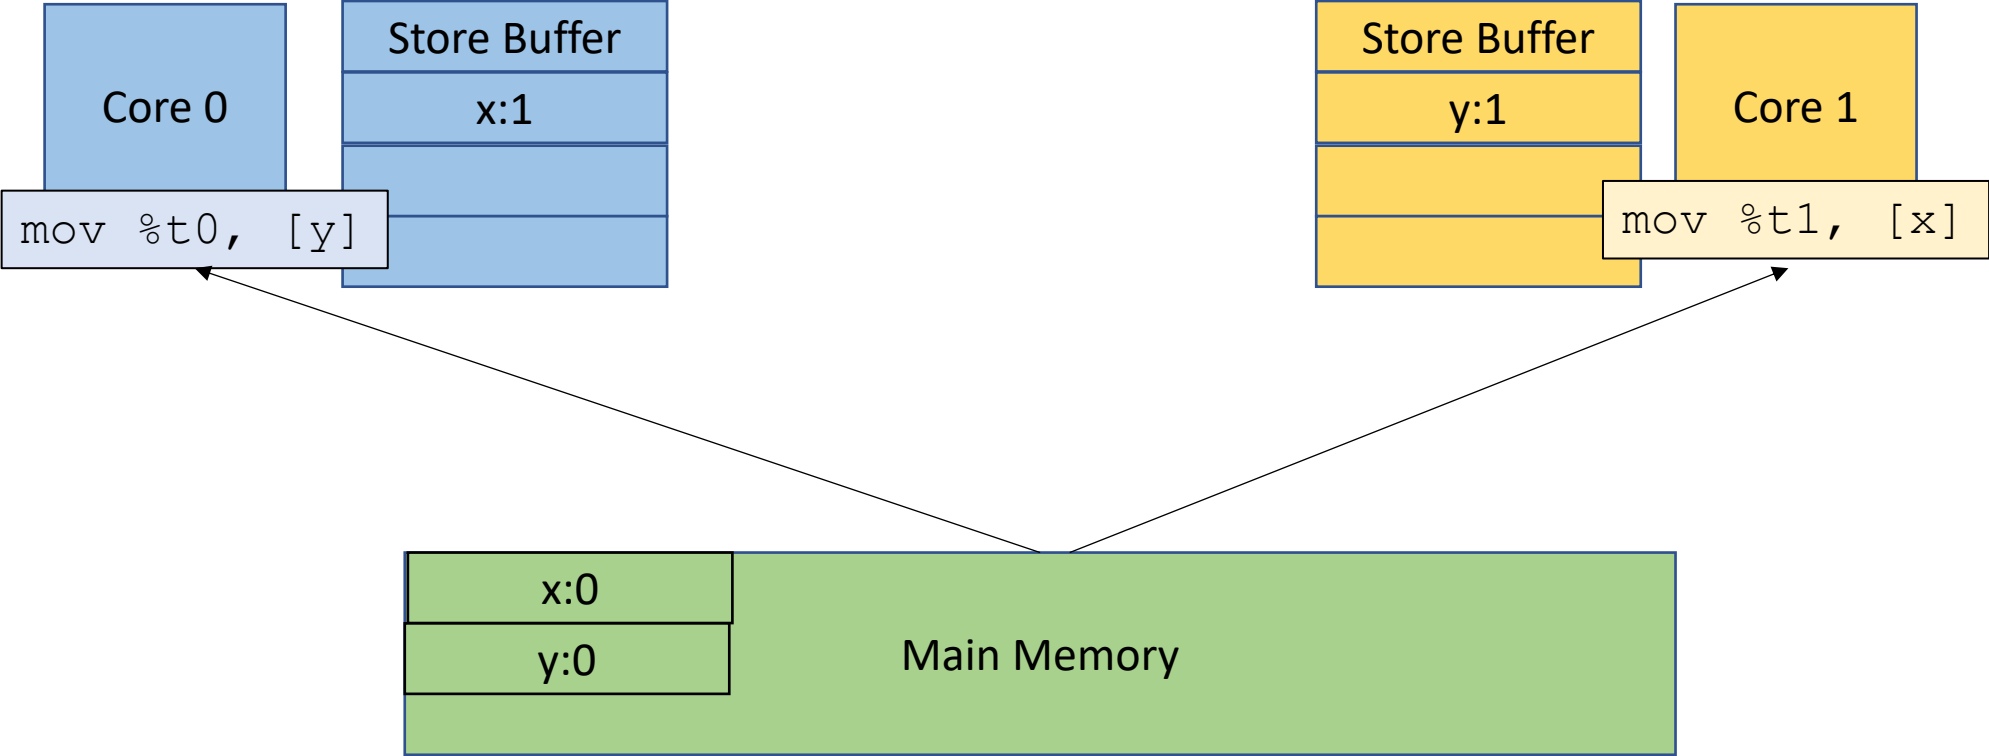
Values get loaded from memory



Thread 0:

Thread 1:

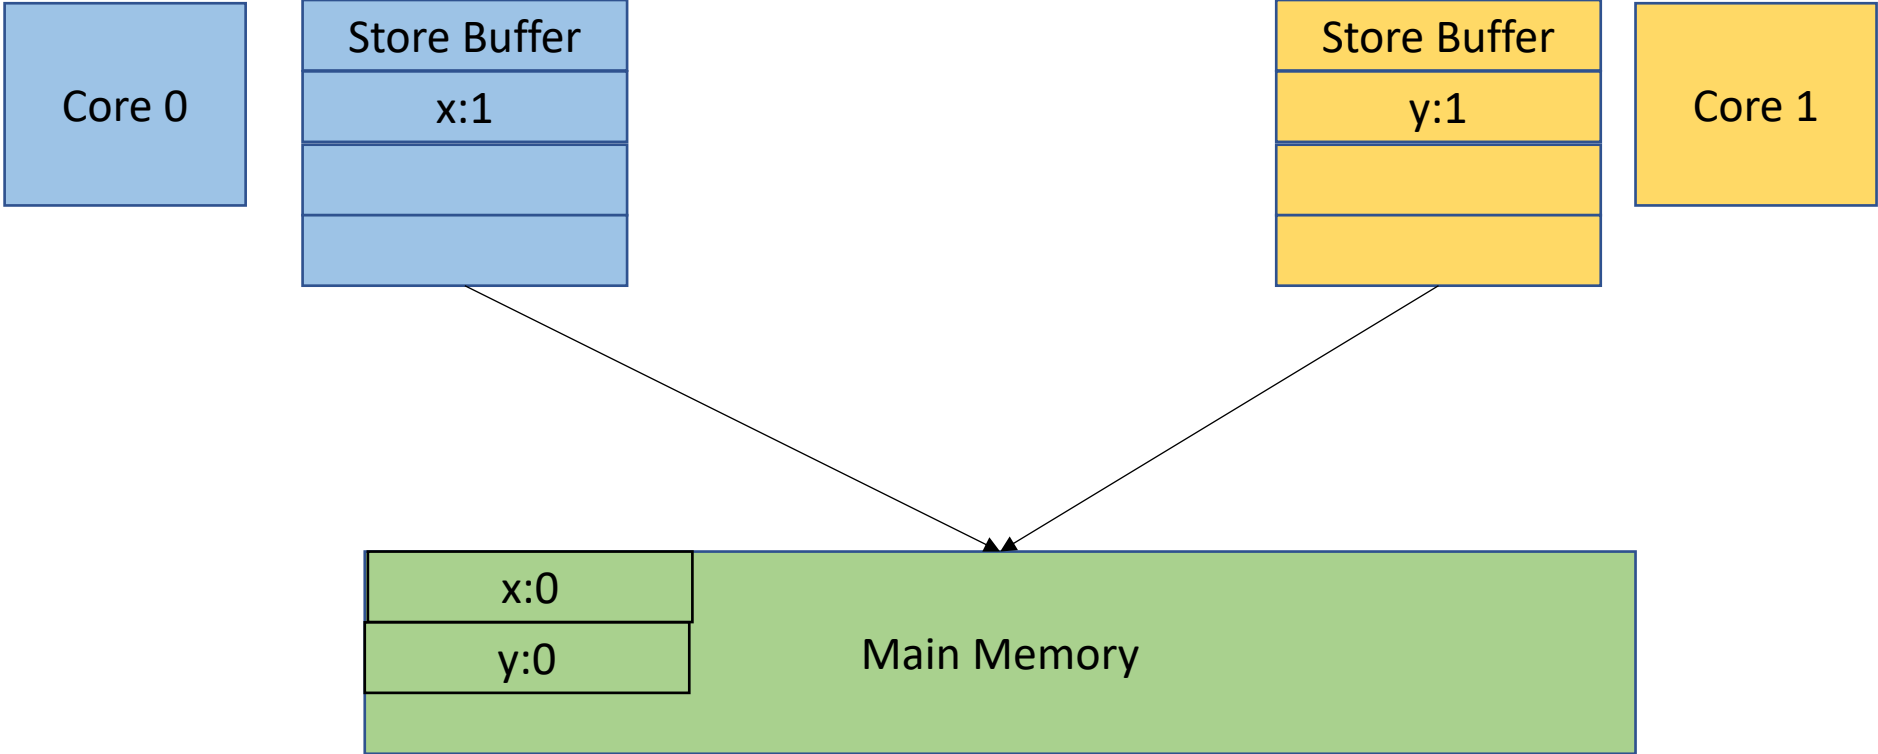
we see `t0 == t1 == 0!`



Thread 0:

Thread 1:

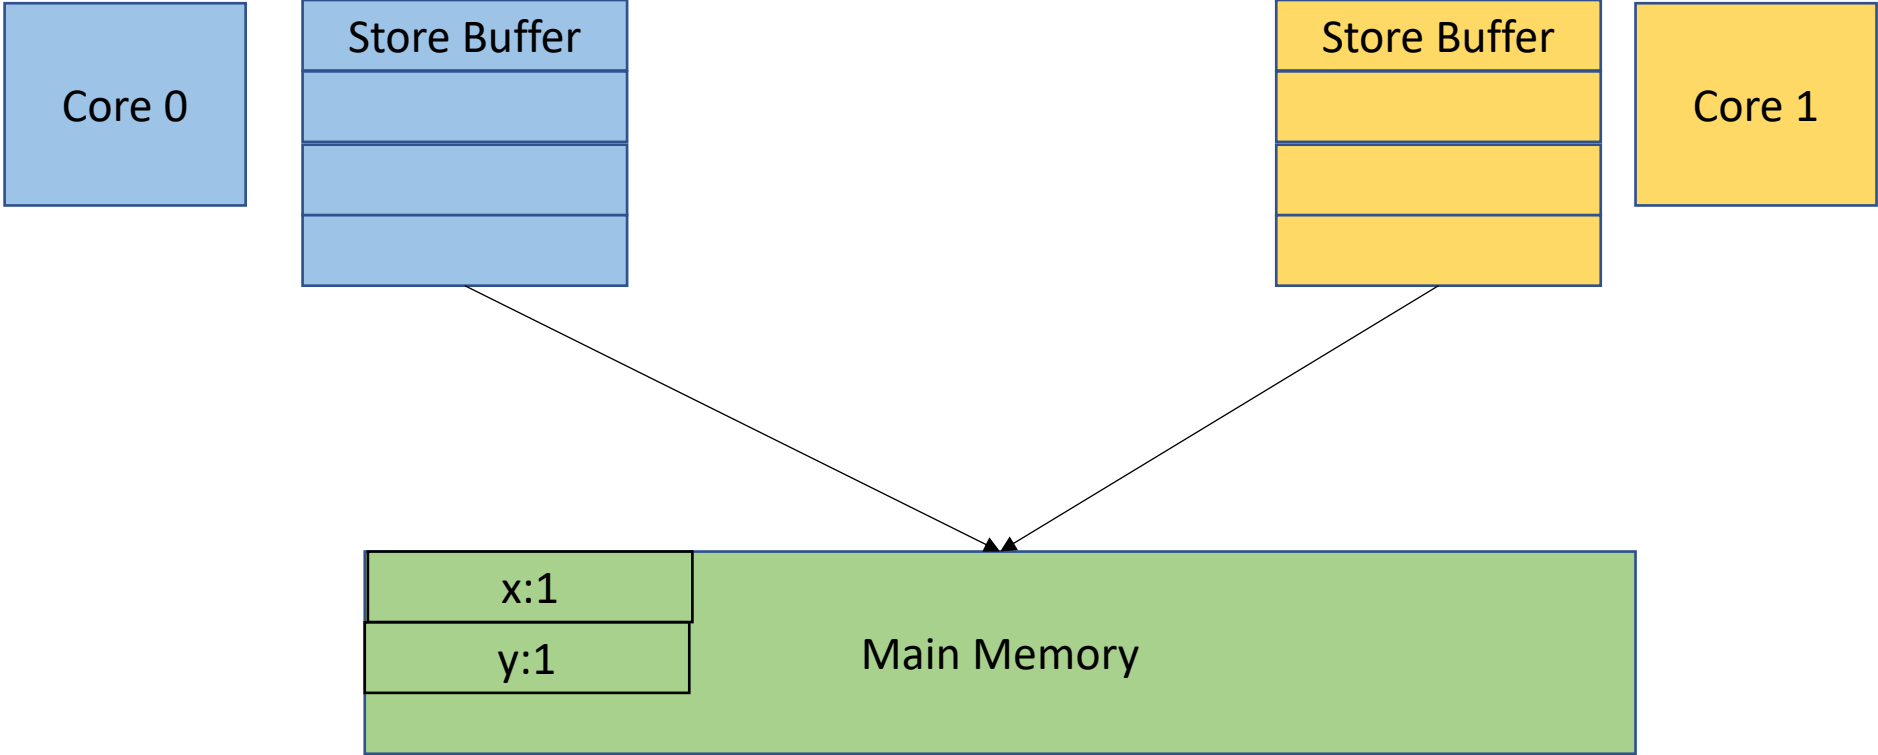
Store buffers are drained eventually



Thread 0:

Thread 1:

Store buffers are drained eventually  
but we've already done our loads



# Our first relaxed memory execution!

- also known as weak memory behaviors
- An execution that is NOT allowed by sequential consistency
- A memory model that allows relaxed memory executions is known as a relaxed memory model
  - X86 has a relaxed memory model due to store buffering
  - If you restrict yourself to use only default atomic operations, C++ has does NOT have a weak memory model

# Litmus tests

- Small concurrent programs that check for relaxed memory behaviors
- Vendors have a long history of under documented memory consistency models
- Academics have empirically explored the memory models
  - Many vendors have unofficially endorsed academic models
  - X86 behaviors were documented by researchers before Intel!



# Litmus tests

This test is called “store buffering”

## Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

## Thread 1:

```
mov [y], 1  
mov %t1, [x]
```

Can `t0 == t1 == 0`?

# Restoring sequential consistency



- It is typical that relaxed memory models provide special instructions which can be used to disallow weak behaviors.
- These instructions are called Fences
- The X86 fence is called `mfence`. It flushes the store buffer.

Thread 0:

mov [x], 1

mfence

mov %t0, [y]





Thread 1:

mov [y], 1

mfence


mov %t1, [x]



Thread 0:

mfence

mov %t0, [y]




Execute first instruction

Thread 1:

mfence

mov %t1, [x]



Thread 0:

mfence

mov %t0, [y]

Core 0

Store Buffer
x:1

Thread 1:

mfence

mov %t1, [x]

Core 1

Store Buffer
y:1

Values go into the store buffer

x:0	Main Memory
y:0	

Thread 0:

Thread 1:

Execute next instruction

```
mov %t0, [y]
```

```
mov %t1, [x]
```

Core 0

mfence

Store Buffer

x:1

Store Buffer

y:1

Core 1

mfence

x:0

y:0

Main Memory






Thread 0:

Thread 1:

store buffers are flushed

```
mov %t0, [y]
```

```
mov %t1, [x]
```







Thread 0:

Thread 1:

store buffers are flushed

```
mov %t0, [y]
```

```
mov %t1, [x]
```








Thread 0:

Thread 1:

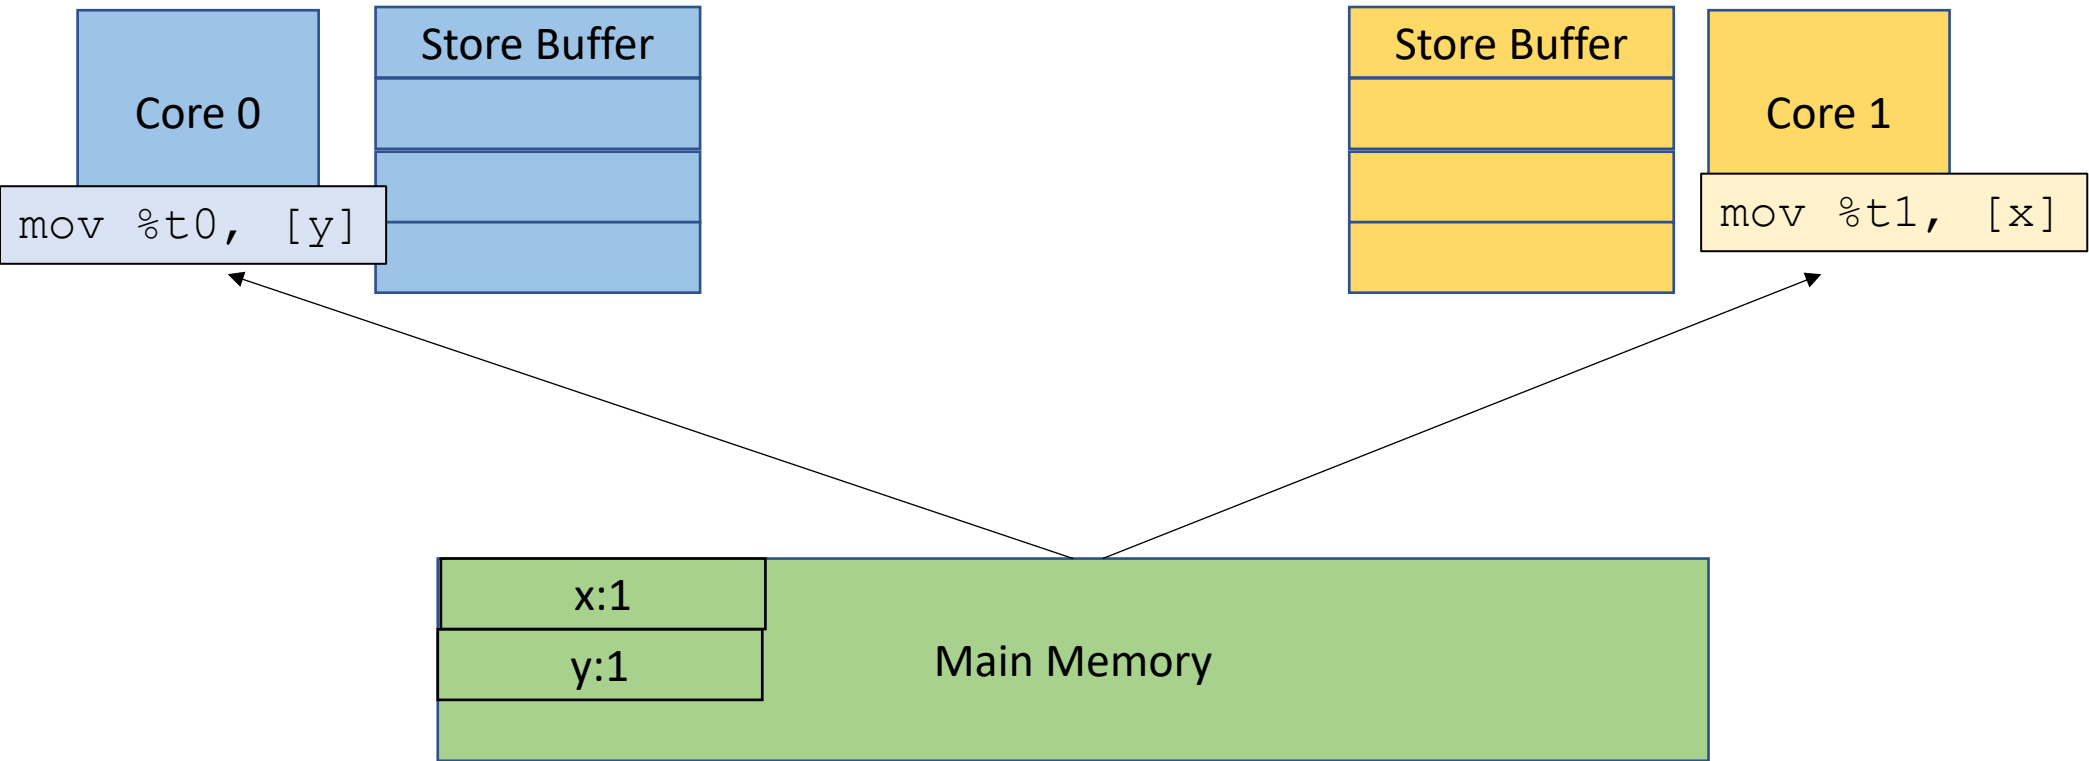
execute next instruction



Thread 0:

Thread 1:

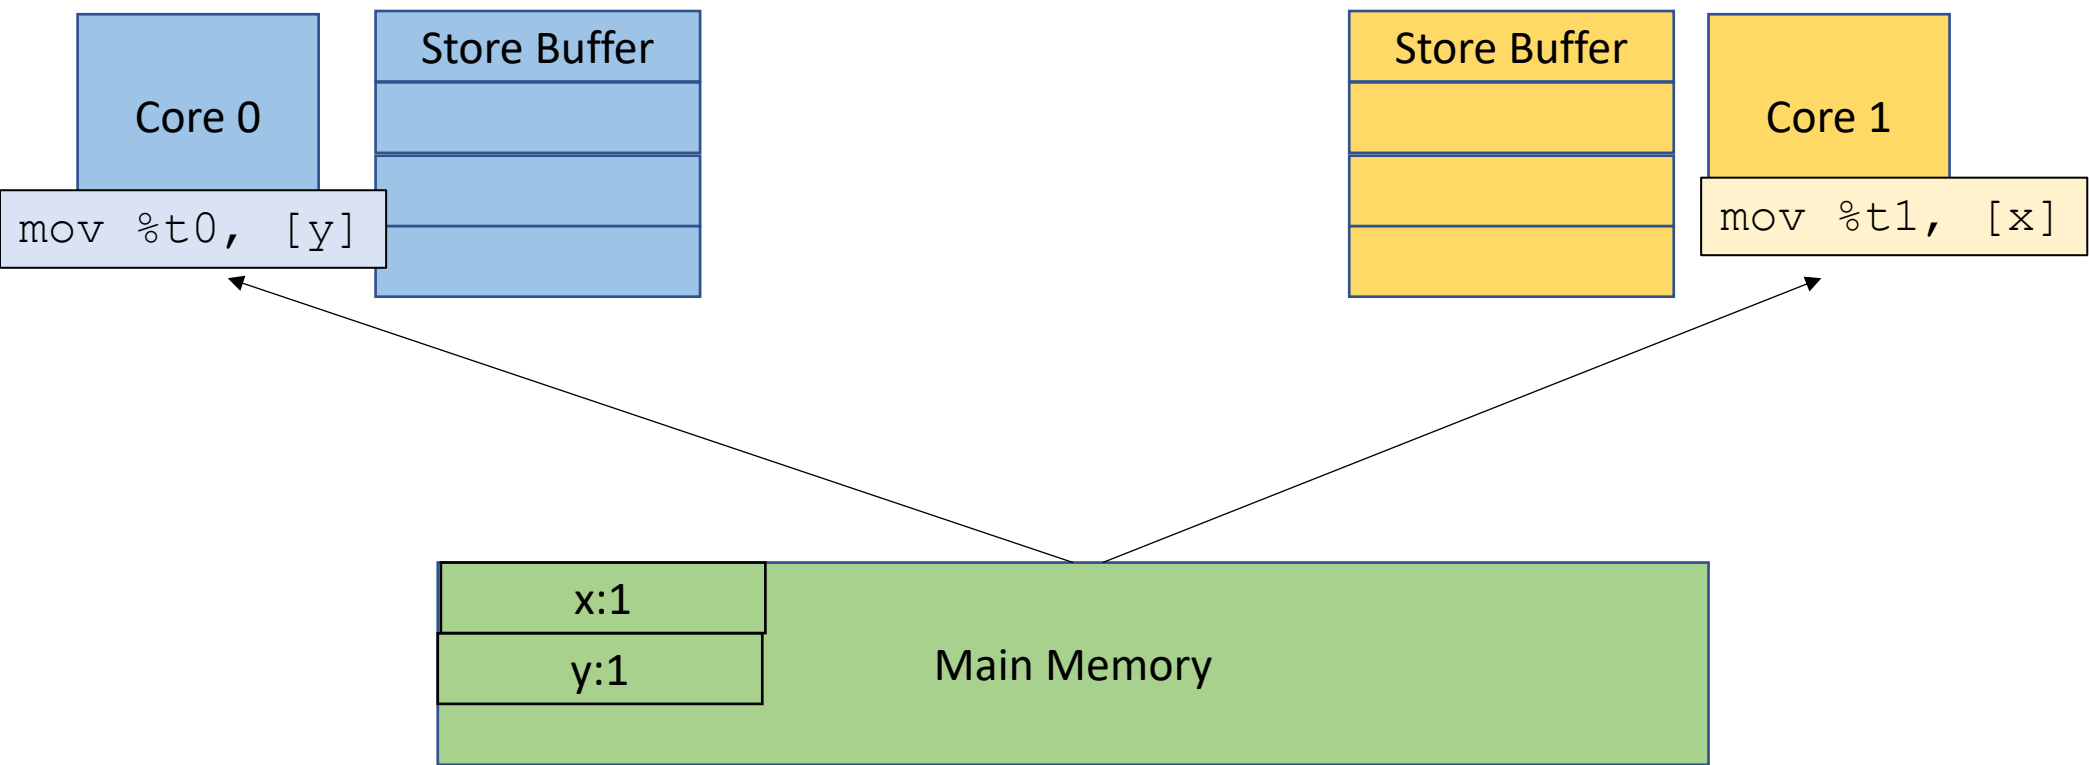
values are loaded from memory



Thread 0:

Thread 1:

We don't get the problematic behavior: `t0 == t1 == 0`



Next example

Thread 0:

```
mov [x], 1
```

```
mov %t0, [x]
```



single thread  
same address

possible outcomes:

t0 = 1

t0 = 0

Which one do you expect?





Thread 0:

```
mov [x], 1
```

```
mov %t0, [x]
```


How does this execute?



Thread 0:

execute first instruction

mov %t0, [x]



Thread 0:

Store the value in the store buffer


```
mov %t0, [x]
```





Thread 0:

Next instruction





Thread 0:

Where to load??

Store buffer?

Main memory?





Thread 0:

Where to load??

Threads check store buffer before going to main memory

It is close and cheap to check.



# Memory Consistency

- How to specify a relaxed memory model?
- We can do it operationally
  - by constructing a high-level machine and reasoning about operations through the machine.
  - or we can talk about instructions that are allowed to "break" program order.

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
mov [x], 1  
mov %t0, [y]
```

Thread 1:

```
mov [y], 1  
mov %t1, [x]
```



*We will annotate instructions with S for store, and L for loads*

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```



*We will annotate instructions with S for store, and L for loads*

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

```
S:mov [x], 1
```

```
L:mov %t0, [y]
```

Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

```
L:mov %t1, [x]
```



### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test


Can `t0 == t1 == 0`?

#### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

#### Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```



```
S:mov [y], 1
```

Now we make a new rule:

S(tores) followed by a L(oad)  
do not have to follow program order



### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?

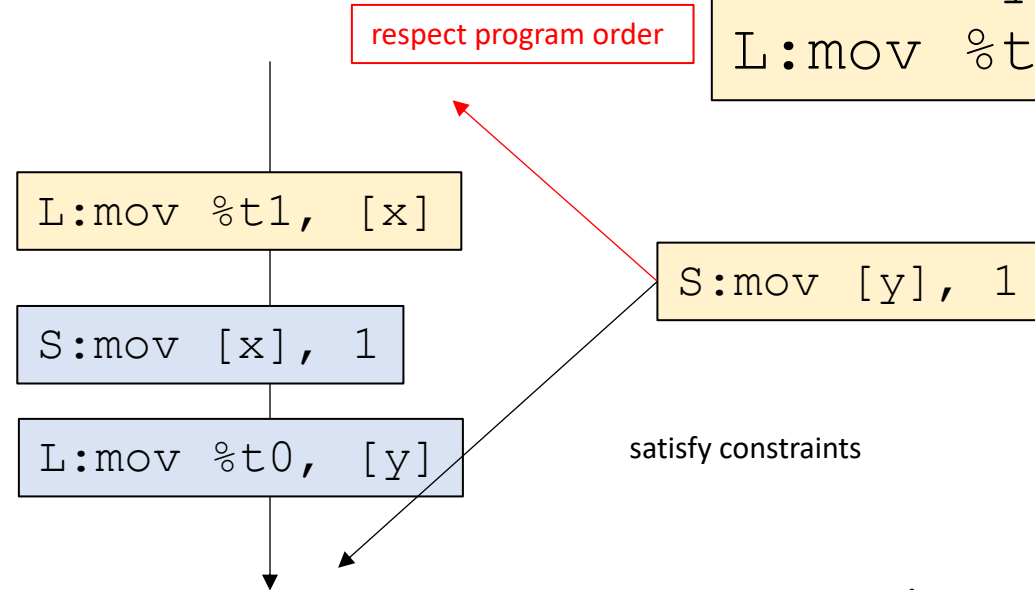
#### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

#### Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

we can ignore this condition!!



Now we make a new rule:

S(tores) followed by a L(oad)  
do not have to follow program order

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == t1 == 0`?


#### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

#### Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

we can ignore this condition!!



Now we can satisfy the condition!

## Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

Lets peak under the hood here


Another test

Can `t0 == t1 == 0`?

### Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

we can ignore this condition!!



```
int x[1] = {0};
int y[1] = {0};
```

```
int x[1] = {0};
int y[1] = {0};
```

## Can $t_0 == t_1 == 0$ ?

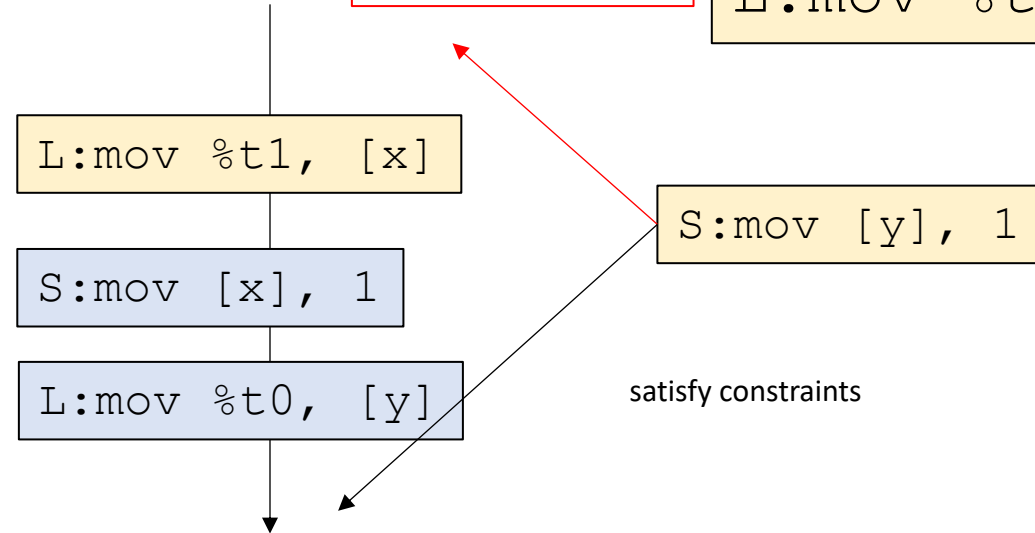
```
S:mov    [x], 1
L:mov    %t0, [y]
```

```
S:mov    [x], 1
L:mov    %t0, [y]
```

```
S:mov  [y], 1
L:mov  %t1, [x]
```

```
S:mov  [y], 1
L:mov  %t1, [x]
```

respect program order



Global timeline is when the Store operation becomes visible to other threads

Global timeline is when the Store operation becomes visible to other threads

## Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

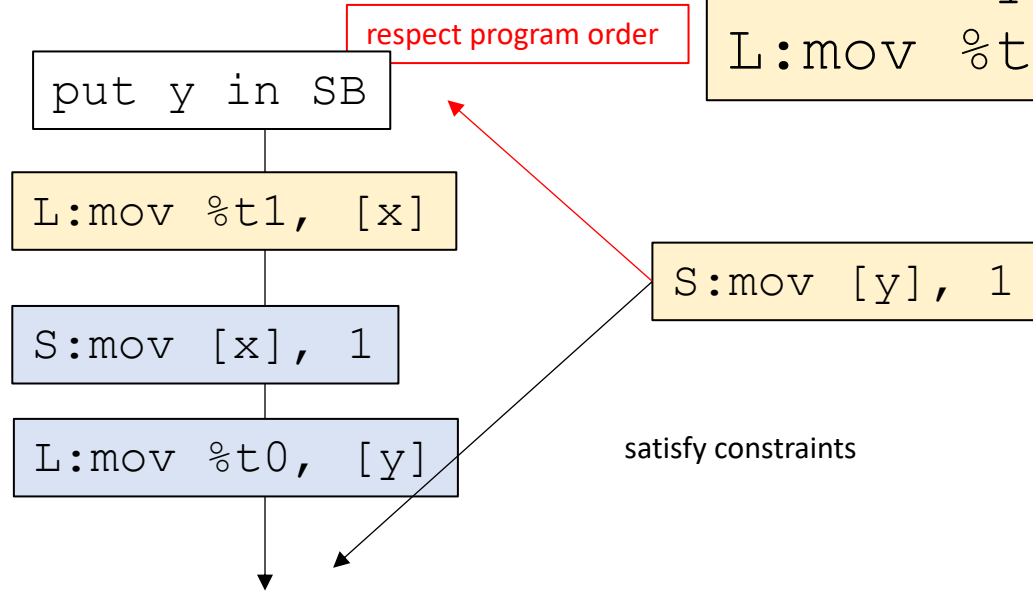
Lets peak under the hood here

Global timeline is when the  
Store operation becomes visible  
to other threads

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

#### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [y]
```

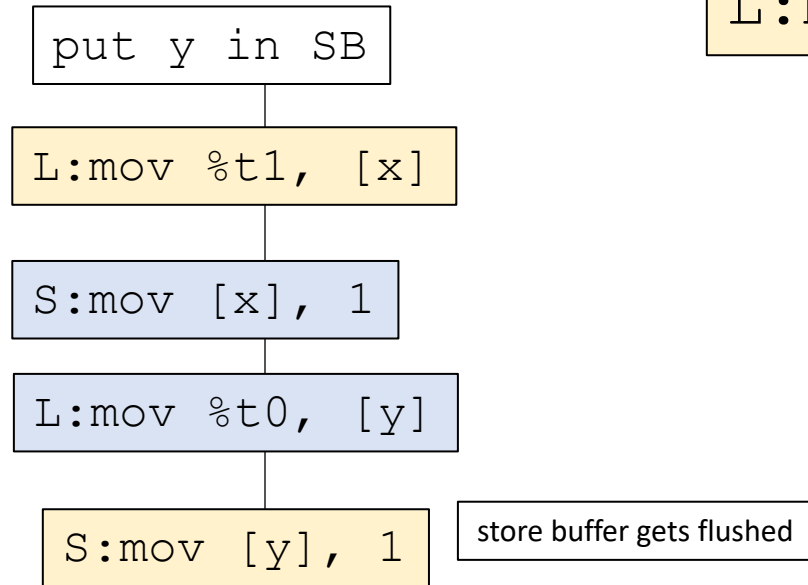
Lets peak under the hood here

Global timeline is when the  
Store operation becomes visible  
to other threads

Another test

Can `t0 == t1 == 0`?

we can ignore this condition!!



#### Thread 1:

```
S:mov [y], 1  
L:mov %t1, [x]
```

# Questions


- Can stores be reordered with stores?

Thread 0:

mov [x], 1

mov [y], 1

Core 0






Thread 0:

mov [y], 1

execute the first instruction



Thread 0:

mov [y], 1

value goes into store buffer

Core 0

Store Buffer
x:1

x:0	Main Memory
y:0	

Thread 0:

mov [y], 1

execute next instruction



Core 0

Store Buffer
x:1

x:0	Main Memory
y:0	


Thread 0:

execute next instruction




Thread 0:

value goes into the store buffer




Thread 0:

Core 0




On x86, the store buffer trains in a FIFO way:  
thus stores cannot be reordered




Thread 0:

Core 0




On x86, the store buffer trains in a FIFO way:  
thus stores cannot be reordered




Thread 0:

Core 0



On x86, the store buffer trains in a FIFO way:  
thus stores cannot be reordered





# Questions

- Can stores be reordered with stores?
- How do we make rules about mfence?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
mfence  
L:mov %t0, [y]
```

```
S:mov [x], 1
```

```
mfence
```

```
L:mov %t0, [y]
```

Another test

Can `t0 == t1 == 0`?

Thread 1:

```
S:mov [y], 1  
mfence  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

```
mfence
```

```
L:mov %t1, [x]
```

Rules: S(tores) followed by a L(oad)  
do not have to follow program order.

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

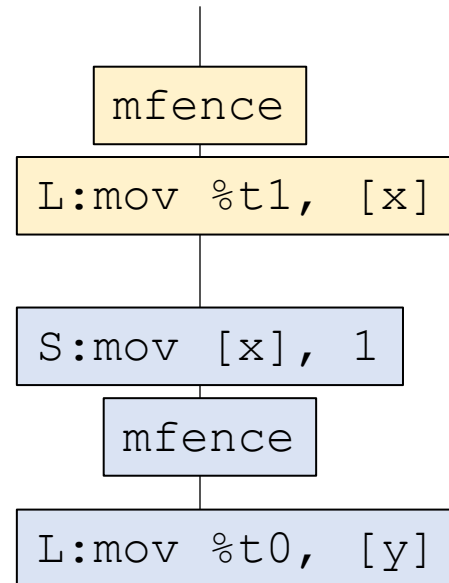
#### Thread 0:

```
S:mov [x], 1  
mfence  
L:mov %t0, [y]
```

*So we can't  
reorder  
this instruction  
at all!*

Another test

Can `t0 == t1 == 0`?



#### Thread 1:

```
S:mov [y], 1  
mfence  
L:mov %t1, [x]
```

```
S:mov [y], 1
```

Rules:

S(tores) followed by a L(oad)  
do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

# Rules

- Are we done?

Rules:

S(tores) followed by a L(oad)

do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Another test

Can `t0 == 0`?

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

```
S:mov [x], 1
```

```
L:mov %t0, [x]
```



Rules:

S(tores) followed by a L(oad)

do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

Another test  
Can `t0 == 0`?

S:mov [x], 1

where to put this store?

L:mov %t0, [x]

Rules:  
S(tores) followed by a L(oad)  
do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

### Global variable:


```
int x[1] = {0};  
int y[1] = {0};
```

### Thread 0:

```
S:mov [x], 1  
L:mov %t0, [x]
```

Another test

Can `t0 == 0`?



where to put this store?

Rules:

S(tores) followed by a L(oad)  
do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

S(tores) cannot be reordered past L(oads)  
from the same address

# TSO - Total Store Order

## **Rules:**

S(tores) followed by a L(oad)  
do not have to follow program order.

S(tores) cannot be reordered past a fence  
in program order

S(tores) cannot be reordered past L(oads)  
from the same address




# Schedule

- Memory consistency models:
  - Total store order
  - **Relaxed memory consistency**
  - Examples

# Other memory models?

- We can specify them in terms of what reorderings are allowed



If memory access 0 appears before memory access 1 in program order, can it bypass program order?

# Other memory models?

- We can specify them in terms of what reorderings are allowed

		memory access 0	
		L	S
memory access 1	L	NO	NO
	S	NO	NO

## Sequential Consistency

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

# Other memory models?

- We can specify them in terms of what reorderings are allowed

		memory access 0	
		L	S
memory access 1	L	NO	Different address
	S	NO	NO

## **TSO - total store order**

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

# Other memory models?

- We can specify them in terms of what reorderings are allowed

		memory access 0	
		L	S
memory access 1	L	?	?
	S	?	?

## Weaker models?

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

# Other memory models?

- We can specify them in terms of what reorderings are allowed

		memory access 0	
		L	S
memory access 1	L	NO	Different address
	S	NO	Different address

## PSO - partial store order

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

*Allows stores to drain from the store buffer in any order*

# Other memory models?

- We can specify them in terms of what reorderings are allowed

		memory access 0	
		L	S
memory access 1	L	YES	Different address
	S	Different address	Different address

## RMO - Relaxed Memory Order

If memory access 0 appears before memory access 1 in program order, can it bypass program order?

*Very relaxed model!*

# Other memory models?

- FENCE: can always restore order using fences. Accesses cannot be reordered past fences!

		memory access 0	
		L	S
memory access 1	L	NO	NO
	S	NO	NO

## Any Memory Model

If memory access 0 appears before memory access 1 in program order, and there is a FENCE between the two accesses, can it bypass program order?



# Schedule

- Memory consistency models:
  - Total store order
  - Relaxed memory consistency
  - **Examples**

*Global variable:*

```
int x[1] = {0};  
int y[1] = {0};
```

First thing: change our syntax to pseudo code

*Thread 0:*

```
L:mov %t0, [y]  
S:mov [x], 1
```

*Thread 1:*

```
L:mov %t1, [x]  
S:mov [y], 1
```

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

First thing: change our syntax to pseudo code  
You should be able to find natural mappings  
to any ISA

#### Thread 0:

```
L:%t0 = load(y)  
S:store(x,1)
```

#### Thread 1:

```
L:%t1 = load(x)  
S:store(y,1)
```

*Global variable:*

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

*Thread 0:*

```
L:%t0 = load(y)  
S:store(x,1)
```

*Thread 1:*

```
L:%t1 = load(x)  
S:store(y,1)
```

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

Get out our lego bricks and try for sequential consistency

Thread 0:

```
L:%t0 = load(y)  
S:store(x,1)
```

```
L:%t0 = load(y)
```

```
S:store(x,1)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y,1)
```

```
L:%t1 = load(x)
```

```
S:store(y,1)
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can  $t0 == t1 == 1$ ?

Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
S:store(x,1)
```

```
L:%t0 = load(y)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y,1)
```

respect program order

```
S:store(x,1)
```

```
L:%t1 = load(x)
```

```
S:store(y,1)
```

satisfy constraints

Not allowed under sequential consistency!

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?


Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



What about TSO?

memory access 0	
L	S
L	NO
S	NO

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

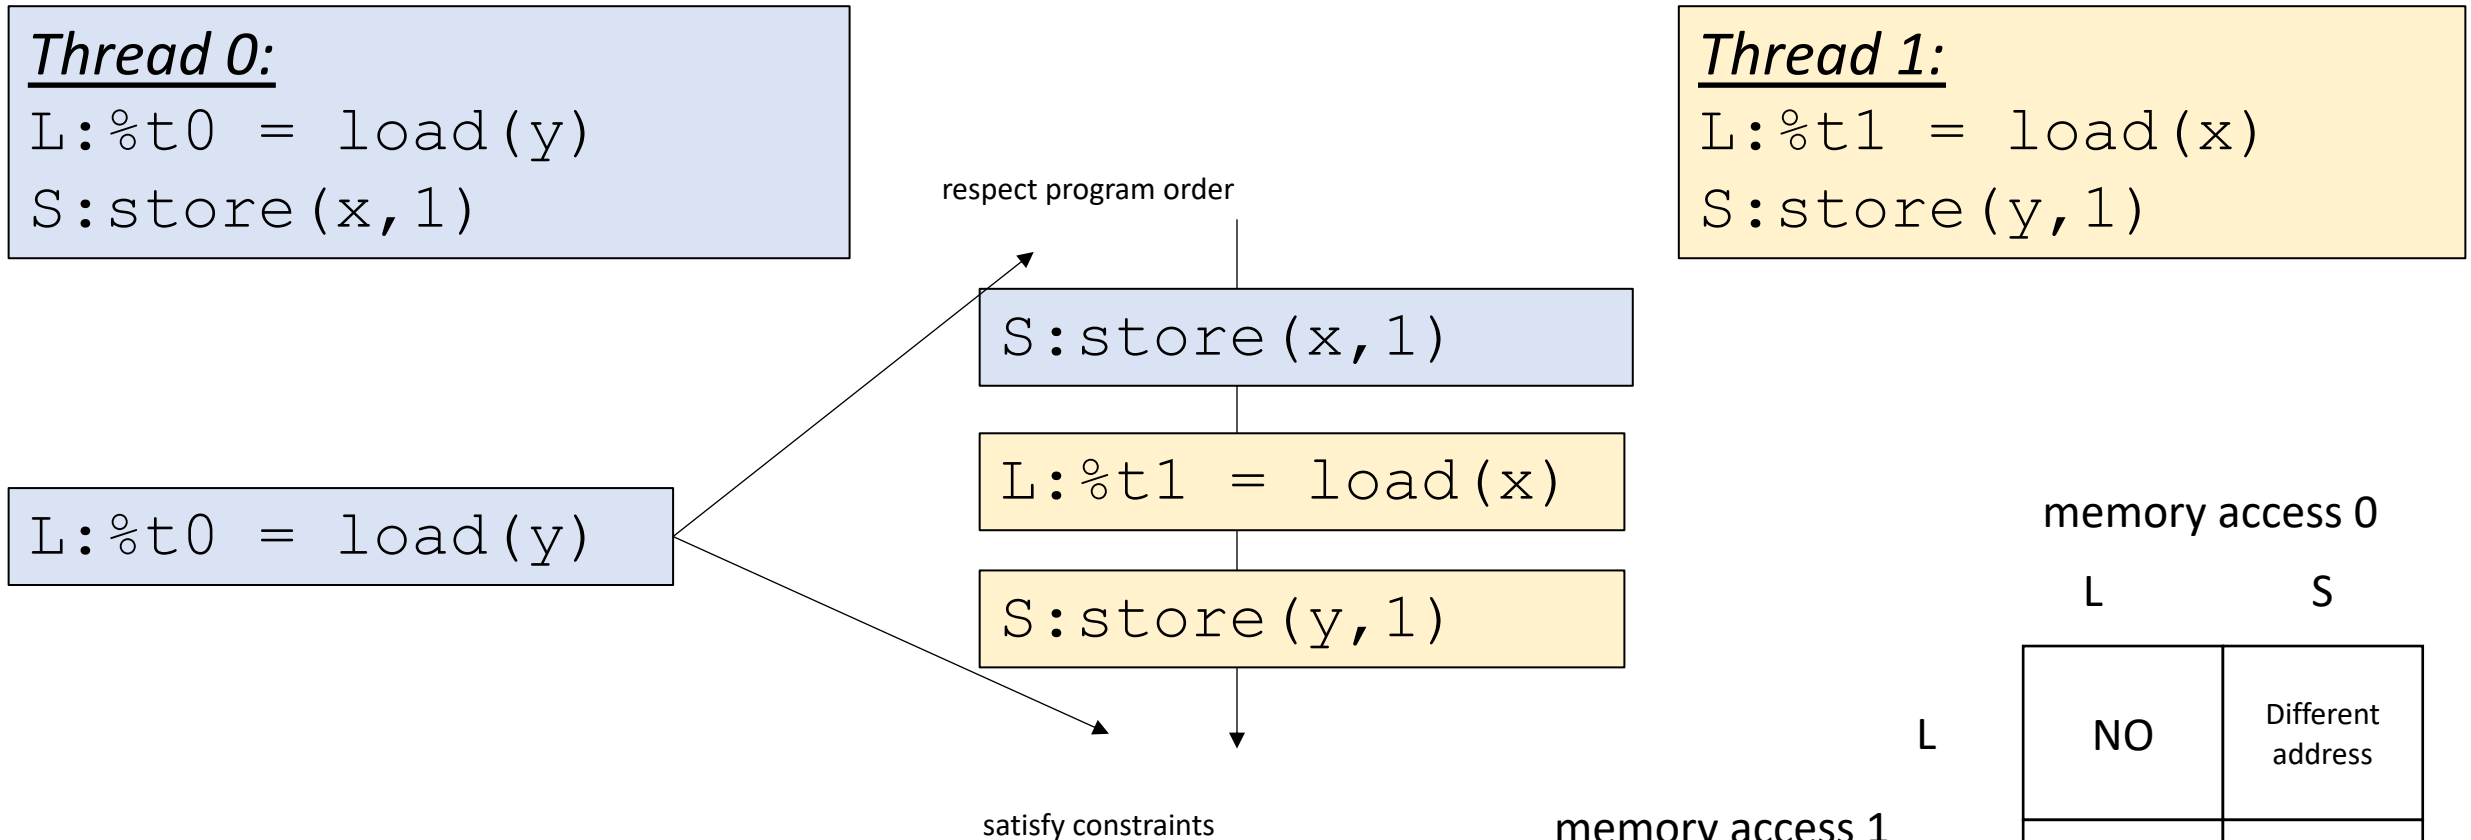
Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



What about TSO? NOT ALLOWED!

memory access 0	
L	S
L	NO
S	NO



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?


Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



What about PSO?

memory access 0	
L	S
L NO	S Different address
S NO	S Different address

## Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

Get out our lego bricks

### Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

```
L:%t0 = load(y)
```

respect program order

```
S:store(x, 1)
```

```
L:%t1 = load(x)
```

```
S:store(y, 1)
```

satisfy constraints

### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```

memory access 0

L S

	L	S
L	NO	Different address
S	NO	Different address

memory access 1

What about PSO? NO!

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

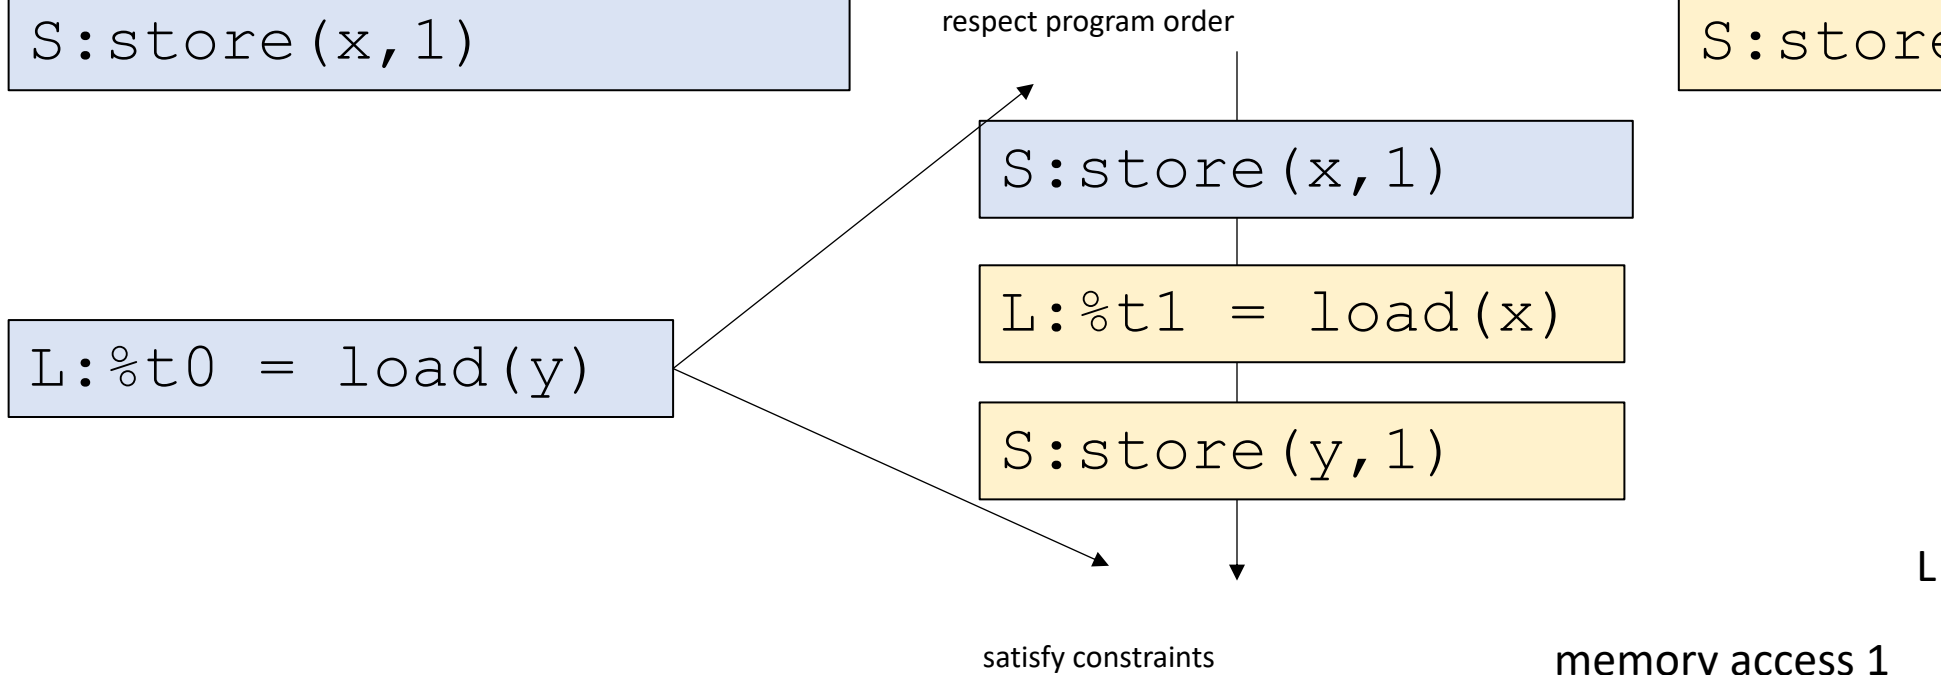
Get out our lego bricks

#### Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

#### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



What about RMO?

memory access 0

L S

L

YES

Different  
address

S

different  
address

Different  
address

memory access 1

## Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

Get out our lego bricks

### Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

```
L:%t0 = load(y)
```

respect program order

```
S:store(x, 1)
```

```
L:%t1 = load(x)
```

```
S:store(y, 1)
```

satisfy constraints

### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```

memory access 0

L S

L

YES

Different  
address

S

different  
address

Different  
address

memory access 1

What about RMO?

## Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

Get out our lego bricks

### Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

```
L:%t0 = load(y)
```

respect program order

```
S:store(x, 1)
```

```
L:%t1 = load(x)
```

```
S:store(y, 1)
```

satisfy constraints

### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```

memory access 0

L S

L

YES

Different  
address

S

different  
address

Different  
address

memory access 1

What about RMO? YES!

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

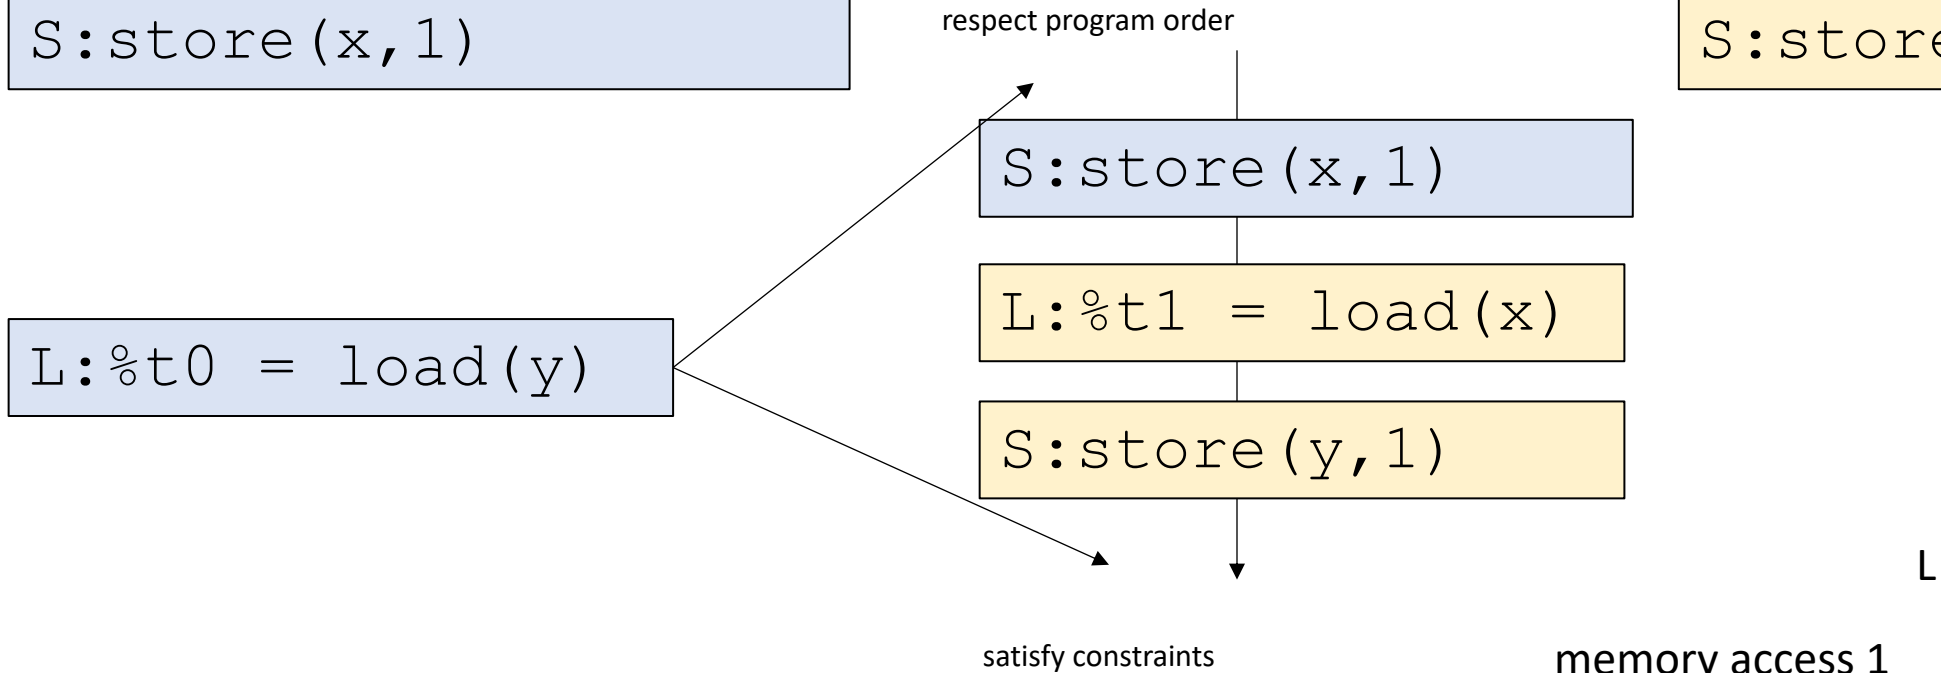
Get out our lego bricks

#### Thread 0:

```
L:%t0 = load(y)  
S:store(x, 1)
```

#### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



How do we disallow the behavior in RMO?

memory access 0	
L	S
L	YES Different address
S	different address Different address

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
fence  
S:store(x, 1)
```

```
L:%t0 = load(y)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```

respect program order

```
S:store(x, 1)
```

```
L:%t1 = load(x)
```

```
S:store(y, 1)
```

satisfy constraints

How do we disallow the behavior in RMO?

memory access 0

L S

L

YES

Different  
address

S

different  
address

Different  
address

memory access 1

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

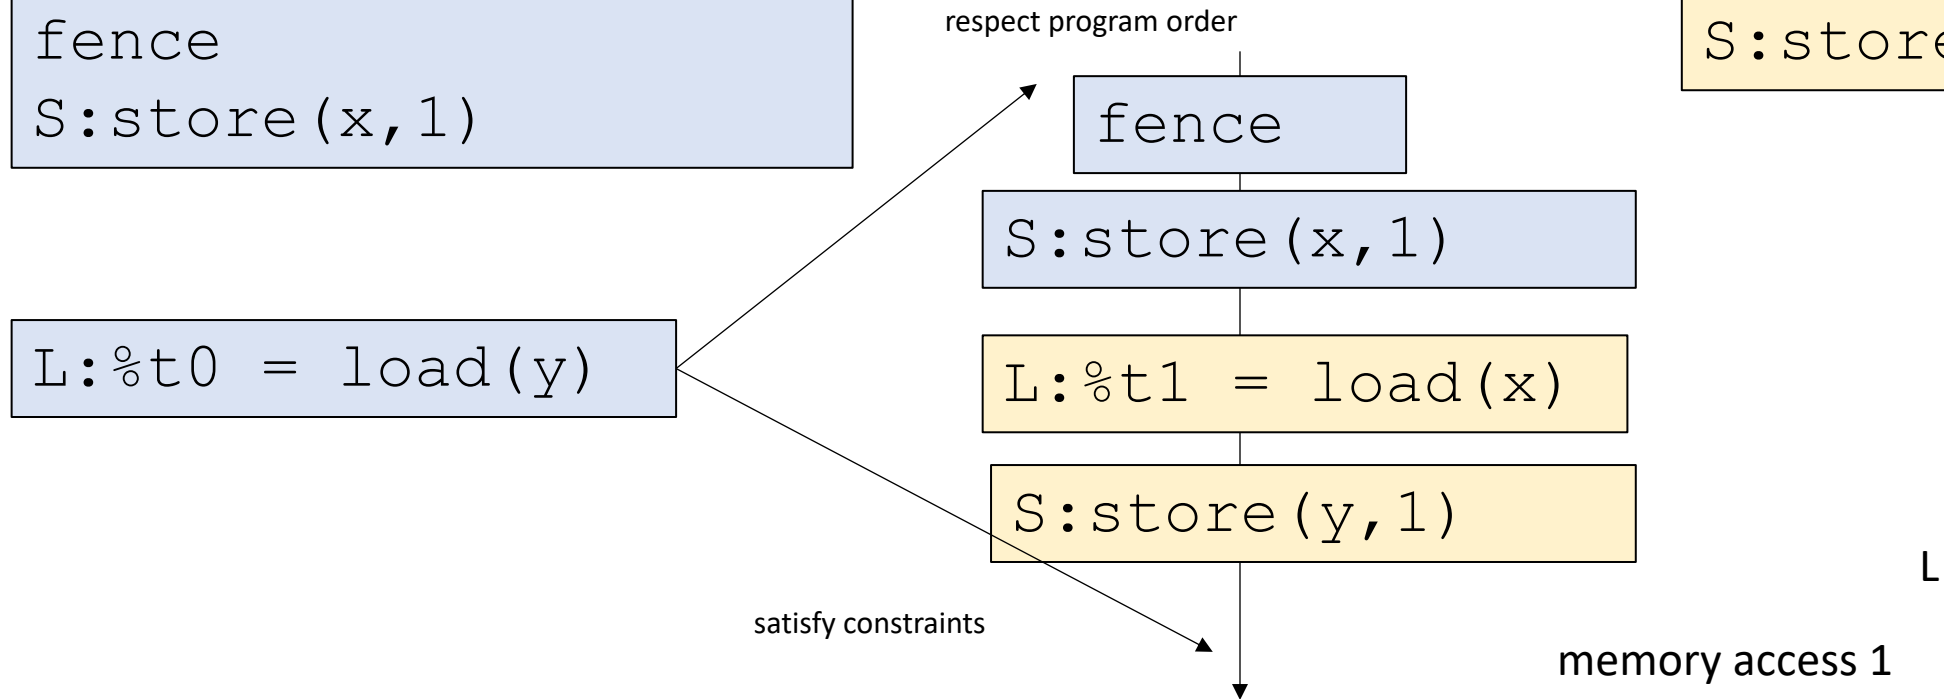
Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
fence  
S:store(x, 1)
```

Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



memory access 0

L S

L

S

	L	S
L	YES	Different address
S	different address	Different address

memory access 1

How do we disallow the behavior in RMO?



### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

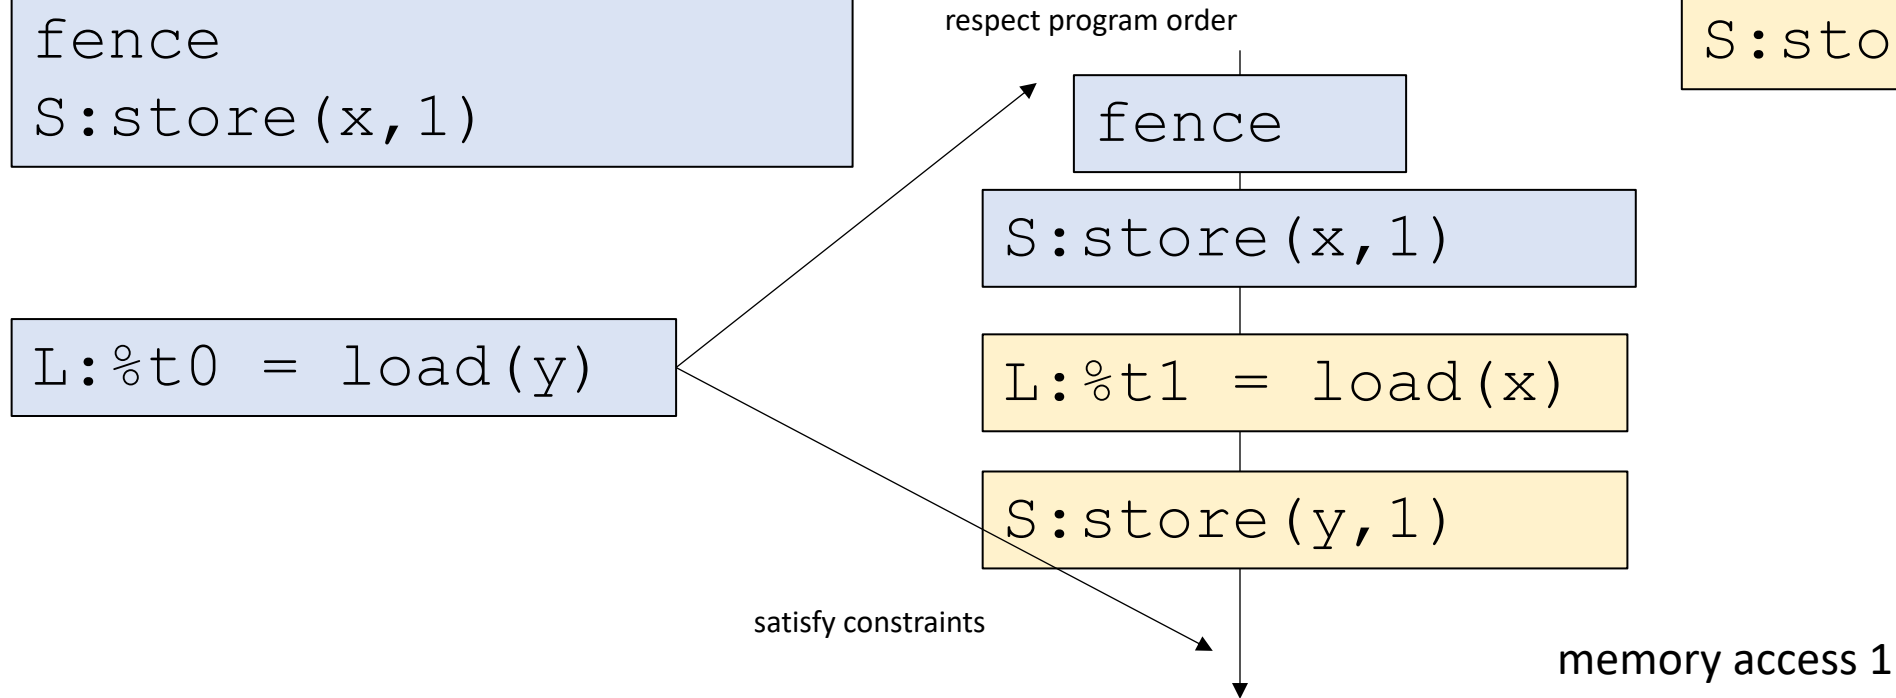
Get out our lego bricks

#### Thread 0:

```
L:%t0 = load(y)  
fence  
S:store(x, 1)
```

#### Thread 1:

```
L:%t1 = load(x)  
S:store(y, 1)
```



Now we cannot break program order past the fence!  
Are we done?

	memory access 0	
	L	S
L	YES	Different address
S	different address	Different address

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

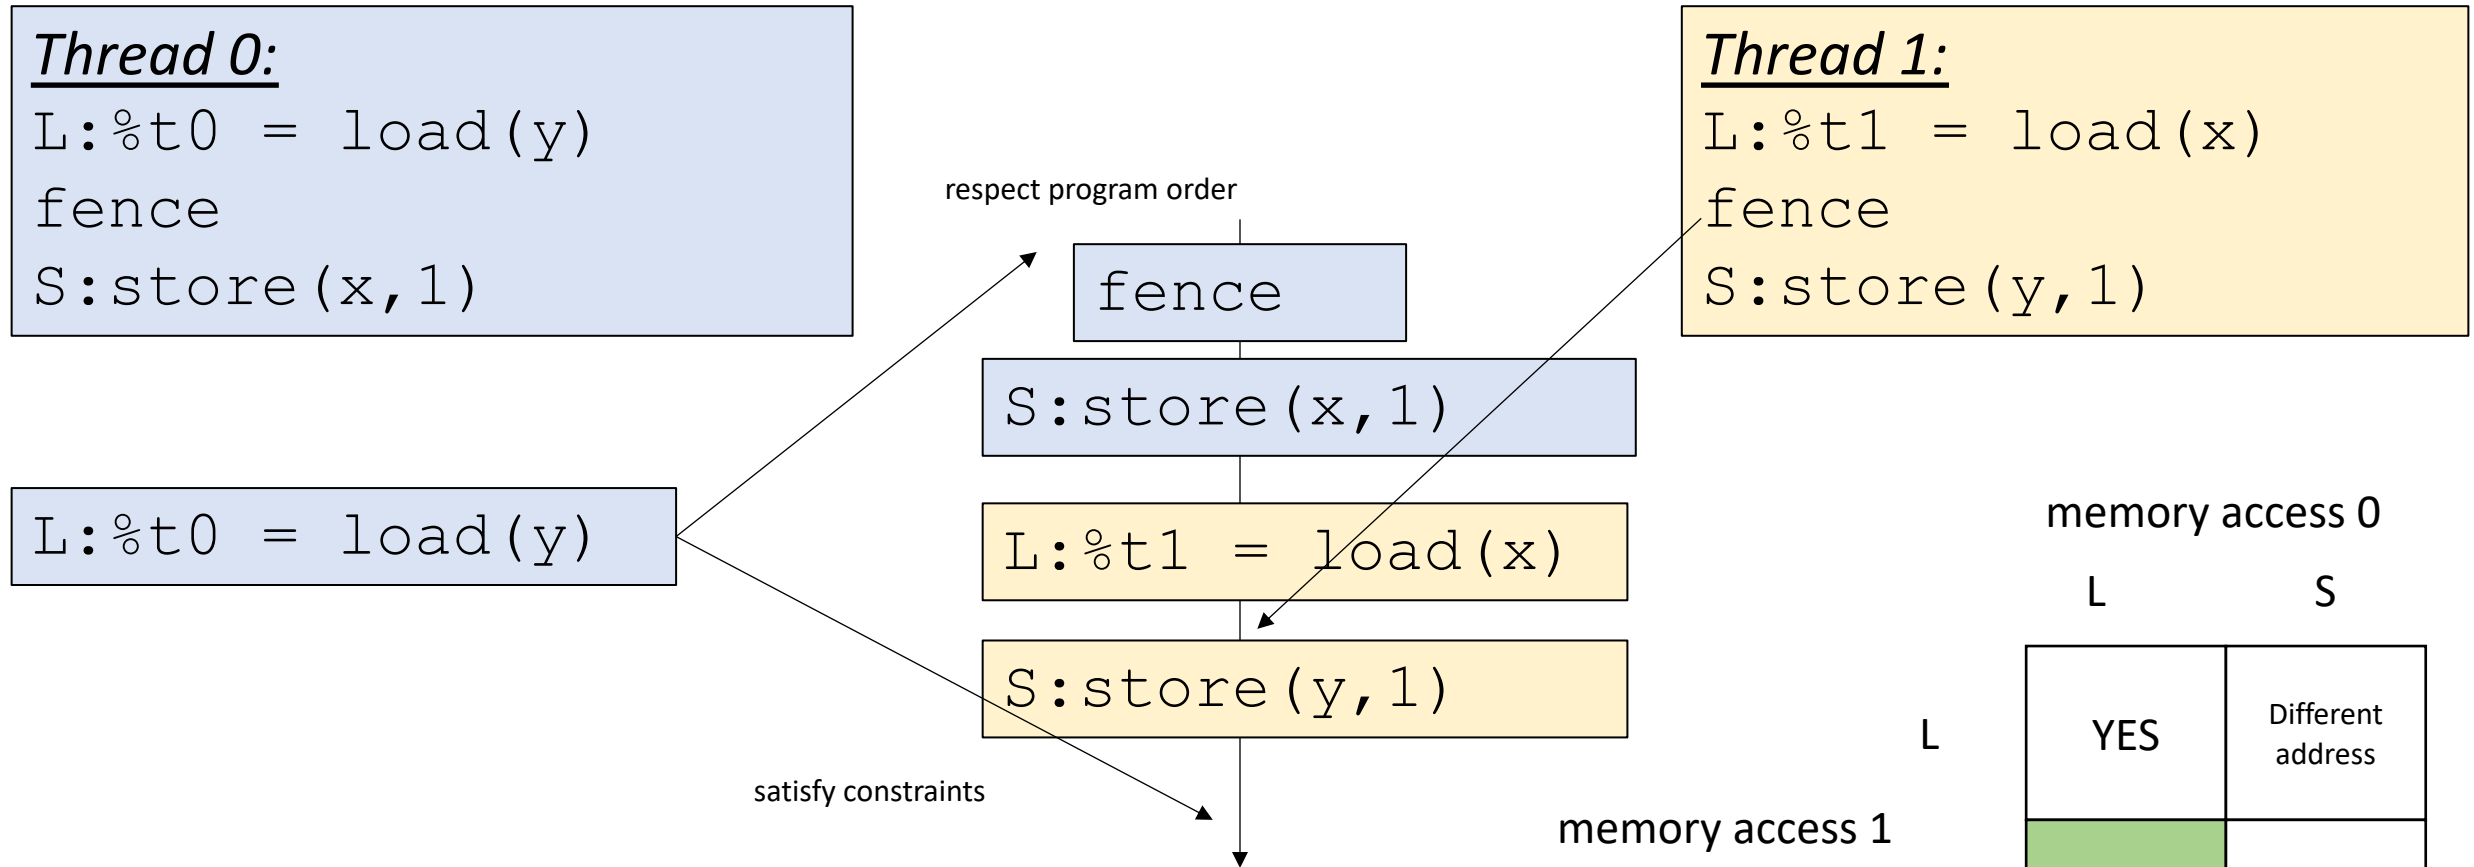
Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
fence  
S:store(x,1)
```

Thread 1:

```
L:%t1 = load(x)  
fence  
S:store(y,1)
```



	L	S
L	YES	Different address
S	different address	Different address

Now we cannot break program order past the fence!  
Are we done?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == t1 == 1`?

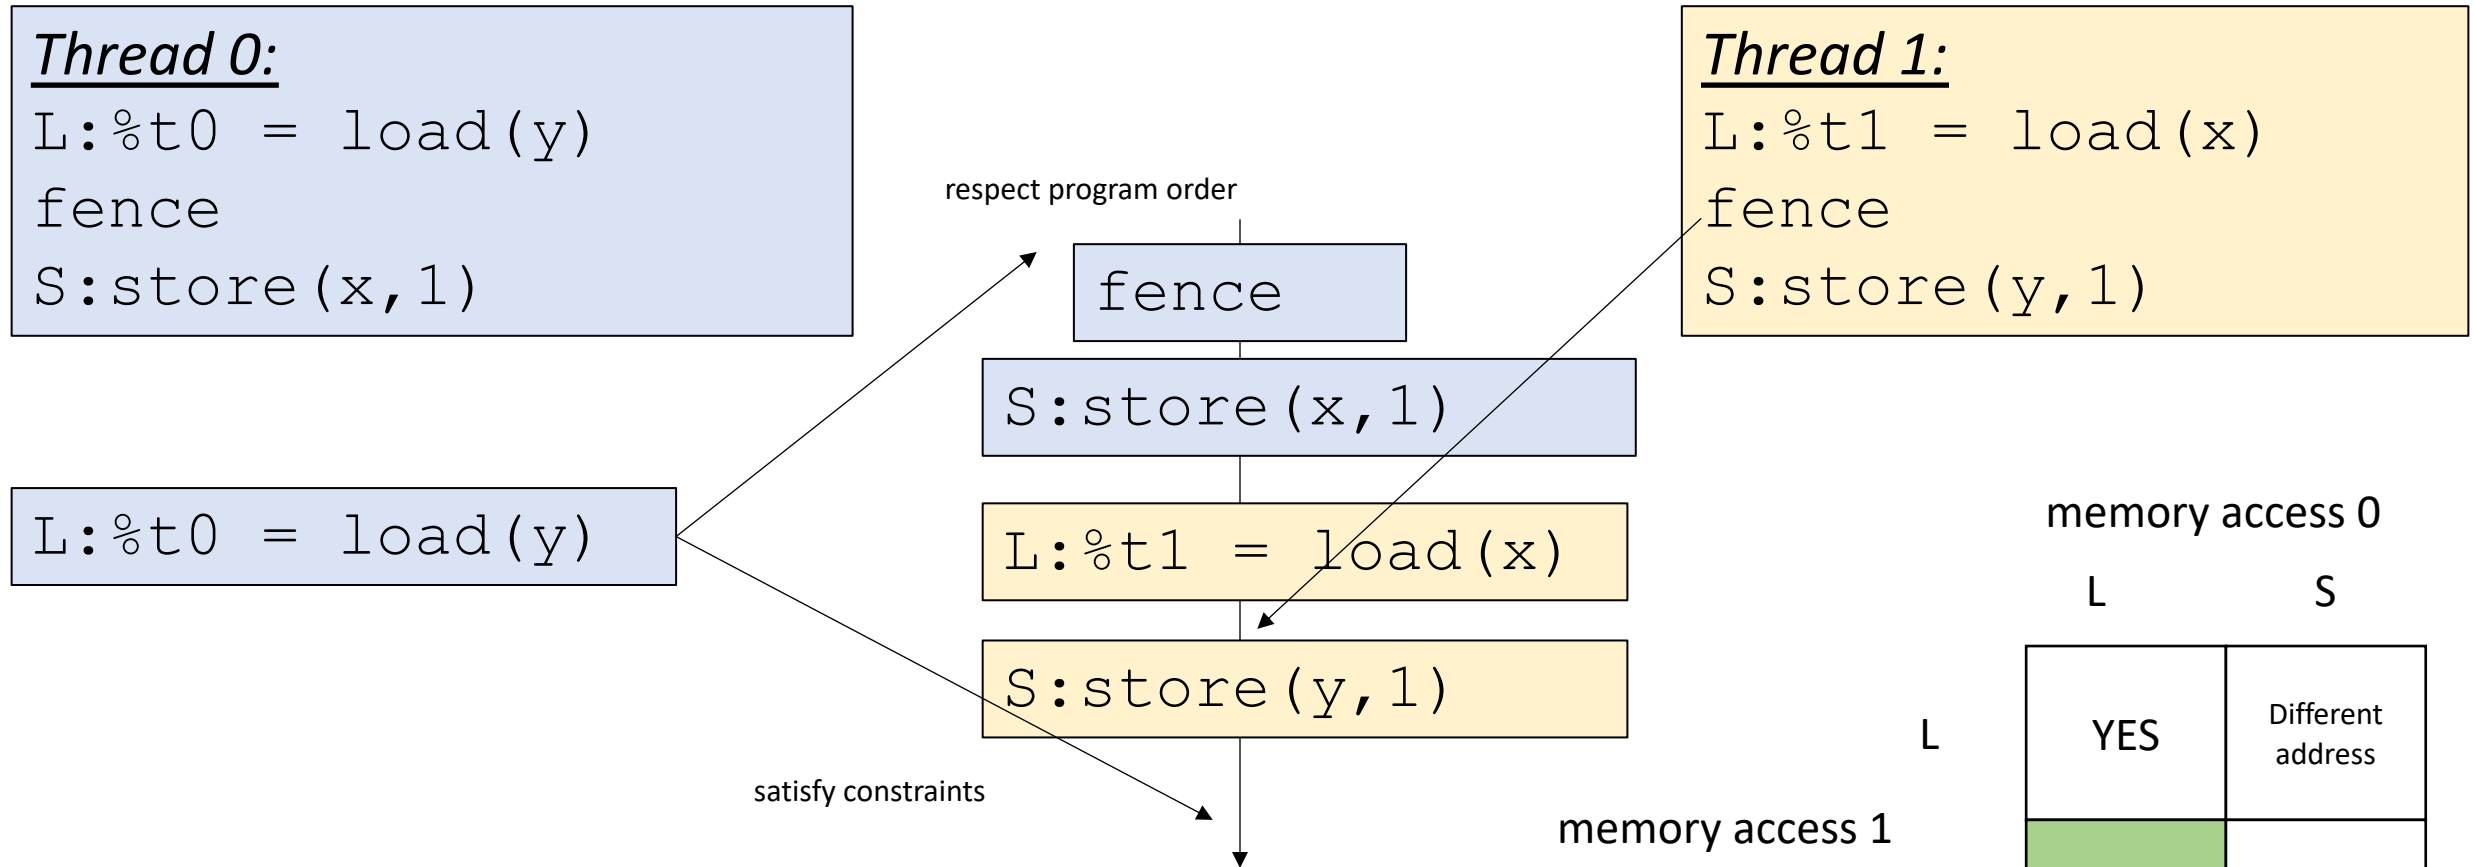
Get out our lego bricks

Thread 0:

```
L:%t0 = load(y)  
fence  
S:store(x,1)
```

Thread 1:

```
L:%t1 = load(x)  
fence  
S:store(y,1)
```



	L	S
L	YES	Different address
S	different address	Different address

Now we cannot break program order past the fence!  
Are we done? The behavior is no longer allowed

One more example

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

```
S:store(x,1)
```

```
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```

```
L:%t0 = load(y)
```

```
L:%t1 = load(x)
```



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Thread 0:


```
S:store(x,1)  
S:store(y,1)
```

```
S:store(x,1)
```

```
S:store(y,1)
```

Question: can `t0 == 1` and `t1 == 0`?

start off thinking  
about sequential  
consistency



Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```

```
L:%t0 = load(y)
```

```
L:%t1 = load(x)
```

Global variable:

```
int x[1] = {0};
```

```
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)
```

```
S:store(y,1)
```

start off thinking  
about sequential  
consistency

Thread 1:

```
L:%t0 = load(y)
```

```
S:%t1 = load(x)
```

```
S:store(x,1)
```

```
S:store(y,1)
```

```
L:%t0 = load(y)
```

```
L:%t1 = load(x)
```

respect program order

satisfy constraints



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```


Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

S

L	S
NO	Different address
NO	NO

What about TSO?



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```


Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

NO

Different  
address

S

NO

NO

What about TSO? NO

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

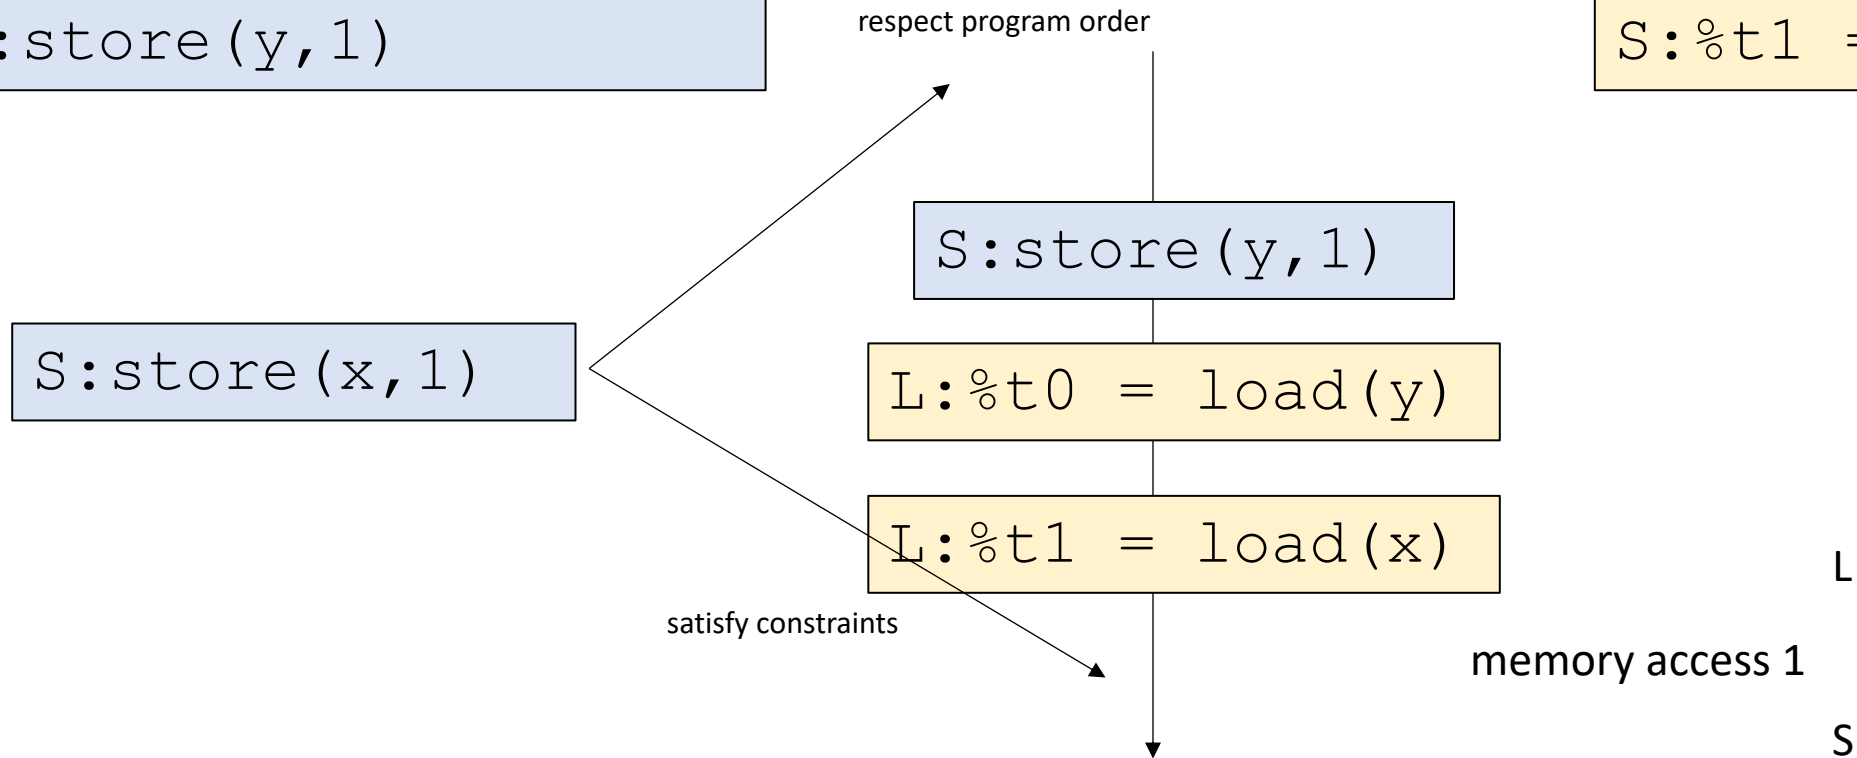
Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

NO

Different  
address

S

NO

Different  
address

memory access 1

What about PSO?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```


Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

NO

Different  
address

S

NO

Different  
address

What about PSO?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

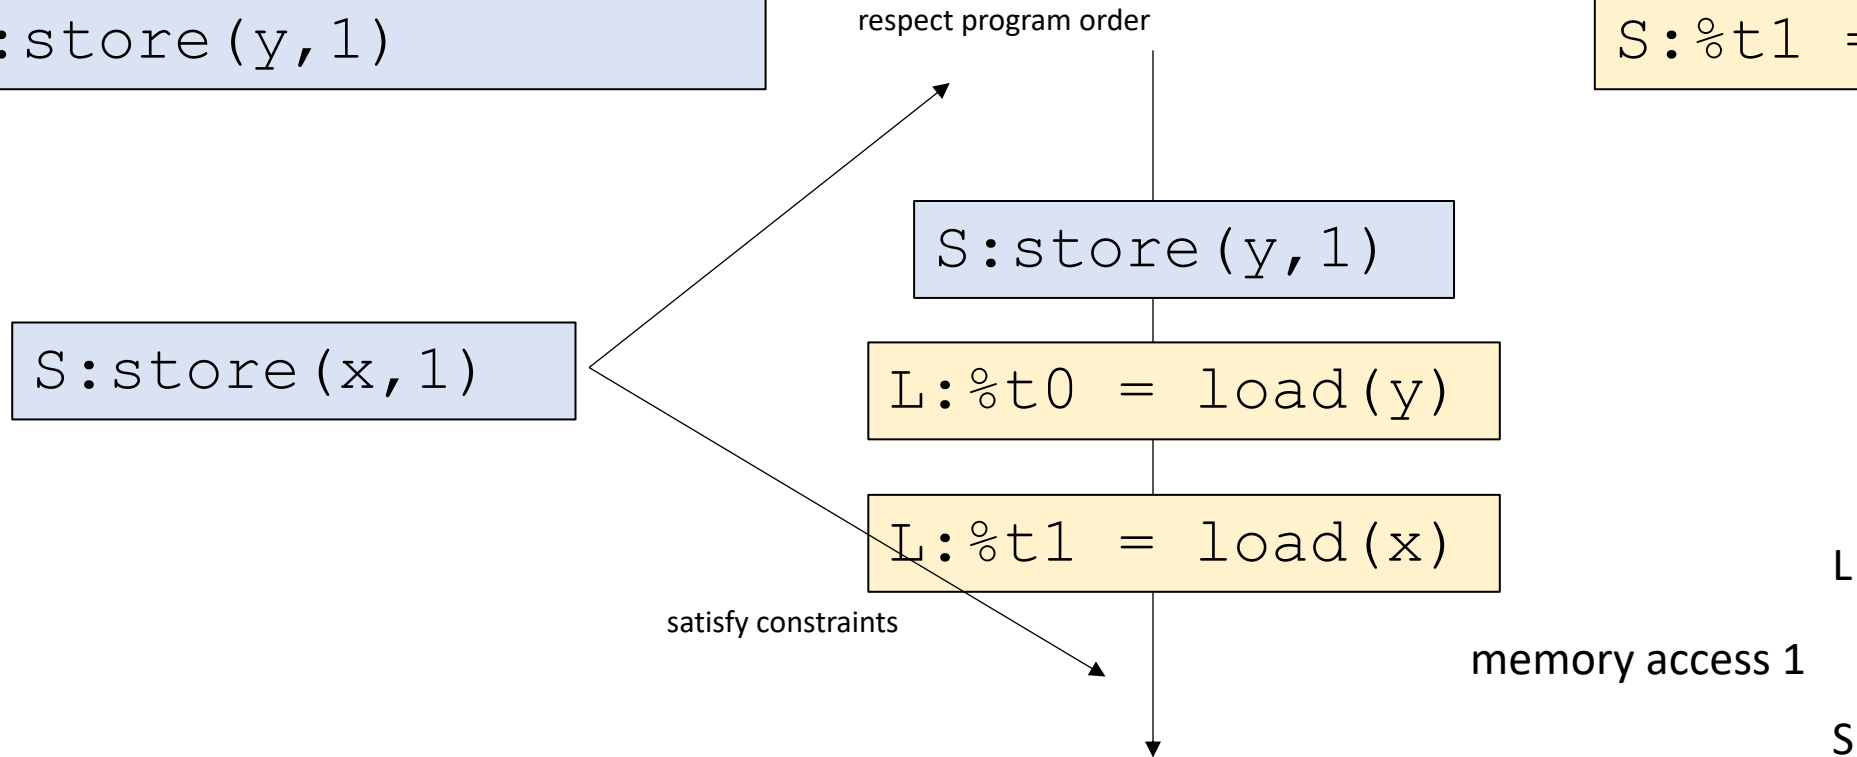
Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

NO

Different address

S

NO

Different address

What about PSO? YES

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

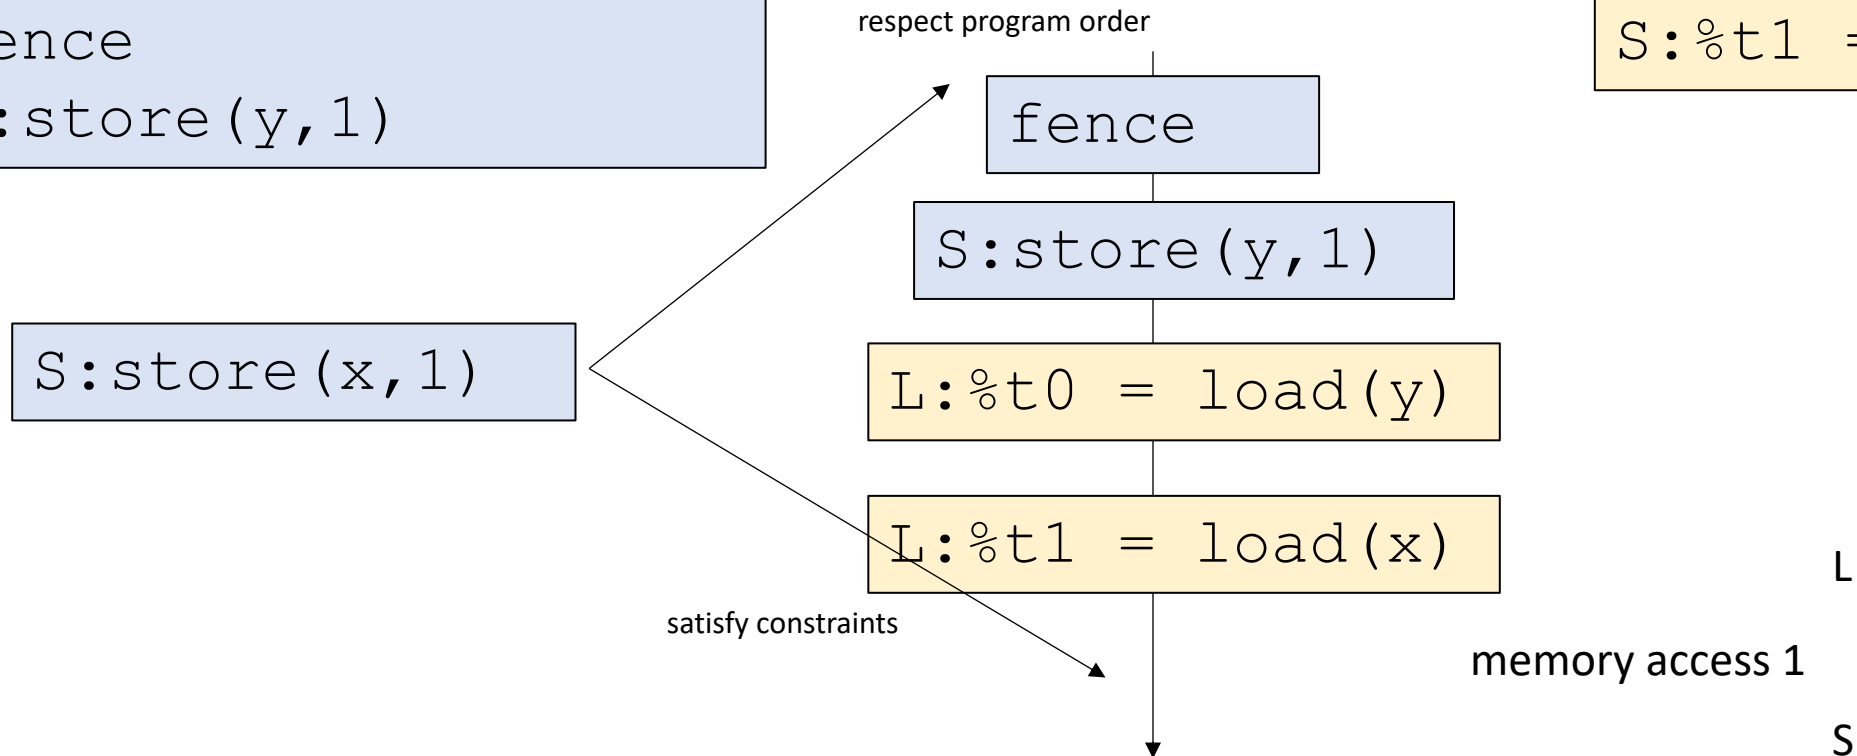
Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L

NO

Different address

S

NO

Different address

memory access 1

Now it is disallowed in PSO

### Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```


Question: can `t0 == 1` and `t1 == 0`?

#### Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```

#### Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```



memory access 0

L S

L	S
YES	Different address
Different address	Different address

What about RMO?

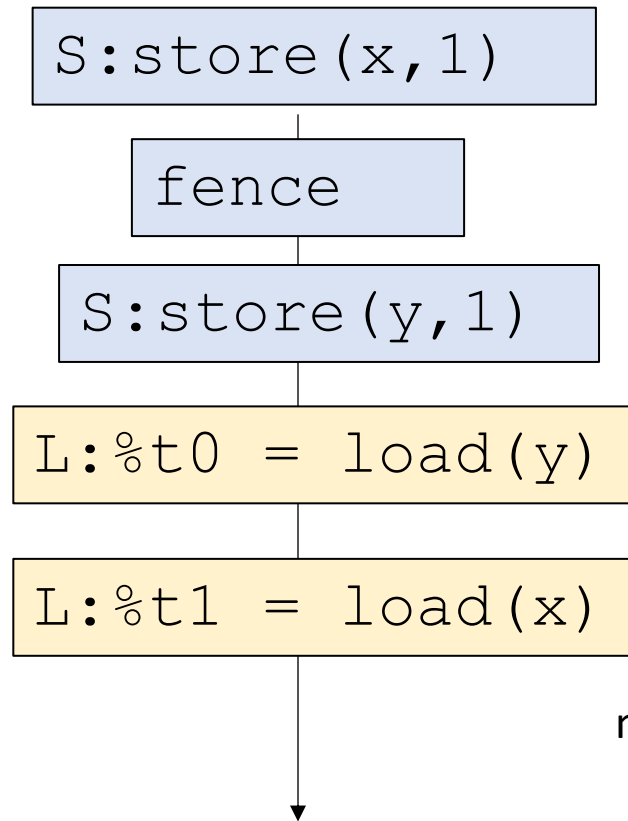
Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```



Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```

		memory access 0	
		L	S
memory access 1	L	YES	Different address
	S	Different address	Different address

What about RMO?

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```

```
L:%t1 = load(x)
```

```
S:store(x,1)
```

```
fence
```

```
S:store(y,1)
```

```
L:%t0 = load(y)
```

Thread 1:

```
L:%t0 = load(y)  
S:%t1 = load(x)
```

memory access 0

L S

	L	S
L	YES	Different address
S	Different address	Different address

memory access 1

What about RMO? The loads can be reordered also!



Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?


Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```

```
L:%t1 = load(x)
```

Thread 1:

```
L:%t0 = load(y)  
fence  
S:%t1 = load(x)
```



memory access 1

memory access 0

	L	S
L	YES	Different address
S	Different address	Different address

What about RMO? add a fence

Global variable:

```
int x[1] = {0};  
int y[1] = {0};
```

Question: can `t0 == 1` and `t1 == 0`?

Thread 0:

```
S:store(x,1)  
fence  
S:store(y,1)
```

S:store(x,1)

fence

S:store(y,1)

L:%t0 = load(y)

fence

L:%t1 = load(x)

memory access 1

Thread 1:

```
L:%t0 = load(y)  
fence  
S:%t1 = load(x)
```

memory access 0

L S

L

S

	L	S
L	YES	Different address
S	Different address	Different address

Now the relaxed behavior is disallowed

# Memory consistency in the real world

- Historic Chips:
  - X86: TSO
    - Surprising robust
    - mutexes and concurrent data structures generally seem to work
    - watch out for store buffering
  - IBM Power and ARM
    - Very relaxed. Similar to RMO with even more rules
    - Mutexes and data structures must be written with care
    - ARM recently strengthened theirs

# Memory consistency in the real world

- Historic Chips:
  - X86: TSO
    - Surprising robust
    - mutexes and concurrent data structures generally seem to work
    - watch out for store buffering
  - IBM Power and ARM
    - Very relaxed. Similar to RMO with even more rules
    - Mutexes and data structures must be written with care
    - ARM recently strengthened theirs

Companies have a history of providing insufficient documentation about their rules: academics have then gone and figured it out!

Getting better these days

# Memory consistency in the real world

- Modern Chips:
  - RISC-V : two specs: one similar to TSO, one similar to RMO
  - Apple M1: toggles between TSO and weaker
- Vulkan does not provide any fences that provide S - L ordering

# Memory consistency in the real world

- PSO and RMO were never implemented widely
  - I have not met anyone who knows of any RMO taped out chip
  - They are part of SPARC ISAs (i.e. RISC-V before it was cool)
  - These memory models might have been part of specialized chips
- Interestingly:
  - Early Nvidia GPUs appeared to informally implement RMO
- Other chips have very strange memory models:
  - Alpha DEC - basically no rules

# Where do programming languages fit in?

- One of the highest priorities of a programming language
  - Write once, run everywhere

# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language

C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

target machine

	L	S
L	?	?
S	?	?



# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language			
C++11 (sequential consistency)			
	L	S	
L	NO	NO	
S	NO	NO	

target machine			
TSO (x86)			
	L	S	
L	NO	different address	
S	NO	No	

# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language  
C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

find mismatch

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language  
C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

find mismatch

Two options:

make sure stores  
are not reordered  
with later loads

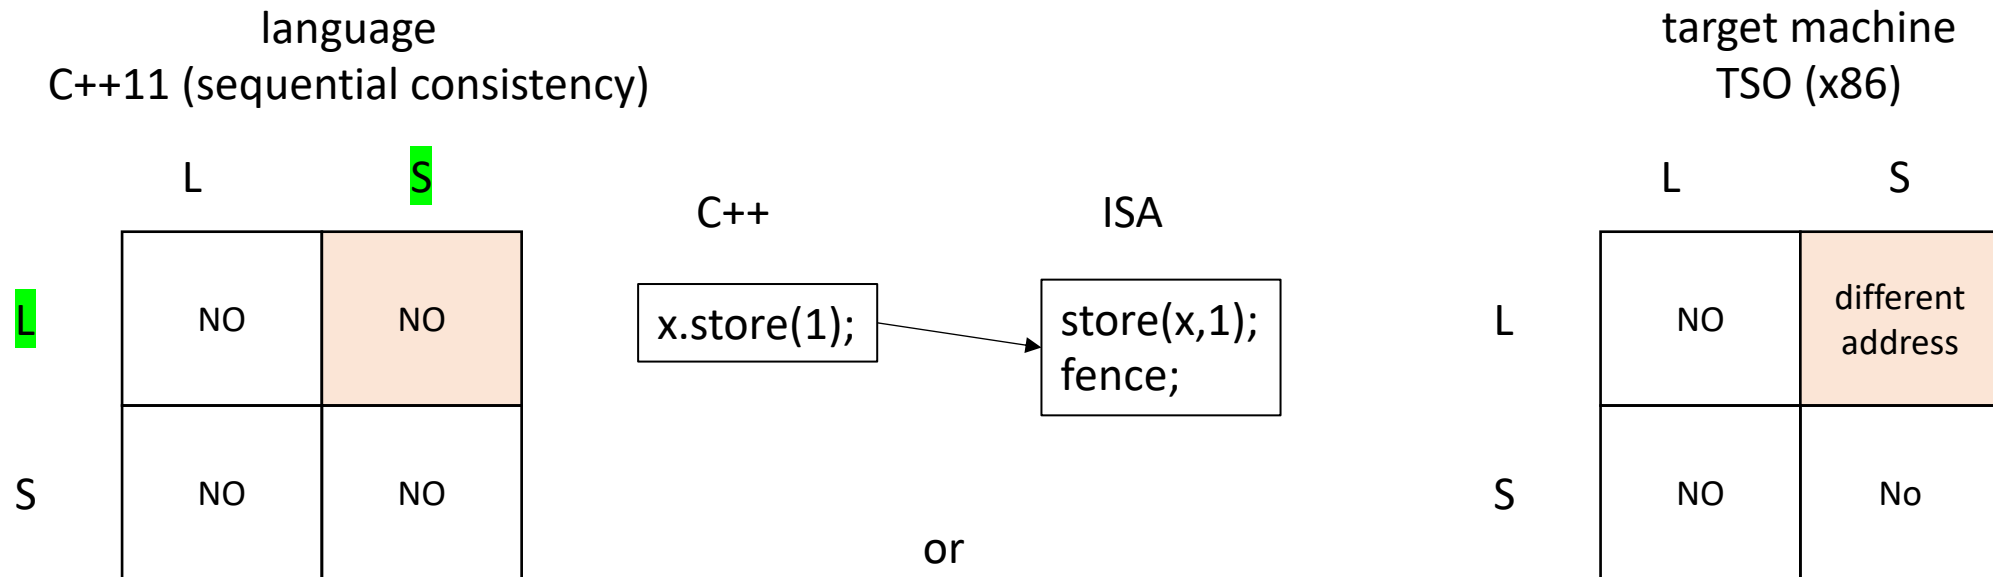
make sure loads  
are not reordered  
with earlier stores

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

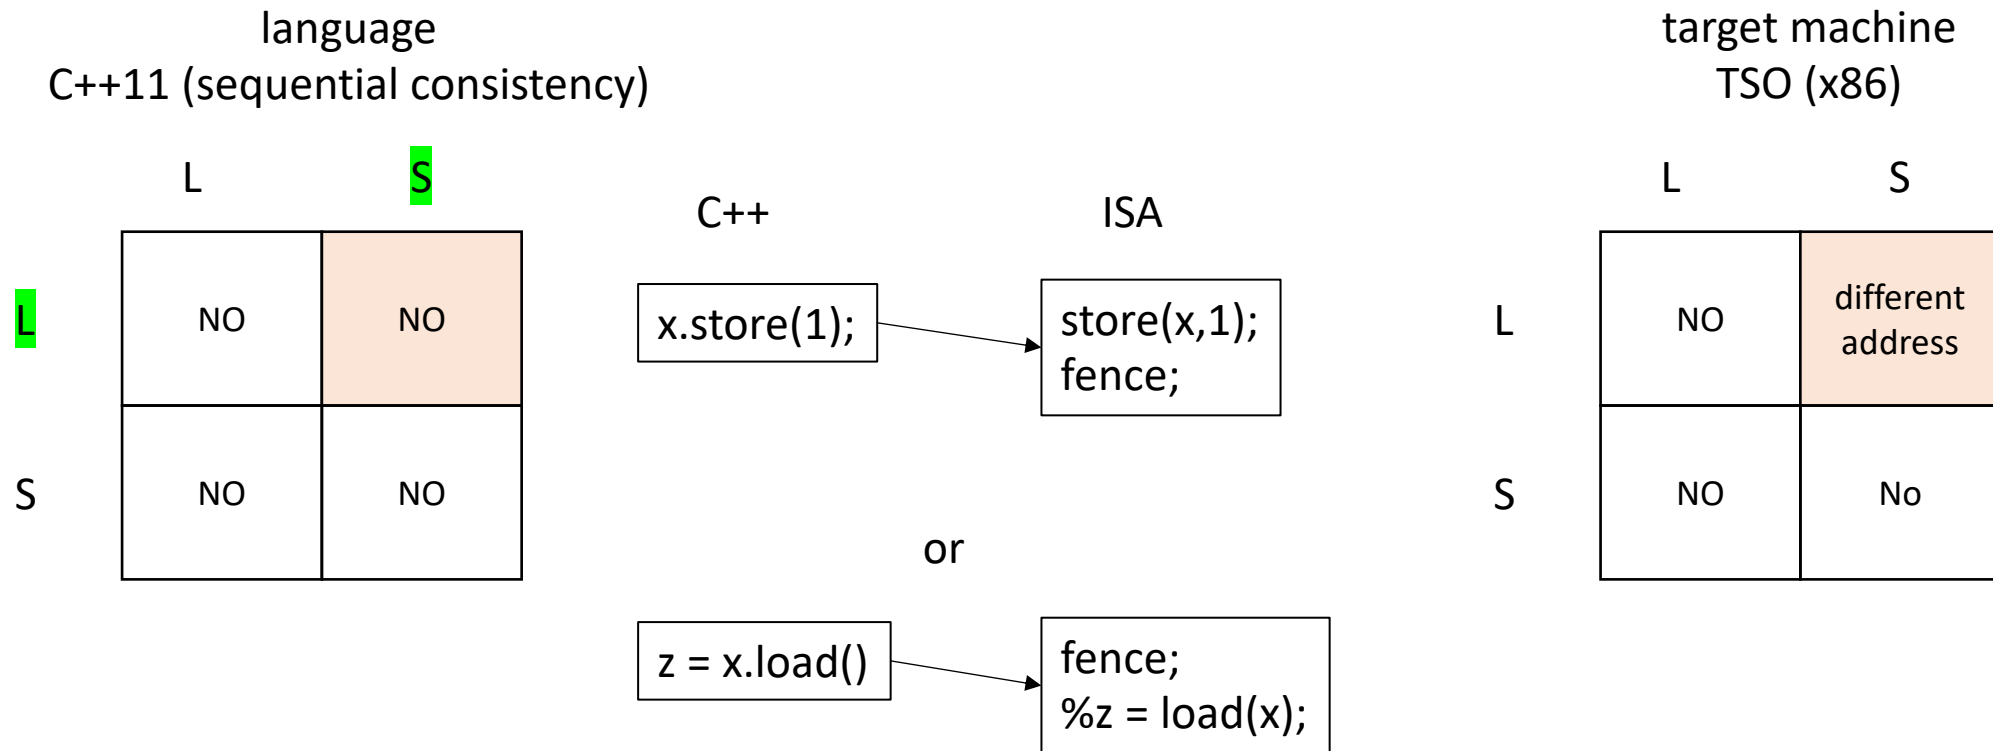
# C++11 atomic operation compilation

start with both both of the grids for the two different memory models



# C++11 atomic operation compilation

start with both both of the grids for the two different memory models



# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language  
C++11 (sequential consistency)

	L	<b>S</b>
<b>L</b>	NO	NO
S	NO	NO

C++

x.store(1);

ISA

store(x,1);  
fence;

or

z = x.load();

fence;  
%z = load(x);

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

*This should help you see why you want to reduce the number of atomic load/stores in your program*

# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

language  
C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

*How about this one?*

target machine  
PSO

	L	S
L	NO	different address
S	NO	different address

# C++11 atomic operation compilation

start with both both of the grids for the two different memory models

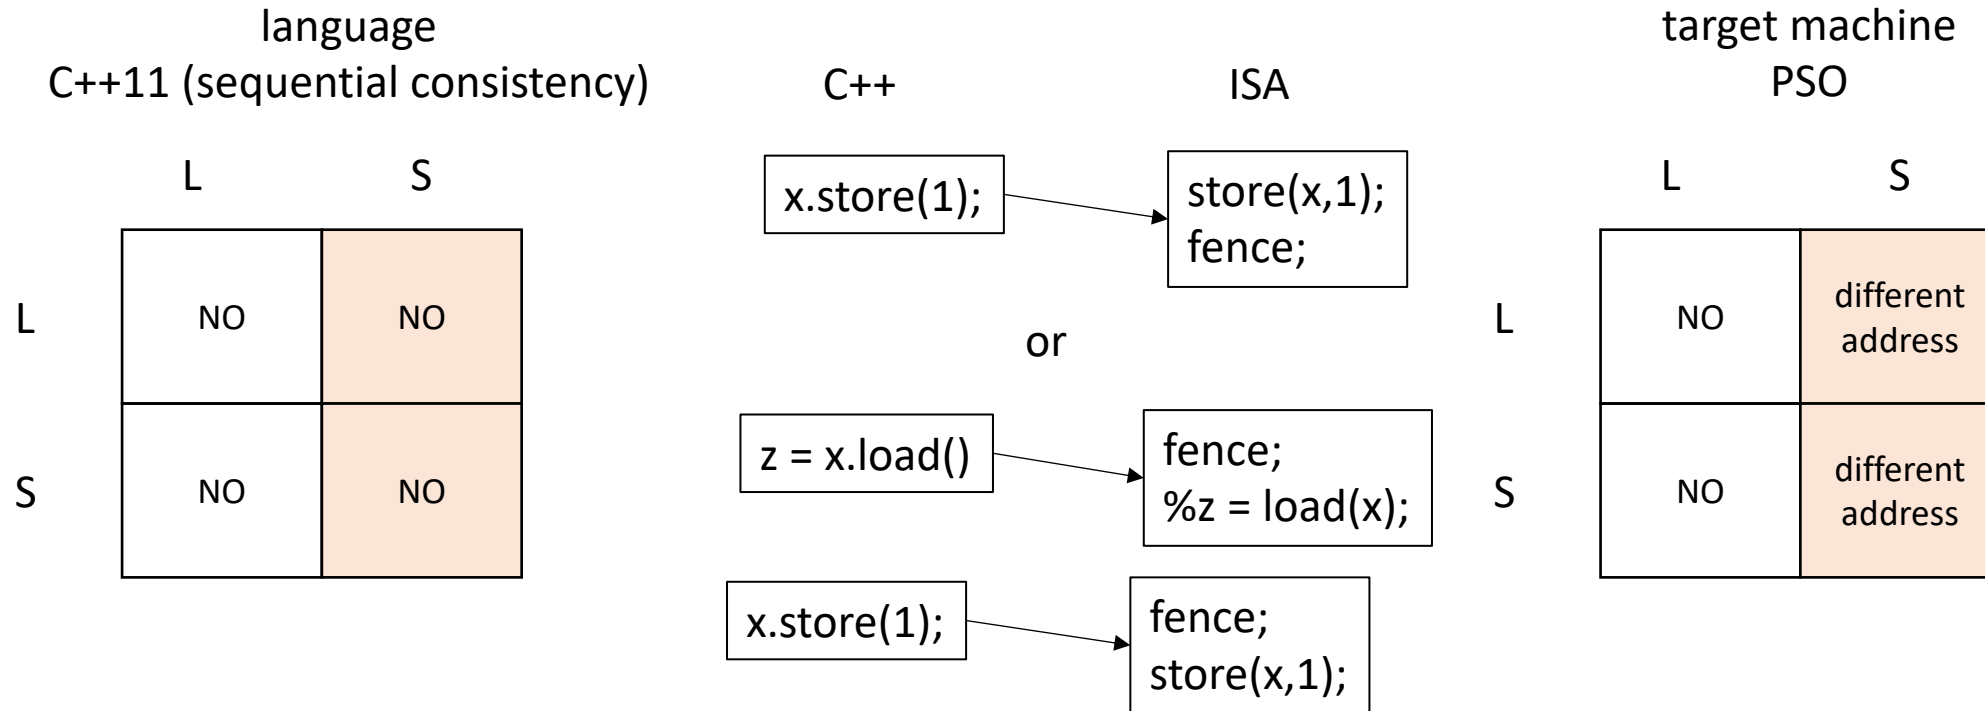
language			
C++11 (sequential consistency)			
	L	S	
L	NO	NO	
S	NO	NO	

target machine			
PSO			
	L	S	
L	NO	different address	
S	NO	different address	



# C++11 atomic operation compilation

start with both both of the grids for the two different memory models



# Memory orders

- Atomic operations take an additional “memory order” argument
  - `memory_order_seq_cst` - default
  - `memory_order_relaxed` - weakest

Where have we seen `memory_order_relaxed`?

# Relaxed memory order

language  
C++11 (sequential consistency)

	L	S
L	NO	NO
S	NO	NO

language  
C++11 (memory\_order\_relaxed)

	L	S
L	different address	different address
S	different address	different address

basically no orderings except for accesses to  
the same address

# Compiling memory order relaxed

language  
C++11 (memory\_order\_relaxed)

	L	S
L	different address	different address
S	different address	different address

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

# Compiling memory order relaxed

language  
C++11 (memory\_order\_relaxed)

	L	S
L	different address	different address
S	different address	different address

lots of mismatches!

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

# Compiling memory order relaxed

language  
C++11 (memory\_order\_relaxed)

	L	S
L	different address	different address
S	different address	different address

lots of mismatches!

But language is more relaxed than machine

*so no fences are needed*

target machine  
TSO (x86)

	L	S
L	NO	different address
S	NO	No

# Compiling memory order relaxed

*Do any of the ISA memory models need any fences for relaxed memory order?*

language  
C++11 (memory\_order\_relaxed)

	L	S
L	different address	different address
S	different address	different address

	L	S
L	NO	Different address
S	NO	NO

TSO

	L	S
L	NO	Different address
S	NO	Different address

PSO

	L	S
L	YES	Different address
S	Different address	Different address

RMO

# Memory order relaxed

- Very few use-cases! Be very careful when using it
  - Peeking at values (later accessed using a heavier memory order)
  - Counting (e.g. number of finished threads in work stealing)
  - ***DO NOT USE FOR QUEUE INDEXES***



# More memory orders: we will not discuss in class

- Atomic operations take an additional “memory order” argument
  - `memory_order_seq_cst` - default
  - `memory_order_relaxed` - weakest
- More memory orders (useful for mutex implementations):
  - `memory_order_acquire`
  - `memory_order_release`
- EVEN MORE memory orders (complicated: in most research it is omitted)
  - `memory_order_consume`

A cautionary tale

*Consider the following example: a graphics program where each thread wants to display a triangle;  
the display is a queue (not thread safe)*

**Thread 0:**

```
m.lock();  
display.enq(triangle0);  
m.unlock();
```

**Thread 1:**

```
m.lock();  
display.enq(triangle1);  
m.unlock();
```

*Consider the following example: a graphics program where each thread wants to display a triangle;  
the display is a queue (not thread safe)*

**Thread 0:**

```
m.lock();  
display.enq(triangle0);  
m.unlock();
```

**Thread 1:**

```
m.lock();  
display.enq(triangle1);  
m.unlock();
```

We know how lock and unlock are implemented

*Consider the following example: a graphics program where each thread wants to display a triangle;  
the display is a queue (not thread safe)*

**Thread 0:**

```
SPIN:CAS(mutex, 0, 1);  
display.enq(triangle0);  
store(mutex, 0);
```

**Thread 1:**

```
SPIN:CAS(mutex, 0, 1);  
display.enq(triangle1);  
store(mutex, 0);
```

We know how lock and unlock are implemented  
We also know how a queue is implemented

*Consider the following example: a graphics program where each thread wants to display a triangle;  
the display is a queue (not thread safe)*

**Thread 0:**

```
SPIN:CAS(mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

**Thread 1:**

```
SPIN:CAS(mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

We know how lock and unlock are implemented

We also know how a queue is implemented

What is an execution?

### Thread 0:

```
SPIN:CAS (mutex, 0, 1) ;  
%i = load(head) ;  
store(buffer+i, triangle0) ;  
store(head, %i+1) ;  
store(mutex, 0) ;
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1) ;  
%i = load(head) ;  
store(buffer+i, triangle1) ;  
store(head, %i+1) ;  
store(mutex, 0) ;
```

CAS (mutex, 0, 1) ;

*if blue goes first  
it gets to complete  
its critical section  
while thread 1 is spinning*



### Thread 0:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

CAS (mutex, 0, 1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex, 0);





### Thread 0:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

CAS (mutex, 0, 1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex, 0);

*now yellow gets a change to go*



### Thread 0:

```
SPIN:CAS(mutex,0,1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex,0);
```

### Thread 1:

```
SPIN:CAS(mutex,0,1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex,0);
```

*now yellow gets a change to go*

CAS(mutex,0,1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex,0);

CAS(mutex,0,1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex,0);



### Thread 0:


```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address



### Thread 0:


```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address



### Thread 0:


```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address




What just happened if this store moves?


# Nvidia in 2015

- Nvidia architects implemented a weak memory model
- Nvidia programmers expected a strong memory model
- Mutexes implemented without fences!

# Nvidia in 2015



(a)




(b)



(c)




(d)



bug found in two  
Nvidia textbooks

We implemented  
a side-channel attack  
that made the bugs  
appear more frequently

These days Nvidia has  
a very well-specified  
memory model!



### Thread 0:

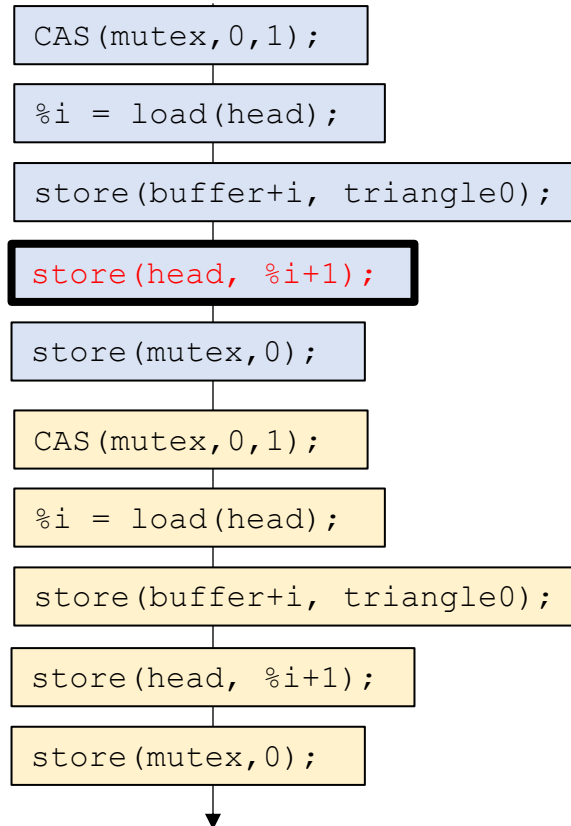
```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
store(mutex, 0);
```

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
store(mutex, 0);
```

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address



How to fix the issue?



### Thread 0:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
fence;  
store(mutex, 0);
```

*unlock contains fence  
before store!*

### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
fence;  
store(mutex, 0);
```

*unlock contains fence  
before store!*

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address

CAS (mutex, 0, 1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex, 0);

CAS (mutex, 0, 1);

%i = load(head);

store(buffer+i, triangle0);

store(head, %i+1);

store(mutex, 0);

How to fix the issue?

your unlock function  
should contain a fence!

### Thread 0:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle0);  
store(head, %i+1);  
fence;  
store(mutex, 0);
```

*unlock contains fence  
before store!*

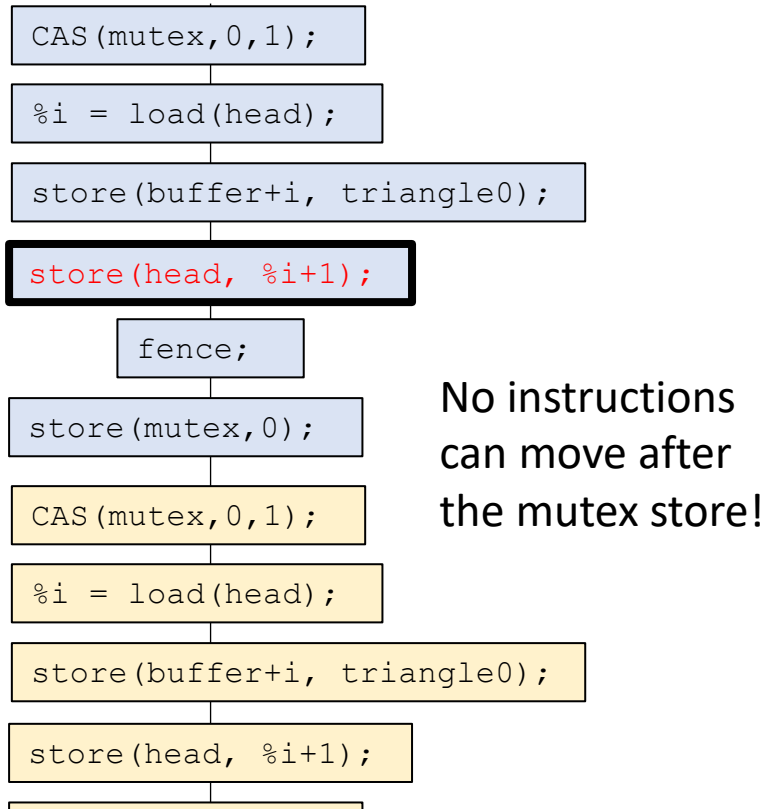
### Thread 1:

```
SPIN:CAS (mutex, 0, 1);  
%i = load(head);  
store(buffer+i, triangle1);  
store(head, %i+1);  
fence;  
store(mutex, 0);
```

*unlock contains fence  
before store!*

*what can happen in a PSO  
memory model?*

	L	S
L	NO	Different address
S	NO	Different address



How to fix the issue?

your unlock function  
should contain a fence!

No instructions  
can move after  
the mutex store!

# Memory Model Strength

- If one memory model M0 allows more relaxed behaviors than another memory model M1, then M0 is more *relaxed* (or *weaker*) than M1.
- It is safe to run a program written for M0 on M1. But not vice versa

	L	S
L	NO	Different address
S	NO	NO

TSO

	L	S
L	NO	Different address
S	NO	Different address

PSO

	L	S
L	YES	Different address
S	Different address	Different address

RMO

# Memory Model Strength

- Many times specifications are weaker than implementations:
  - A chip might document PSO, but implement TSO:
    - Why?

	L	S
L	NO	Different address
S	NO	NO

TSO

	L	S
L	NO	Different address
S	NO	Different address

PSO

	L	S
L	YES	Different address
S	Different address	Different address

RMO

See you all next Monday!